

# Concurrency: Foundations and Algorithms

## Mutual Exclusion



**Prof. P. Felber**

[Pascal.Felber@unine.ch](mailto:Pascal.Felber@unine.ch)

<http://iiun.unine.ch/>


*Based on slides by Maurice Herlihy and Nir Shavit*




# Review

- Model of computation
- Asynchrony
- Mutual exclusion

# Concurrent Computation



# Asynchrony




- Sudden unpredictable delays
  - Cache misses (*short*)
  - Page faults (*long*)
  - Scheduling quantum used up (*really long*)

# Parallel Primality Testing

- Challenge
  - Print primes from 1 to  $10^{10}$
- Given
  - Ten-processor multiprocessor
  - One thread per processor
- Goal
  - Get ten-fold speedup (or close)

# Shared Counter



Each thread  
takes a number

If it's a prime, print it, otherwise get another number

# Procedure for Thread *i*

```

Counter counter = new Counter(1);


void primePrint {
    while (true) {
        long j = counter.getAndIncrement();
        if (j > 1010)
            break;
        if (isPrime(j))
            print(j);
    }
}

```

# Where Things Reside

```
Counter counter = new Counter(1);
void primePrint {
  while (true) {
    long j = counter.getAndIncrement();
    if (j > 1010)
      break;
    if (isPrime(j))
      print(j);
  }
}
```

Code





# Challenge

```
public class Counter {
    private long value;

    public long getAndIncrement() {
        long temp = value;
        value = temp + 1;
        return temp;
    }
}
```

Make these steps  
*atomic* (indivisible)

# Mutual Exclusion



- Today we will try to formalize our understanding of mutual exclusion
- We will also use the opportunity to show you how to argue about and prove various properties in an asynchronous concurrent setting

# Mutual Exclusion



- Formal problem definitions
- Solutions for 2 threads
- Solutions for  $n$  threads
- Fair solutions
- Inherent costs

# Warning

- You will never use these protocols
  - Get over it
- You are advised to understand them
  - The same issues show up everywhere
  - Except hidden and more complex

# Why is Concurrent Programming so Hard?

- Try preparing a seven-course banquet
  - By yourself
  - With one friend
  - With twenty-seven friends...
  
- Before we can talk about programs
  - Need a language
  - Describing time and concurrency

# What is Time?

- *“Absolute, true and mathematical time, of itself and from its own nature, flows equably without relation to anything external.” (I. Newton, 1689)*
- *“Time is, like, Nature’s way of making sure that everything doesn’t happen all at once.” (Anonymous, circa 1968)*



Time


# Events

- An **event**  $a_0$  of thread A is
  - Instantaneous
  - No simultaneous events (break ties)



# Threads

- A **thread**  $A$  is (formally) a sequence  $a_0, a_1, \dots$  of events  
processor instructions. read/write/calculate....
- “Trace” model
- Notation:  $a_0 \rightarrow a_1$  indicates order  
precedence






# Example Thread Events

- Assign to shared variable
- Assign to local variable
- Invoke method
- Return from method
- Lots of other things...

# Threads are State Machines

Events are transitions



# States

Data associated with the thread

- Thread state
  - Program counter
  - Local variables

Composed of many threads

- System state
  - Object fields (shared variables)
  - Union of thread states

# Concurrency

- Thread A




- Thread B




# Interleavings

- Events of two or more threads
  - Interleaved
  - Not necessarily independent (why?)




# Intervals

- An **interval**  $A_0=(a_0, a_1)$  is
  - Time between events  $a_0$  and  $a_1$




# Intervals may Overlap

Nonempty intersection



# Intervals may be Disjoint


- Precedence
  - Interval  $A_0$  **precedes** interval  $B_0$






# Precedence

- Notation:  $A_0 \rightarrow B_0$  A0 happens before B0 when the last event of A happens before the first event of B
- Formally:
  - End event of  $A_0$  before start event of  $B_0$
  - Also called “happens before” or “precedes”




# Precedence Ordering

- Remark:  $A_0 \rightarrow B_0$  is just like saying
  - 1066 AD  $\rightarrow$  1492 AD
  - Middle Ages  $\rightarrow$  Renaissance
- Oh wait, what about this week vs. this month?



# Precedence Ordering

- Never true that  $A \rightarrow A$
- If  $A \rightarrow B$  then not true that  $B \rightarrow A$
- If  $A \rightarrow B$  **and**  $B \rightarrow C$  then  $A \rightarrow C$
- Funny thing:  $A \rightarrow B$  **and**  $B \rightarrow A$  might both be false!



# Partial Orders

(you may know this already)

- **Irreflexive**
  - Never true that  $A \rightarrow A$
  
- **Antisymmetric**
  - If  $A \rightarrow B$  then not true that  $B \rightarrow A$
  
- **Transitive**
  - If  $A \rightarrow B$  and  $B \rightarrow C$  then  $A \rightarrow C$

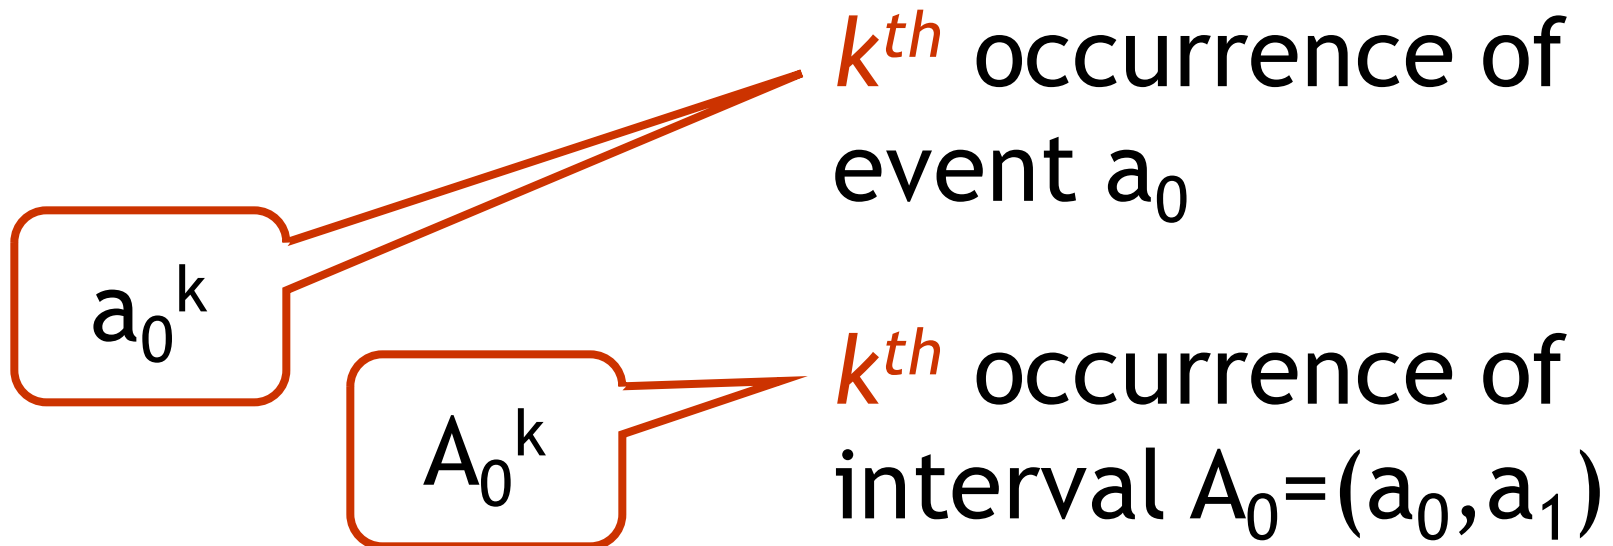
# Total Orders

(you may know this already)

- Also
  - Irreflexive
  - Antisymmetric
  - Transitive
  
- Except that for every distinct A, B
  - Either  $A \rightarrow B$  or  $B \rightarrow A$

# Repeated Events

```
while (mumble) {
    a0; a1;
}
```



# Implementing a Counter

```

public class Counter {
    private long value;

    public long getAndIncrement() {
        long temp = value;
        value = temp + 1;
        return temp;
    }
}

```

Make these steps  
*atomic* (indivisible)

# Locks (Mutual Exclusion)

```

public interface Lock {
    public void lock();
    public void unlock();
}
  
```

Acquire lock

Release lock

Class that implements this Interface will ensure mutual exclusion



# Using Locks

```

public class Counter {
    private long value;
    private Lock lock;
    public long getAndIncrement() {
        int temp;
        lock.lock();
        try {
            temp = value;
            value = value + 1;
        } finally {
            lock.unlock();
        }
        return temp;
    }
}
  
```

Acquire lock

Critical section



Release lock  
(no matter what)


# Using Locks

```
private Lock lock;
...
while{mumble}{
    <non critical section>
    lock.lock();
    <critical section>
    lock.unlock();
}
...
```

- Only **well formed executions!**
- What are the required properties of locks?

# Mutual Exclusion

- Let  $CS_i^k$   be thread  $i$ 's  $k^{\text{th}}$  critical section execution
- And  $CS_j^m$   thread  $j$ 's  $m^{\text{th}}$  critical section execution
- Then either



# Deadlock-Free



- If some thread calls **lock()**
  - And never returns
  - Then other threads must complete **lock()** and **unlock()** calls infinitely often
- System as a whole makes progress
  - Even if individuals starve

# Lockout-Free



- If some thread calls **lock()**
  - It will eventually return
- Individual threads make progress

# Two-Thread vs. $n$ -Thread

- Two-thread solutions first
  - Illustrate most basic ideas Read/write shared variables (registers)
  - Fits on one slide
- Then  $n$ -thread solutions

# Two-Thread Conventions

```

class ... implements Lock {
    ...
    // Thread-local index, 0 or 1
    public void lock() {
        int i = ThreadID.get();
        int j = 1 - i;
    }
}

```

Henceforth: *i* is current thread, *j* is other thread

# LockOne

For 2 threads

volatile: will disable reordering/optimizations from compiler to guarantee order. use it when dealing with shared data.

```

class LockOne implements Lock {
    private volatile boolean[]
        flag = new boolean[2];
    public void lock() {
        flag[i] = true;
        while (flag[j]) {}
    }
    public void unlock() {
        flag[i] = false;
    }
}
  
```

Set my flag

Wait for other flag to go false

No longer interested



# LockOne Satisfies Mutual Exclusion

- Assume  $CS_A^j$  overlaps  $CS_B^k$
- Consider each thread's ( $j^{\text{th}}$  and  $k^{\text{th}}$ ) last read and write in the **lock()** method before entering
- Derive a contradiction

# From the Code

$\text{write}_A(\text{flag}[A]=\text{true}) \rightarrow \text{read}_A(\text{flag}[B]==\text{false}) \rightarrow \text{CS}_A$

$\text{write}_B(\text{flag}[B]=\text{true}) \rightarrow \text{read}_B(\text{flag}[A]==\text{false}) \rightarrow \text{CS}_B$

```
class LockOne implements Lock {
  private volatile boolean[] flag = new boolean[2];
  public void lock() {
    flag[i] = true;
    while (flag[j]) {}
  }
}
```

# From the Assumption

$\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow \text{write}_B(\text{flag}[B] = \text{true})$

$\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow \text{write}_A(\text{flag}[A] = \text{true})$

# Combining

## Assumptions

$\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow \text{write}_B(\text{flag}[B] = \text{true})$

$\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow \text{write}_A(\text{flag}[A] = \text{true})$

## From the code


$\text{write}_A(\text{flag}[A] = \text{true}) \rightarrow \text{read}_A(\text{flag}[B] == \text{false})$

$\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{read}_B(\text{flag}[A] == \text{false})$

**Cycle!**

# Deadlock Freedom

- LockOne fails deadlock-freedom
  - Concurrent execution can deadlock

Time 	<pre>flag[i] = true; while (flag[j]) {}</pre>	<pre>flag[j] = true; while (flag[i]) {}</pre>
--	---	---

- Sequential executions OK

# LockTwo

```

public class LockTwo implements Lock {
    private volatile int victim;
    public void lock() {
        victim = i;
        while (victim == i) {}
    }

    public void unlock() {}
}
  
```

Let other go first

Wait for permission

Nothing to do

# LockTwo Claims

- Satisfies mutual exclusion
  - If thread **i** in CS
  - Then **victim == j**
  - Cannot be both **i** and **j**
- Not deadlock free
  - Sequential execution deadlocks
  - Concurrent execution does not

```

{
    victim = i;
    while (victim == i) {}
}

```

# Peterson's Algorithm

Combines flag and victim

```

public void lock() {
    flag[i] = true;
    victim = i;
    while (flag[j] && victim == i) {}
}
public void unlock() {
    flag[i] = false;
}
  
```

Announce I'm interested

Defer to other

Wait while other interested & I'm the victim

No longer interested



# Mutual Exclusion

```
public void lock() {
    flag[i] = true;
    victim = i;
    while (flag[j] && victim == i) {}
}
```

- If thread **0** in critical section
  - **flag[0] == true**
  - **victim == 1**
- If thread **1** in critical section
  - **flag[1] == true**
  - **victim == 0**

**Cannot both be true**

# Deadlock Free

```
public void lock() {
    ...
    while (flag[j] && victim == i) {}
}
```

- Thread blocked
  - Only at **while** loop
  - Only if it is the victim
- One or the other must not be the victim

# Lockout Free

```


public void lock() {
    flag[i] = true;
    victim = i;
    while (flag[j] && victim == i) {}
}
public void unlock() {
    flag[i] = false;
}

```

- Thread **i** blocked only if **j** repeatedly re-enters with **flag[j] == true** and **victim == i**
- When **j** re-enters, it sets **victim** to **j**, so **i** gets in

# The Filter Algorithm for $n$ Threads

- There are  $n-1$  “waiting rooms” called levels
- At each level
  - At least one enters level
  - At least one blocked if many try
- Only one thread makes it through




# Filter

```

class Filter implements Lock {
    volatile int[] level;    // level[i] for thread i
    volatile int[] victim;  // victim[L] for level L

    public Filter(int n) {
        level = new int[n];
        victim = new int[n];
        for (int i = 0; i < n; i++) {
            level[i] = 0;
        }
    }
    ...
}
  
```



# Filter

```
class Filter implements Lock {
```

```
...
```

```
public void lock() {
```

One level at a time

```
for (int L = 1; L < n; L++) {
```

Announce intention  
to enter level L

```
level[i] = L;
```

```
victim[L] = i;
```

Give priority to anyone but me

```
while ((∃k != i | level[k] >= L) &&  
victim[L] == i) {}
```

Wait as  
long as


someone else is at same or higher level and I'm  
designated victim (enter level L when loop  
complete)

```
public void unlock() { level[i] = 0; }
```

```
}
```

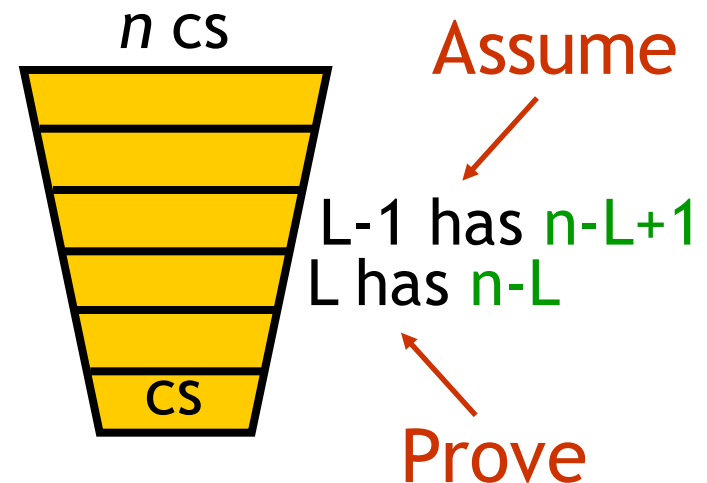
# Claim

- Start at level  $L=0$
- At most  $n-L$  threads enter level  $L$
- Mutual exclusion at level  $L=n-1$



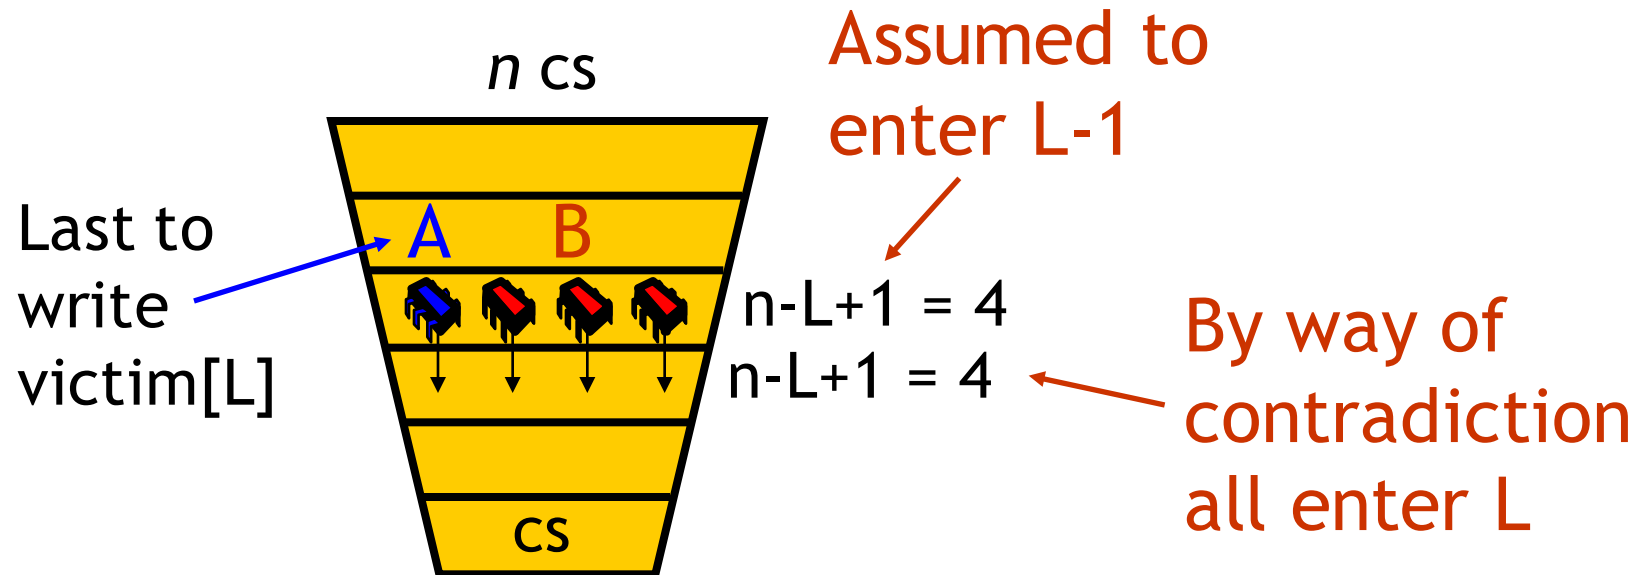
# Induction Hypothesis

- No more than  $n-L+1$  threads at level  $L-1$
- Induction step: by contradiction
- Assume all at level  $L-1$  enter level  $L$
- A last to write `victim[L]`
- B is any other thread at level  $L$





# Proof Structure



- Show that A must have seen B at level  $L$  and since  $\text{victim}[L] == A$  could not have entered

# Observations

1.  $\text{write}_B(\text{level}[B]=L) \rightarrow \text{write}_B(\text{victim}[L]=B)$  (code)
2.  $\text{write}_A(\text{victim}[L]=A) \rightarrow \text{read}_A(\text{level}[B])$  (code)
3.  $\text{write}_B(\text{victim}[L]=B) \rightarrow \text{write}_A(\text{victim}[L]=A)$   
 (by assumption: A is the last thread to write victim[L])

```

public void lock() {
    for (int L = 1; L < n; L++) {
        level[i] = L;
        victim[L] = i;
        while ((∃k != i | level[k] >= L) && victim[L] == i) {}
    }
}
  
```

# Combining Observations

1.  $\text{write}_B(\text{level}[B]=L) \rightarrow \text{write}_B(\text{victim}[L]=B)$  (code)
3.  $\text{write}_B(\text{victim}[L]=B) \rightarrow \text{write}_A(\text{victim}[L]=A)$  (ass.)
2.  $\text{write}_A(\text{victim}[L]=A) \rightarrow \text{read}_A(\text{level}[B])$  (code)

Thus,  $A$  read  $\text{level}[B] \geq L$ ,  $A$  was last to write  $\text{victim}[L]$ , so it could not have entered level  $L$ !

```
public void lock() {
    for (int L = 1; L < n; L++) {
        level[i] = L;
        victim[L] = i;
        while ((∃k != i | level[k] >= L) && victim[L] == i) {}
    }
}
```

# No Lockout

- Filter Lock satisfies “no lockout” properties
  - Just like Peterson algorithm at any level
  - So no one starves (no lockout)
- But what about fairness?
  - Threads can be overtaken by others

# Bounded Waiting

- Want stronger fairness guarantees
- Thread not “overtaken” too much
- Need to adjust definitions...
  
- Divide **lock()** method into 2 parts
  - Doorway interval ( $D_A$ ), always finishes in finite number of steps
  - Waiting interval ( $W_A$ ), may take unbounded number of steps

# $r$ -Bounded Waiting

typically you want  $r$  to be 1

- For threads A and B
  - If  $D_A^k \rightarrow D_B^j$ 
    - A's  $k^{\text{th}}$  doorway precedes B's  $j^{\text{th}}$  doorway
  - Then  $CS_A^k \rightarrow CS_B^{j+r}$ 
    - A's  $k^{\text{th}}$  critical section precedes B's  $(j+r)^{\text{th}}$  critical section
    - B cannot overtake A by more than  $r$  times
  
- First-come-first-served if  $r = 0$

# Fairness Again

- Filter Lock satisfies some properties
  - No one starves (no lockout)
  - But very weak fairness
    - Not  $r$ -bounded for any  $r$ !
  - That's pretty lame...

# Bakery Algorithm


- Provides First-Come-First-Served
- How?
  - Take a “number”
  - Wait until lower numbers have been served
- Lexicographic order
  - $(a, i) > (b, j)$ 
    - If  $a > b$ , or  $a == b$  and  $i > j$



# Bakery Algorithm

```

class Bakery implements Lock {
    volatile boolean[] flag;
    volatile int[] label;
    public Bakery (int n) {
        flag = new boolean[n];
        label = new int[n];
        for (int i = 0; i < n; i++) {
            flag[i] = false; label[i] = 0;
        }
    }
    ...
  
```



# Bakery Algorithm

Doorway lines are completed in a finite number of steps

Doorway

```

class Bakery implements Lock {
  ...
  public void lock() {
    flag[i] = true;
    label[i] = max(label[0], ..., label[n-1]) + 1;
    while (( $\exists k \mid$  flag[k])
      && (label[i], i) > (label[k], k)) {}
  }
  ...
}

```

Take increasing label  
 (read labels in some  
 arbitrary order)

I'm interested

Someone is interested...

...with lower label in lexicographic order

# Bakery Algorithm

```

class Bakery implements Lock {
    ...
    public void unlock() {
        flag[i] = false;
    }
}

```

➤ No longer interested

# No Deadlock

- There is always one thread with earliest label
- Ties are impossible (why?) because we have a total order

# First-Come-First-Served

- If  $D_A \rightarrow D_B$  then A's label is earlier
  - $\text{write}_A(\text{label}[A]) \rightarrow \text{read}_B(\text{label}[A]) \rightarrow \text{write}_B(\text{label}[B]) \rightarrow \text{read}_B(\text{flag}[A])$
- So B is locked out while  **$\text{flag}[A] == \text{true}$**

```

class Bakery implements Lock {
    public void lock() {
        flag[i] = true;
        label[i] = max(label[0], ..., label[n-1]) + 1;
        while ((∃k | flag[k]) && (label[i], i) > (label[k], k)) {}
    }
}

```

# Mutual Exclusion

- Suppose A and B in CS together
- Suppose A has earlier label
- When B entered, it must have seen  
 $\text{flag}[A] == \text{false}$  or  $\text{label}[A] > \text{label}[B]$

```

class Bakery implements Lock {
    public void lock() {
        flag[i] = true;
        label[i] = max(label[0], ..., label[n-1]) + 1;
        while (( $\exists k \mid \text{flag}[k]$ ) && (label[i], i) > (label[k], k)) {}
    }
}
  
```

# Mutual Exclusion

- Labels are strictly increasing so B must have seen **flag[A] == false**
  - $\text{labeling}_B \rightarrow \text{read}_B(\text{flag}[A]) \rightarrow \text{write}_A(\text{flag}[A]) \rightarrow \text{labeling}_A$
- Which contradicts the assumption that A has an earlier label

# Deep Philosophical Question

- The Bakery Algorithm is
  - Succinct
  - Elegant
  - Fair
- Q: So why isn't it practical?
- A: Well, you have to read  $n$  distinct variables

you have to go through all arrays. cost will increase linear to the number of threads



# Shared Memory

- Shared read/write memory locations called **registers** (historical reasons)
- Come in different flavors
  - Multi-Reader-Single-Writer (**flag[]**)
  - Multi-Reader-Multi-Writer (**victim[]**)
  - Not interesting: SRMW and SRSW

# Theorem

- At least  $n$  MRSW (multi-reader/single-writer, like `flag[]`) registers are needed to solve deadlock-free mutual exclusion

# Summary of Lecture

- In the 1960's many **incorrect** solutions to lockout-free mutual exclusion using RW-registers were published...
- Today we know how to solve FIFO  $n$  thread mutual exclusion using  $2n$  RW-registers