

Concurrency: Foundations and Algorithms

Concurrent Objects and Consistency



Prof. P. Felber

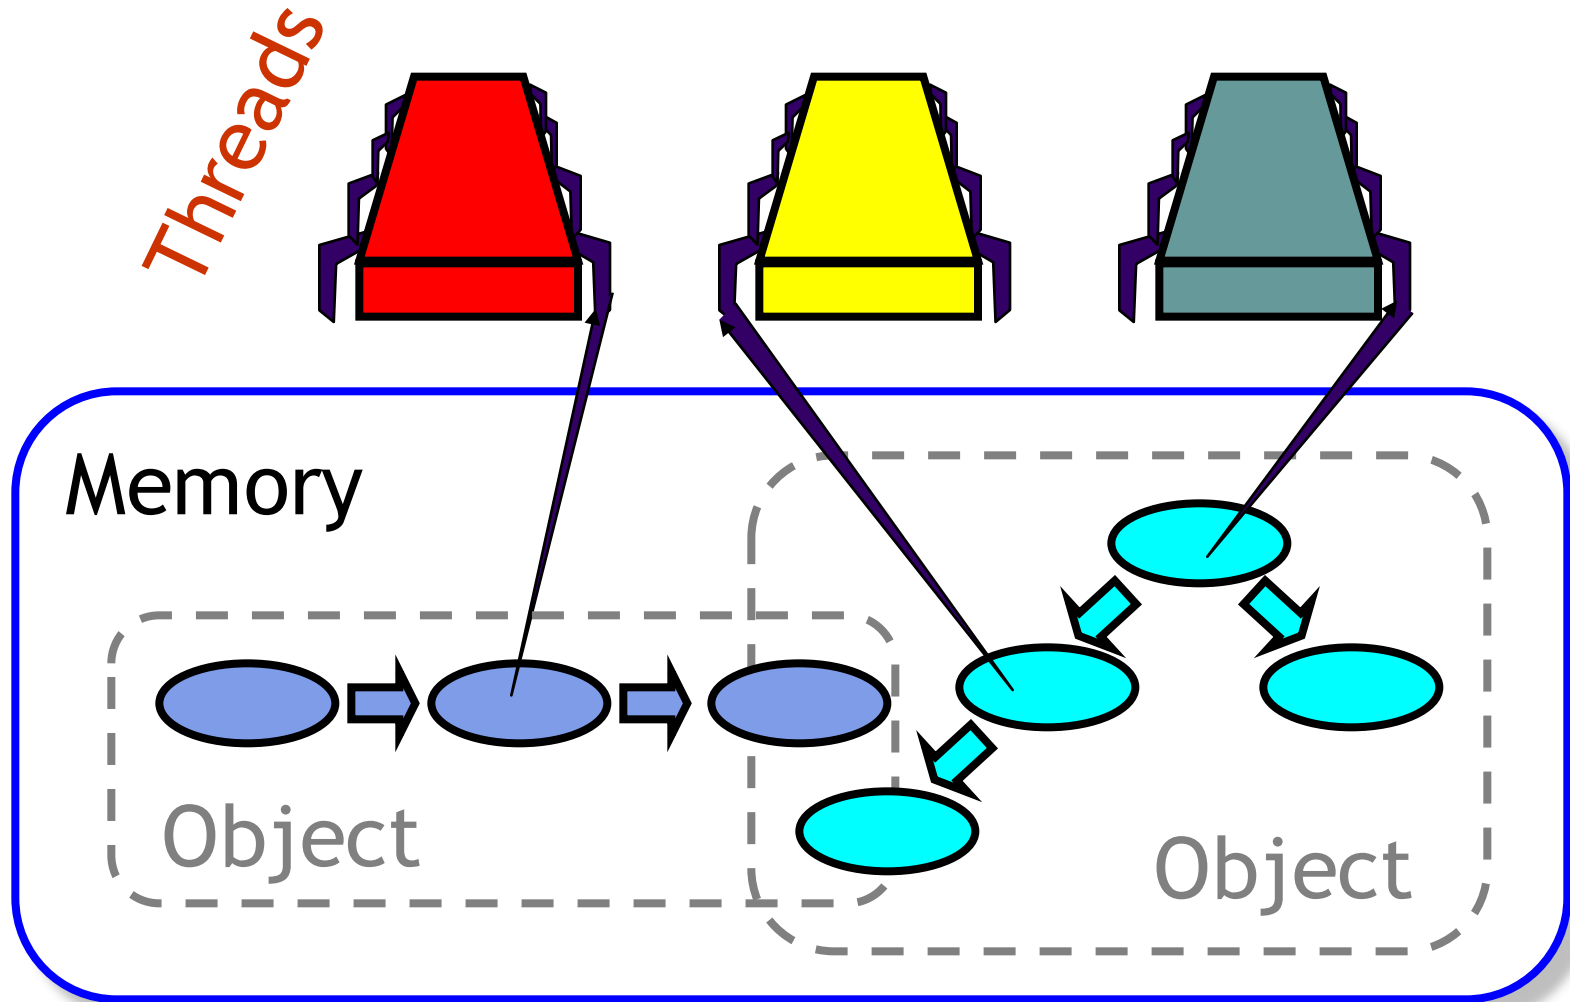
Pascal.Felber@unine.ch

<http://iiun.unine.ch/>

Based on slides by Maurice Herlihy and Nir Shavit




Concurrent Computation




Objectivism

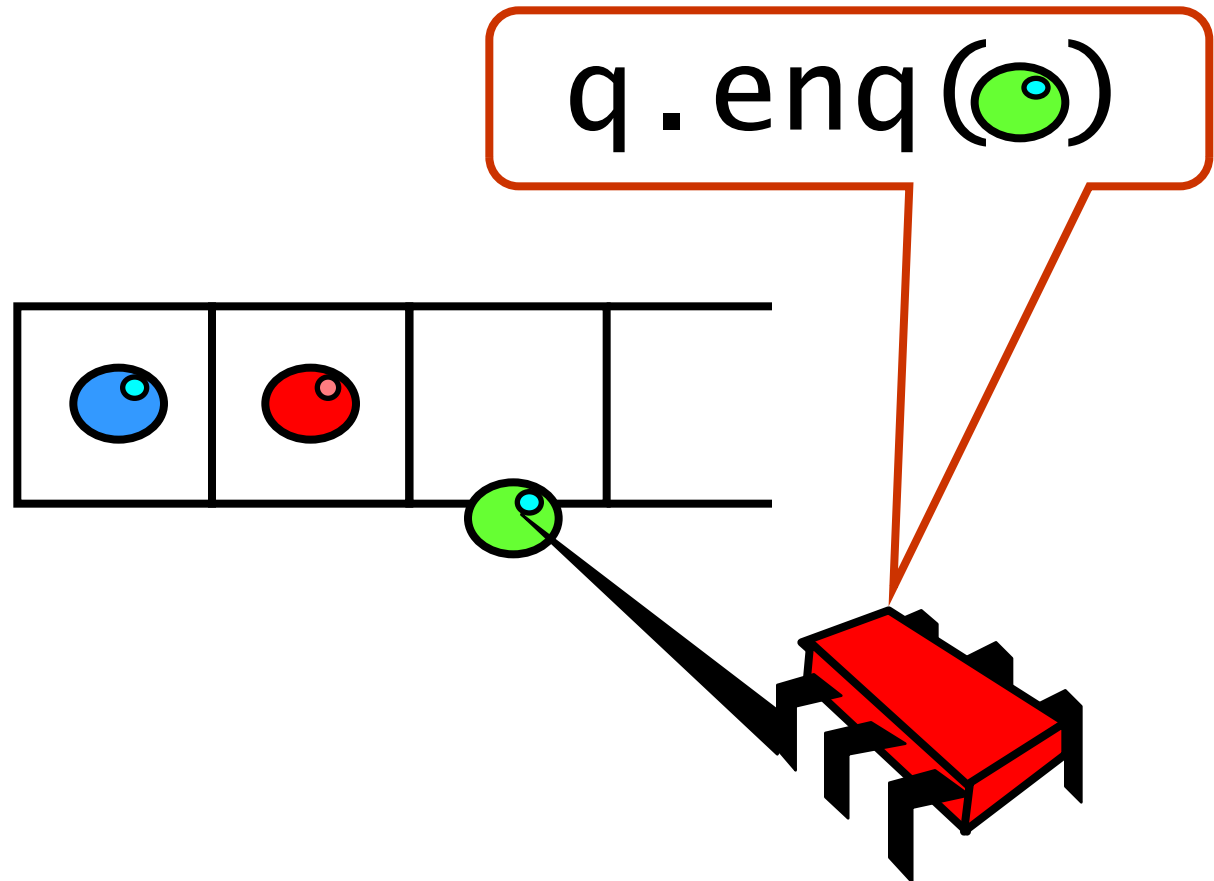
- What is a concurrent object?
 - How do we **describe** one?
 - How do we **implement** one?
 - How do we **tell if we are right**?

This lecture

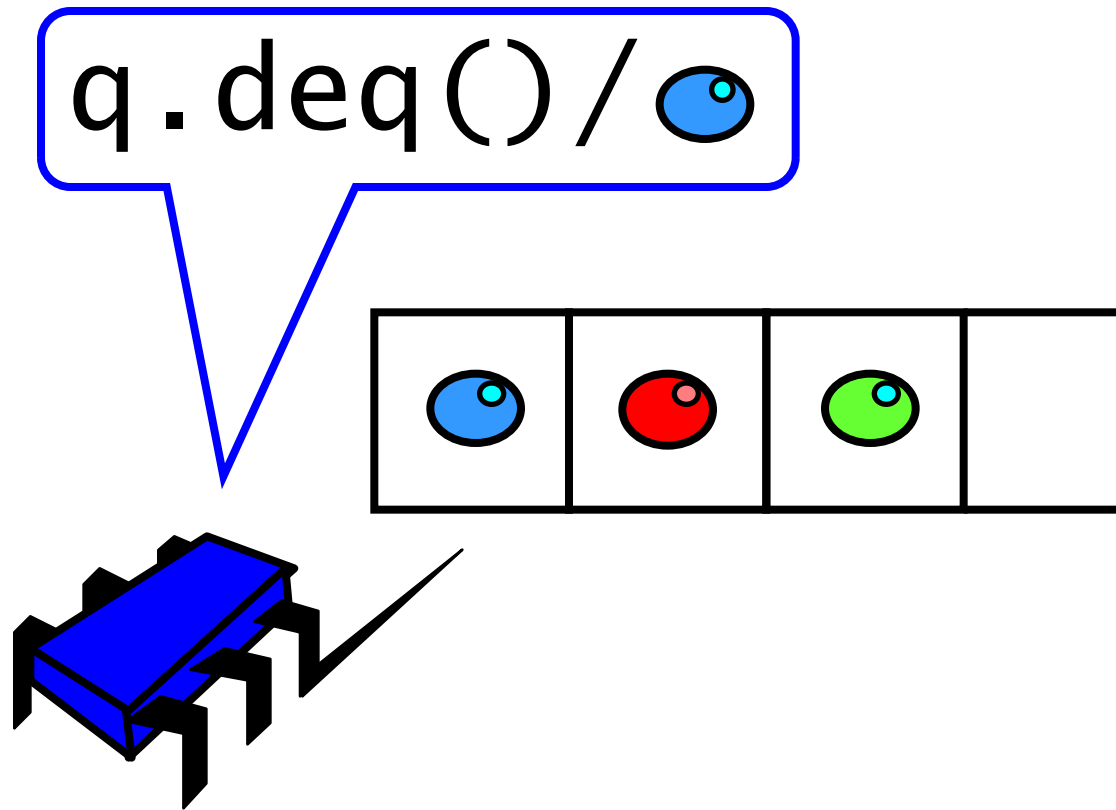




FIFO Queue: Enqueue Method



FIFO Queue: Dequeue Method



Implementation: enq()

```

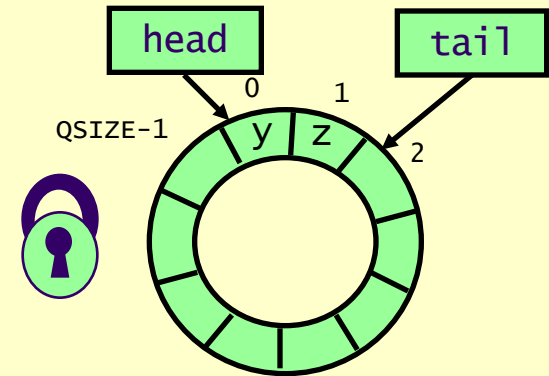
public class Queue {
  int head = 0, tail = 0;
  Object[QSIZE] items;

```

```

  public synchronized
  void enq(Object x) {
    while (tail - head == QSIZE)
      this.wait();
    items[tail % QSIZE] = x;
    tail++;
    this.notifyAll();
  } ...
}

```



Method execution is mutually exclusive. Release lock if need to wait.

Here is where the queue is actually updated!

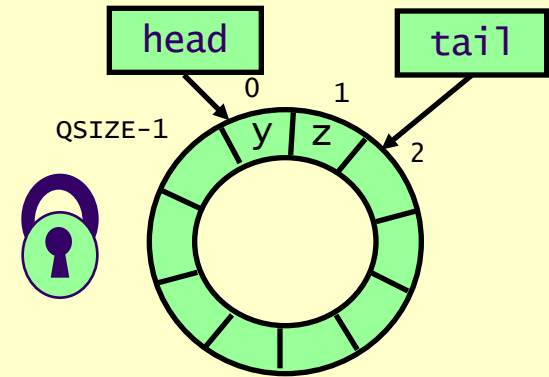
Notify all others that you release the lock

Implementation: `deq()`

```

public class Queue {
  int head = 0, tail = 0;
  Object[QSIZE] items;

  public synchronized
  Object deq() {
    while (tail - head == 0)
      this.wait();
    Object x = items[head % QSIZE];
    head++;
    this.notifyAll();
    return x;
  } ...
}
  
```



We understand it is “correct” because all modifications are mutually exclusive...

A Concurrent Implementation

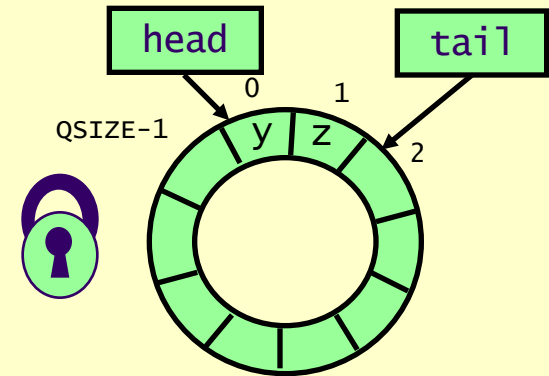
- Now consider the following implementation
 - The same thing without mutual exclusion

- For simplicity, **only two threads**:
 - One thread only calls **enq()**
 - The other only **deq()**

Lock-free 2-Thread Queue

```

public class LockFreeQueue {
    volatile int head = 0,
               tail = 0;
    Object[QSIZE] items;
    public void enq(Item x) {
        while (tail-head == QSIZE); // busy-wait
        items[tail % QSIZE] = x; tail++;
    }
    public Item deq() {
        while (tail == head); // busy-wait
        Item item = items[head % QSIZE]; head++;
        return item;
    }
}
  
```



Queue updated without lock

How do we define “correct” when modifications are not exclusive?

Defining “Correct”

- Need a way to **specify** a concurrent queue object
- Need a way to prove that an algorithm **implements** the object's specification
- Lets talk about object specifications...

Sequential Objects

- Each object has a **state**
 - Usually given by a set of *fields*
 - Queue example: sequence of items

- Each object has a set of **methods**
 - Only way to manipulate state
 - Queue example: **enq()** and **deq()** methods

Sequential Specifications

- If (precondition)...
 - The object is in such-and-such a state...
 - Before you call the method...
- Then (postcondition)...
 - The method will return a particular value...
 - Or throw a particular exception...
- And (postcondition, con't)...
 - The object will be in some other state...
 - When the method returns

Pre- and Postconditions: `deq()`

- Precondition
 - Queue is non-empty
- Postcondition
 - Returns first item in queue
- Postcondition
 - Removes first item in queue

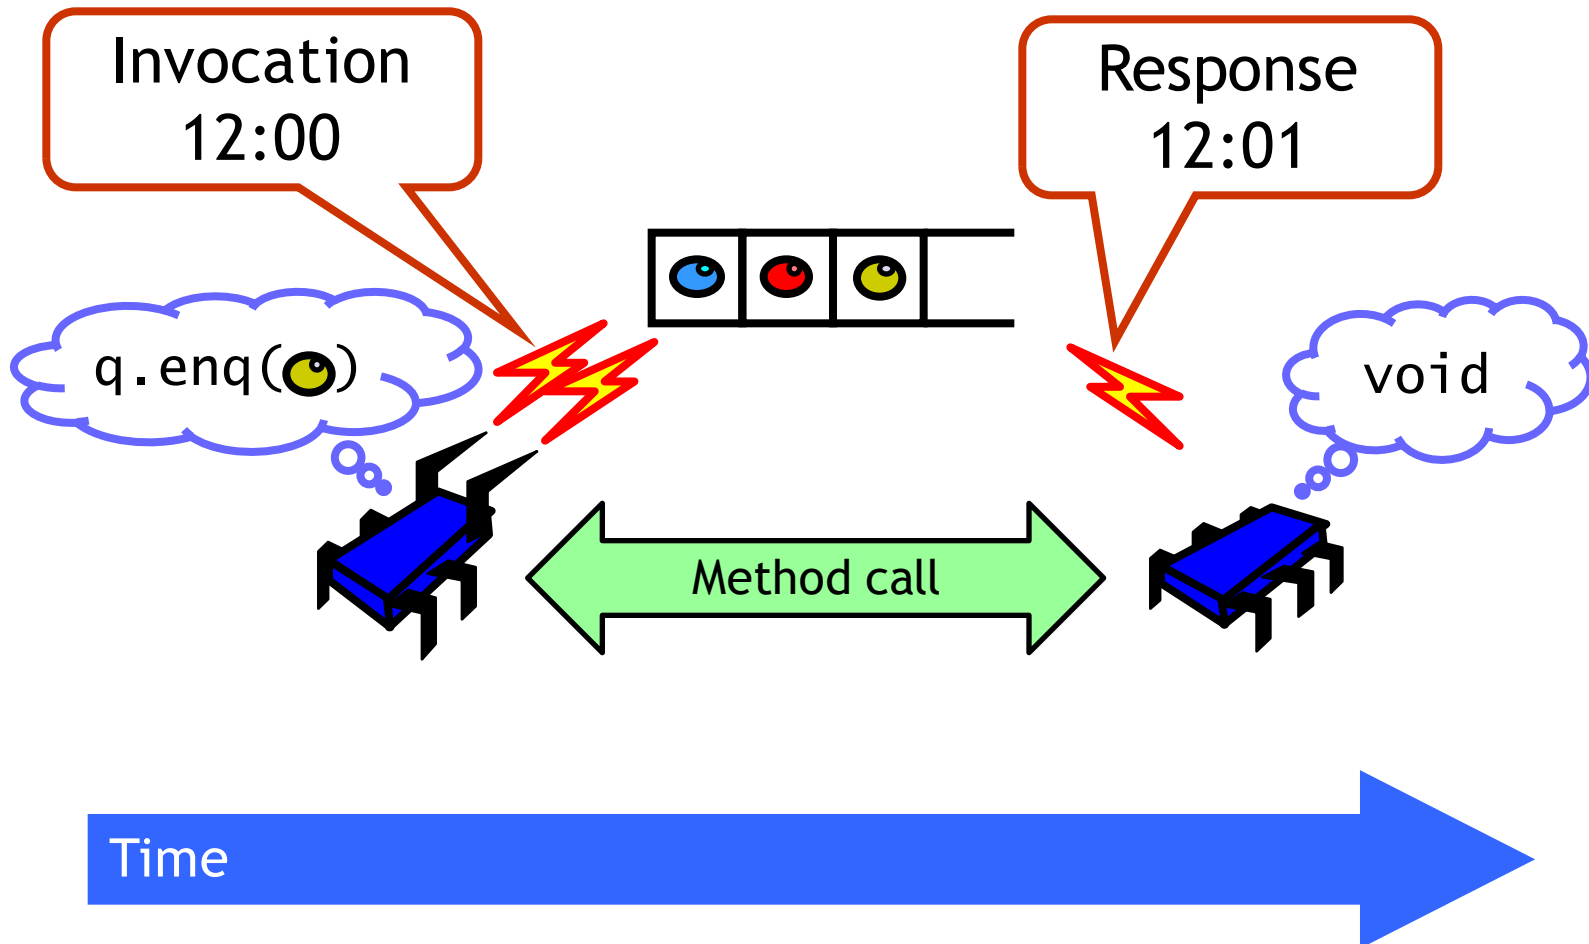
Pre- and Postconditions: `deq()`

- Precondition
 - Queue is empty
- Postcondition
 - Throws Empty exception
- Postcondition
 - Queue state unchanged

Why Sequential Specifications Totally Rock

- Interactions among methods captured by side-effects on object state
 - State meaningful between method calls
- Documentation size linear in # of methods
 - Each method described in isolation
- Can add new methods
 - Without changing descriptions of old methods
- What About Concurrent Specifications?
 - Methods? Documentation? Adding new methods?

Methods Take Time

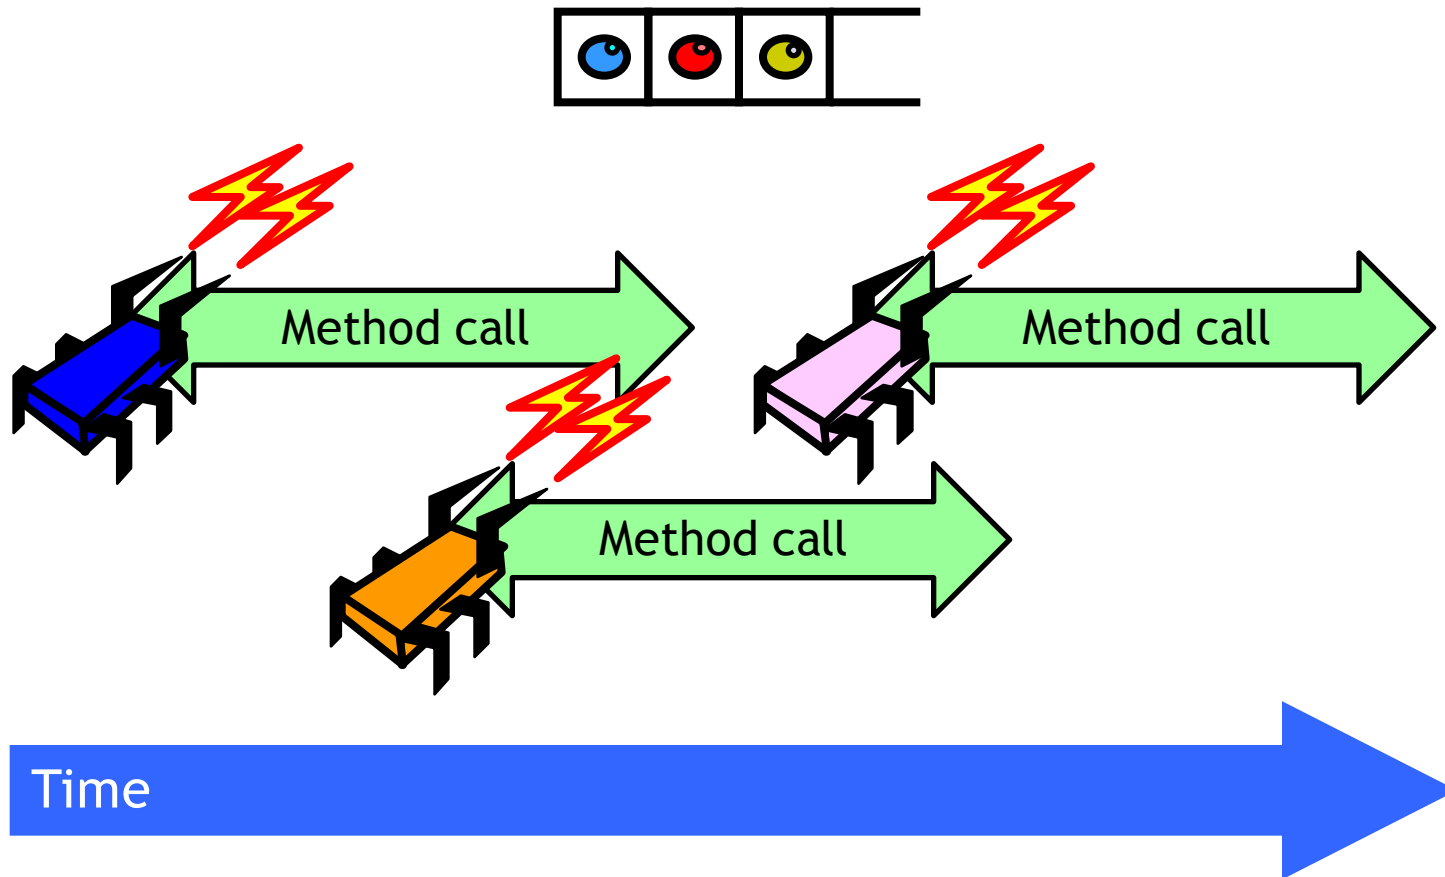


Sequential vs. Concurrent

- Sequential
 - Methods take time? Who knew?

- Concurrent
 - Method call is **not** an event
 - Method call is an interval

Concurrent Methods Take Overlapping Time



Sequential vs. Concurrent

Sequential

- Object needs meaningful state only between method calls
- Each method described in isolation
- Can add new methods without affecting older methods

Concurrent

- Because method calls overlap, object might **never** be between method calls
- Must characterize **all** possible interactions with concurrent calls
 - What if 2 `enq()` overlap? Or 2 `deq()`? Or `enq()` and `deq()`?
- Everything can potentially interact with **everything else**

PANIC!

The Big Question

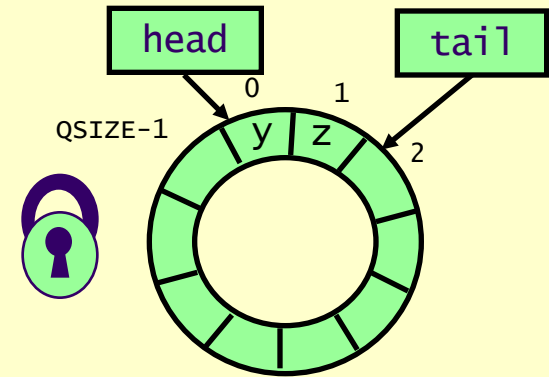
- What does it **mean** for a *concurrent* object to be correct?
 - What *is* a concurrent FIFO queue?
 - FIFO means strict temporal order
 - Concurrent means **ambiguous temporal order**

Intuitively...

```

public class Queue {
    int head = 0, tail = 0;
    Object[QSIZE] items;

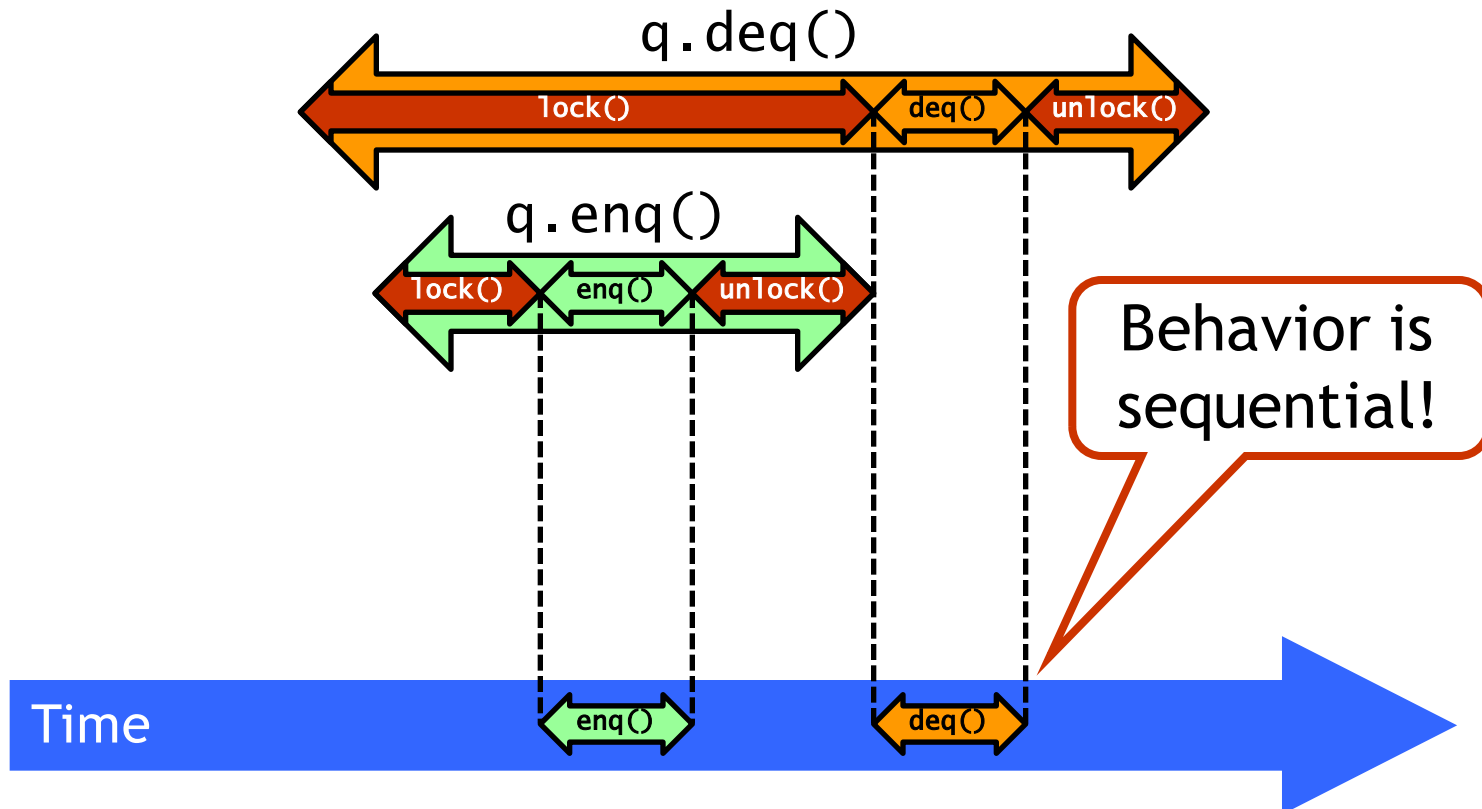
    public synchronized
    void enq(Object x) {
        while (tail - head == QSIZE)
            this.wait();
        items[tail % QSIZE] = x;
        tail++;
        this.notifyAll();
    } ...
}
  
```



Queue is updated while holding lock (mutually exclusive)

Intuitively...

Let's capture the idea of describing the concurrent via the sequential



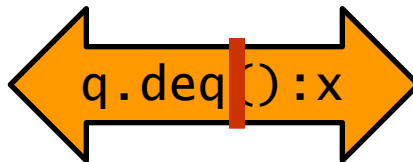
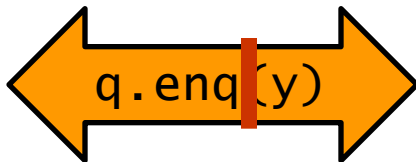
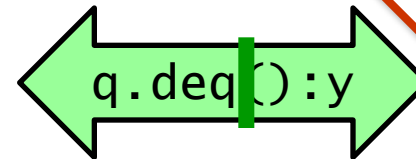
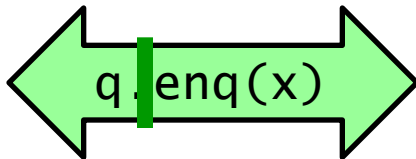
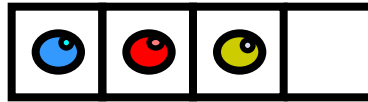
Linearizability

See previous graphic

- Each method should
 - “Take effect”
 - Instantaneously
 - Between invocation and response events
- Object is correct if this “sequential” behavior is correct
- Any such concurrent object is **linearizable**
 - Formally, a **linearizable object** is an object all of whose possible executions are linearizable

Example

If we can find some points in the execution such as the time line is correct, it's linearizable.

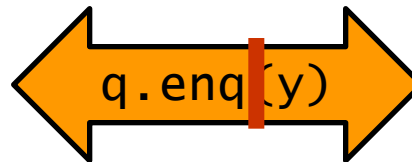
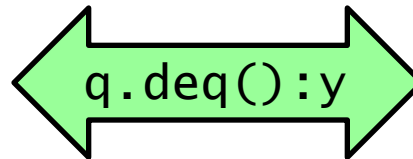
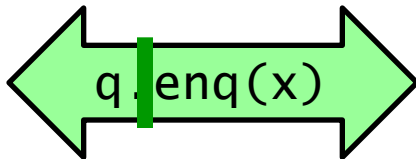
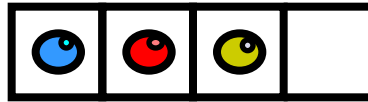


Linearizable

Time



Example

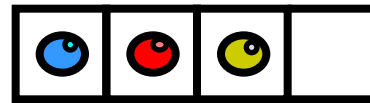


Not linearizable

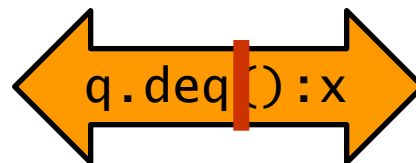
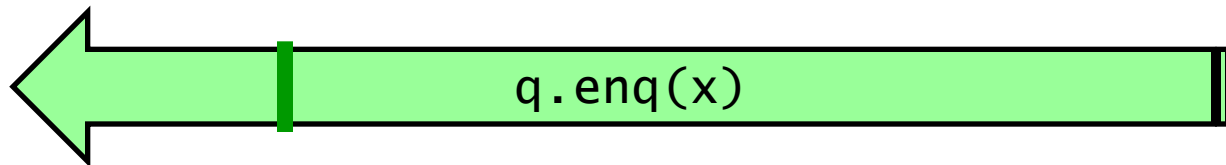
Time



Example

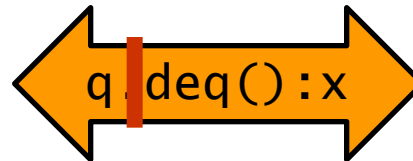
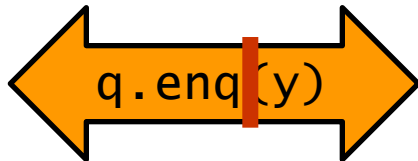
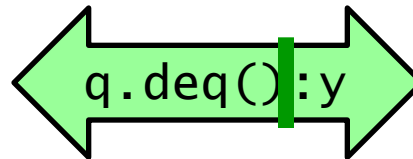
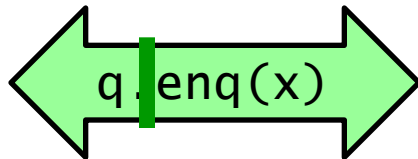
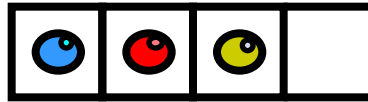


Linearizable



Example

Exam question could be drawing this stuff



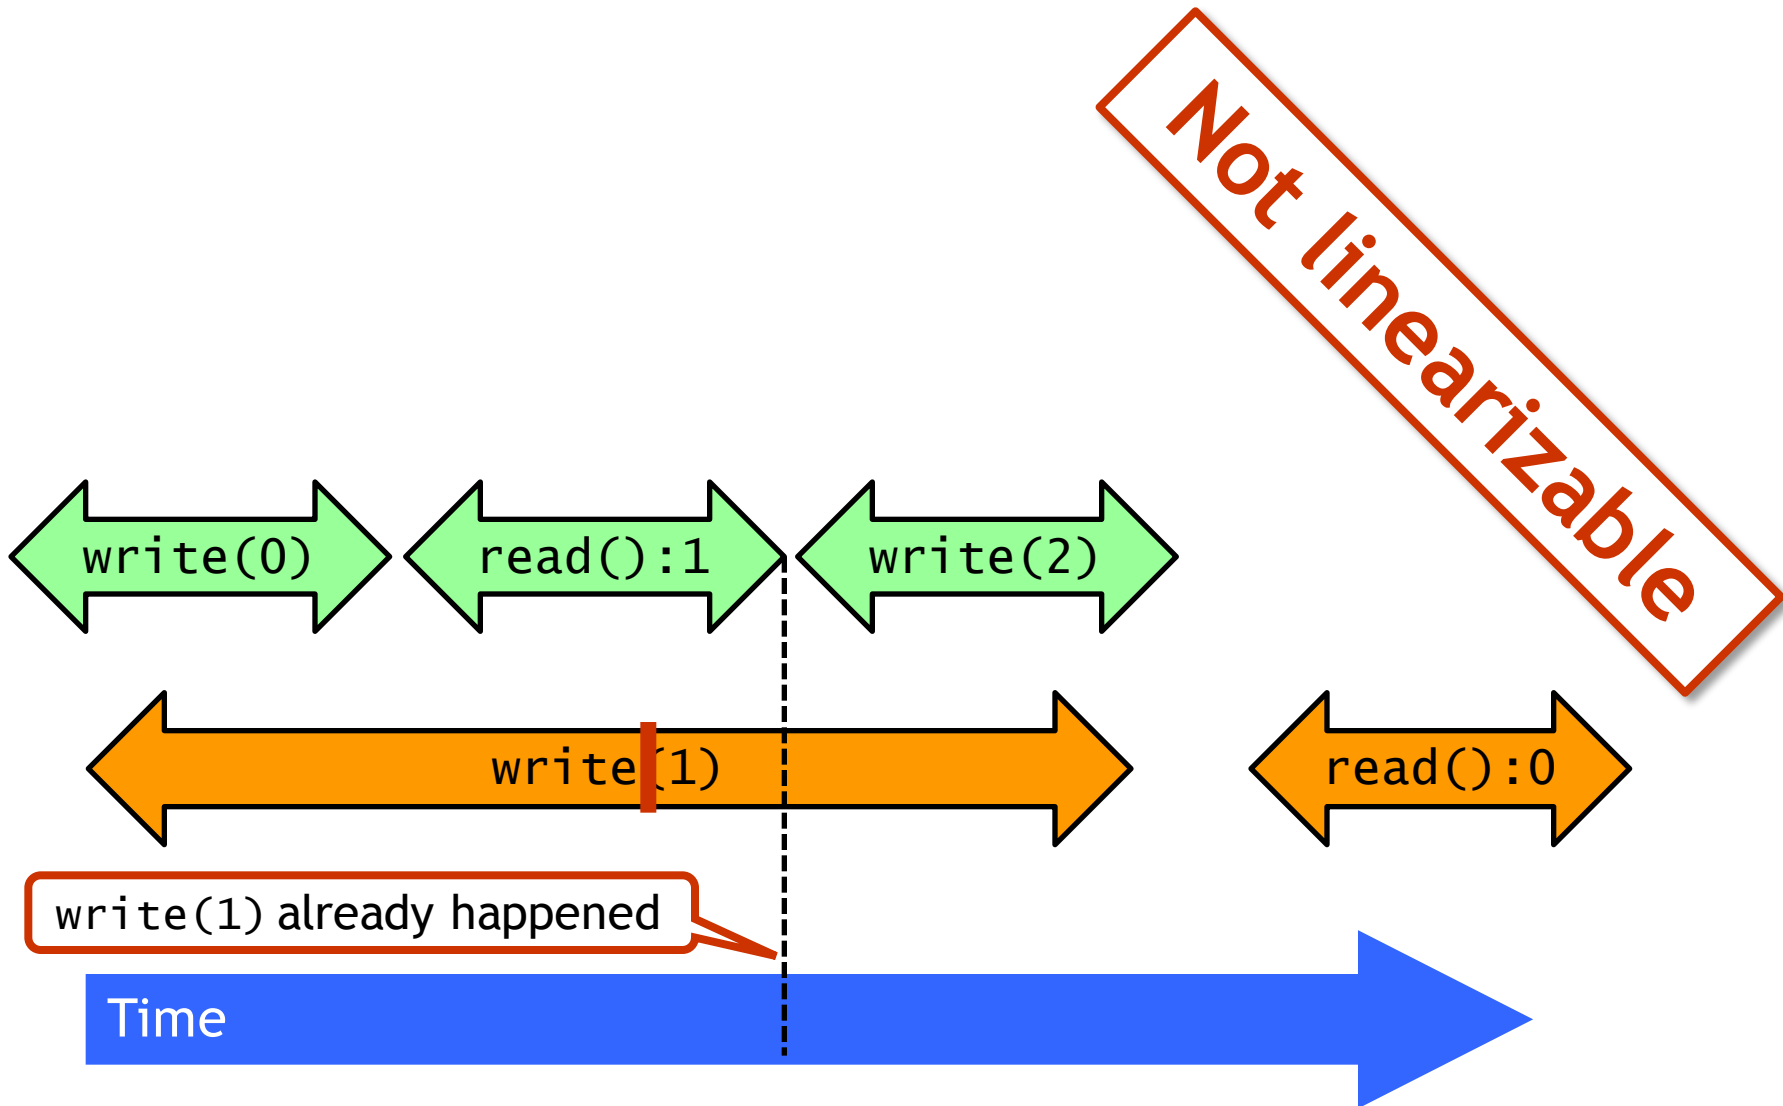
Linearizable

 Multiple orders OK

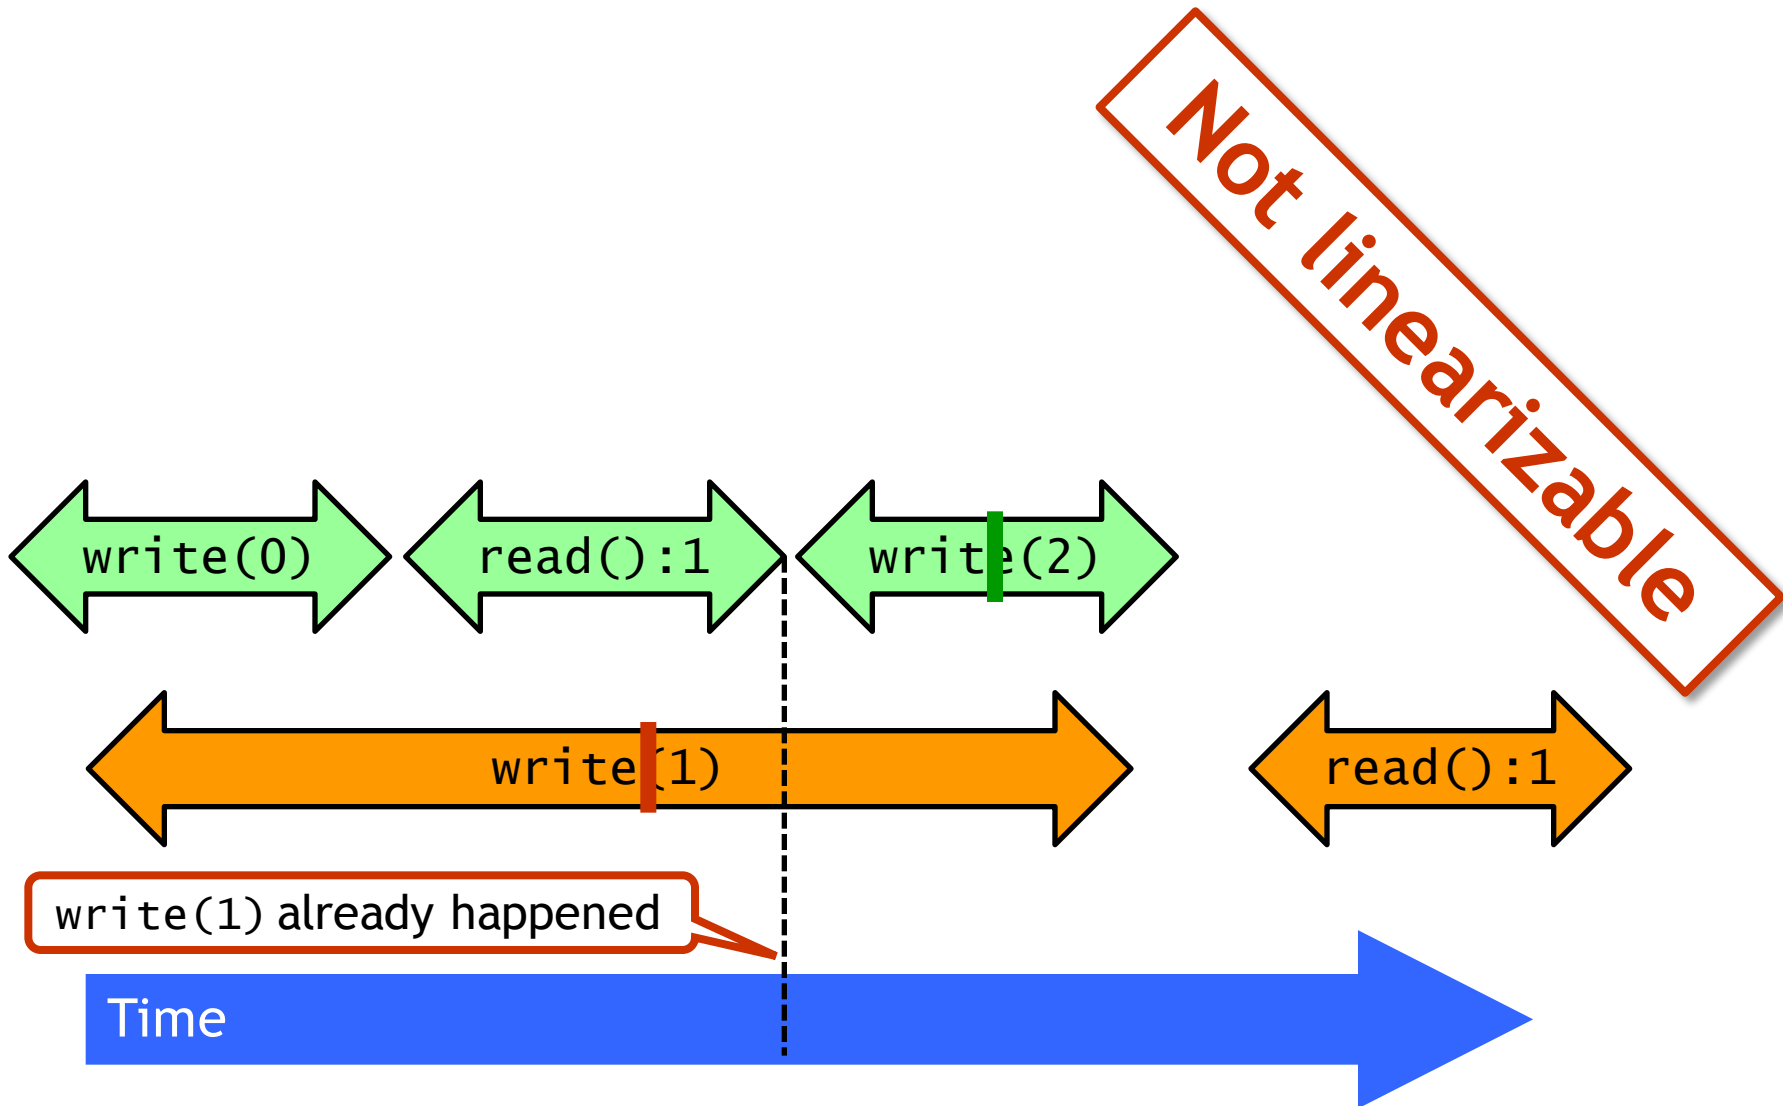
Time



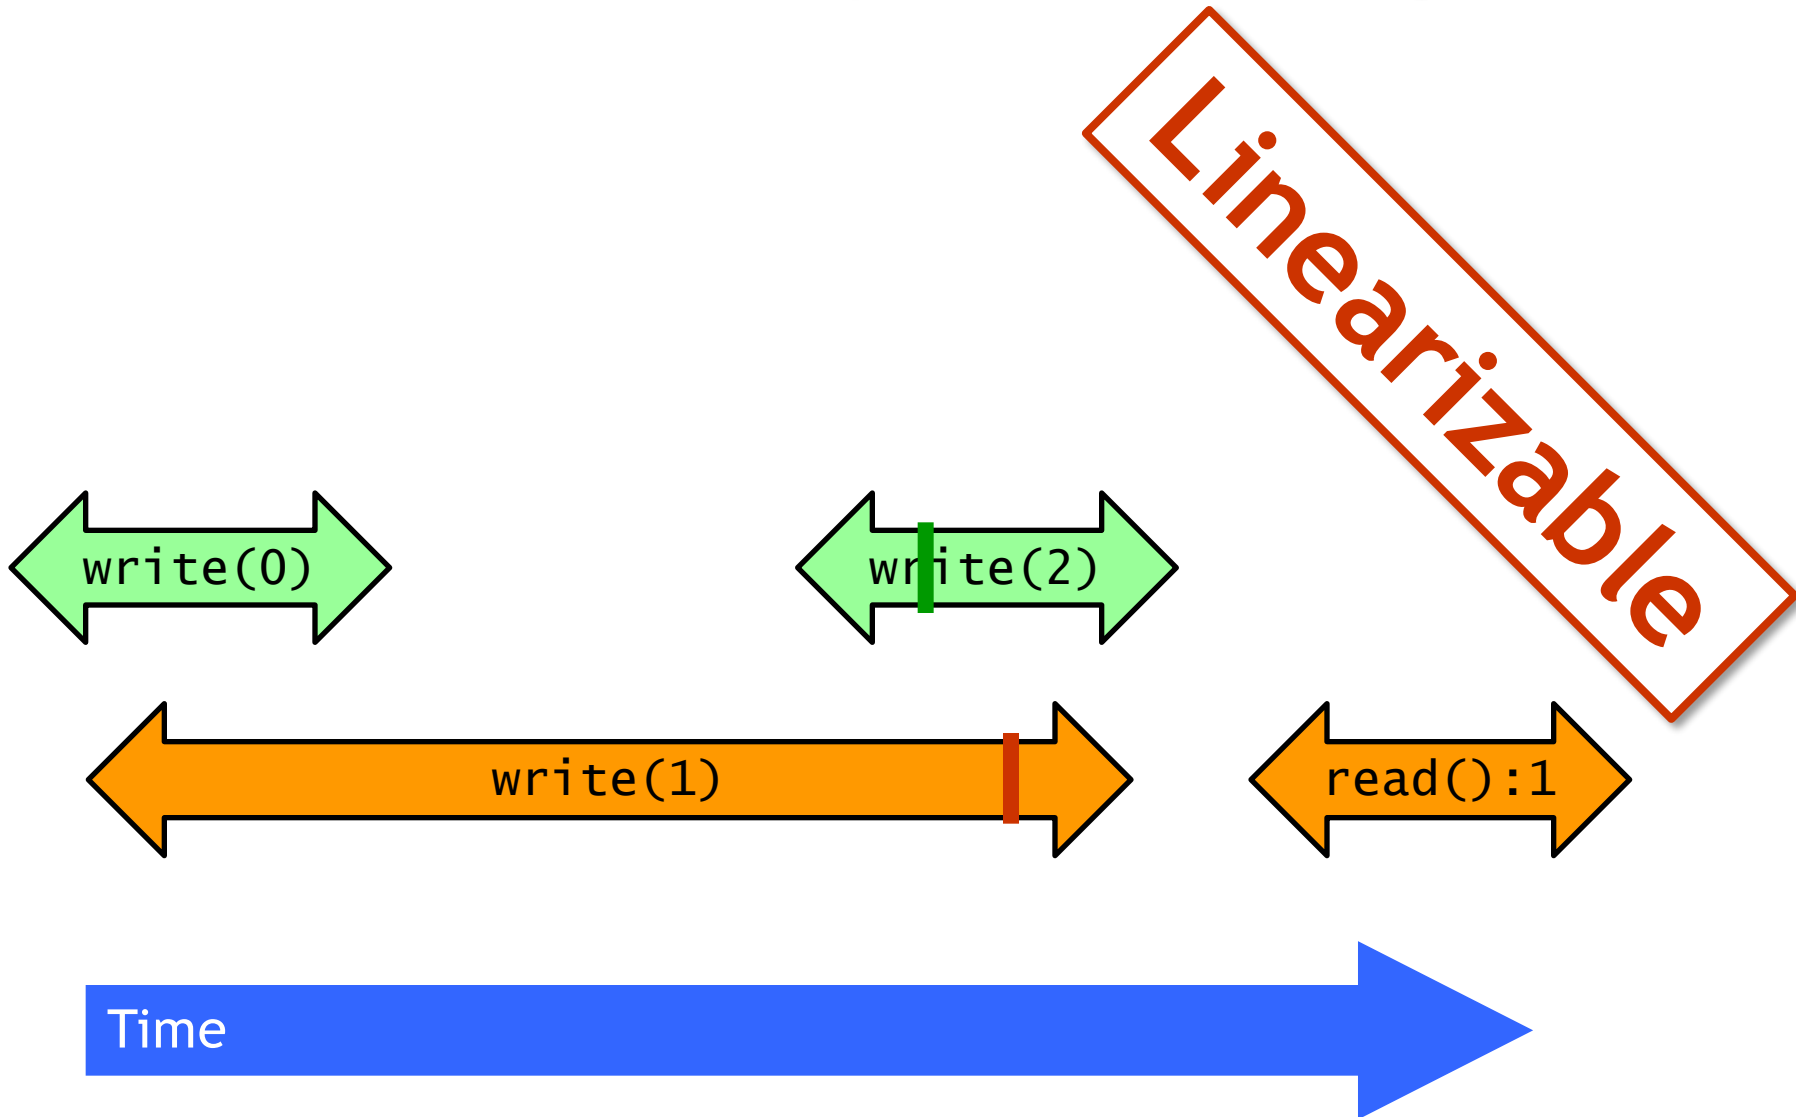
Read/Write Register Example



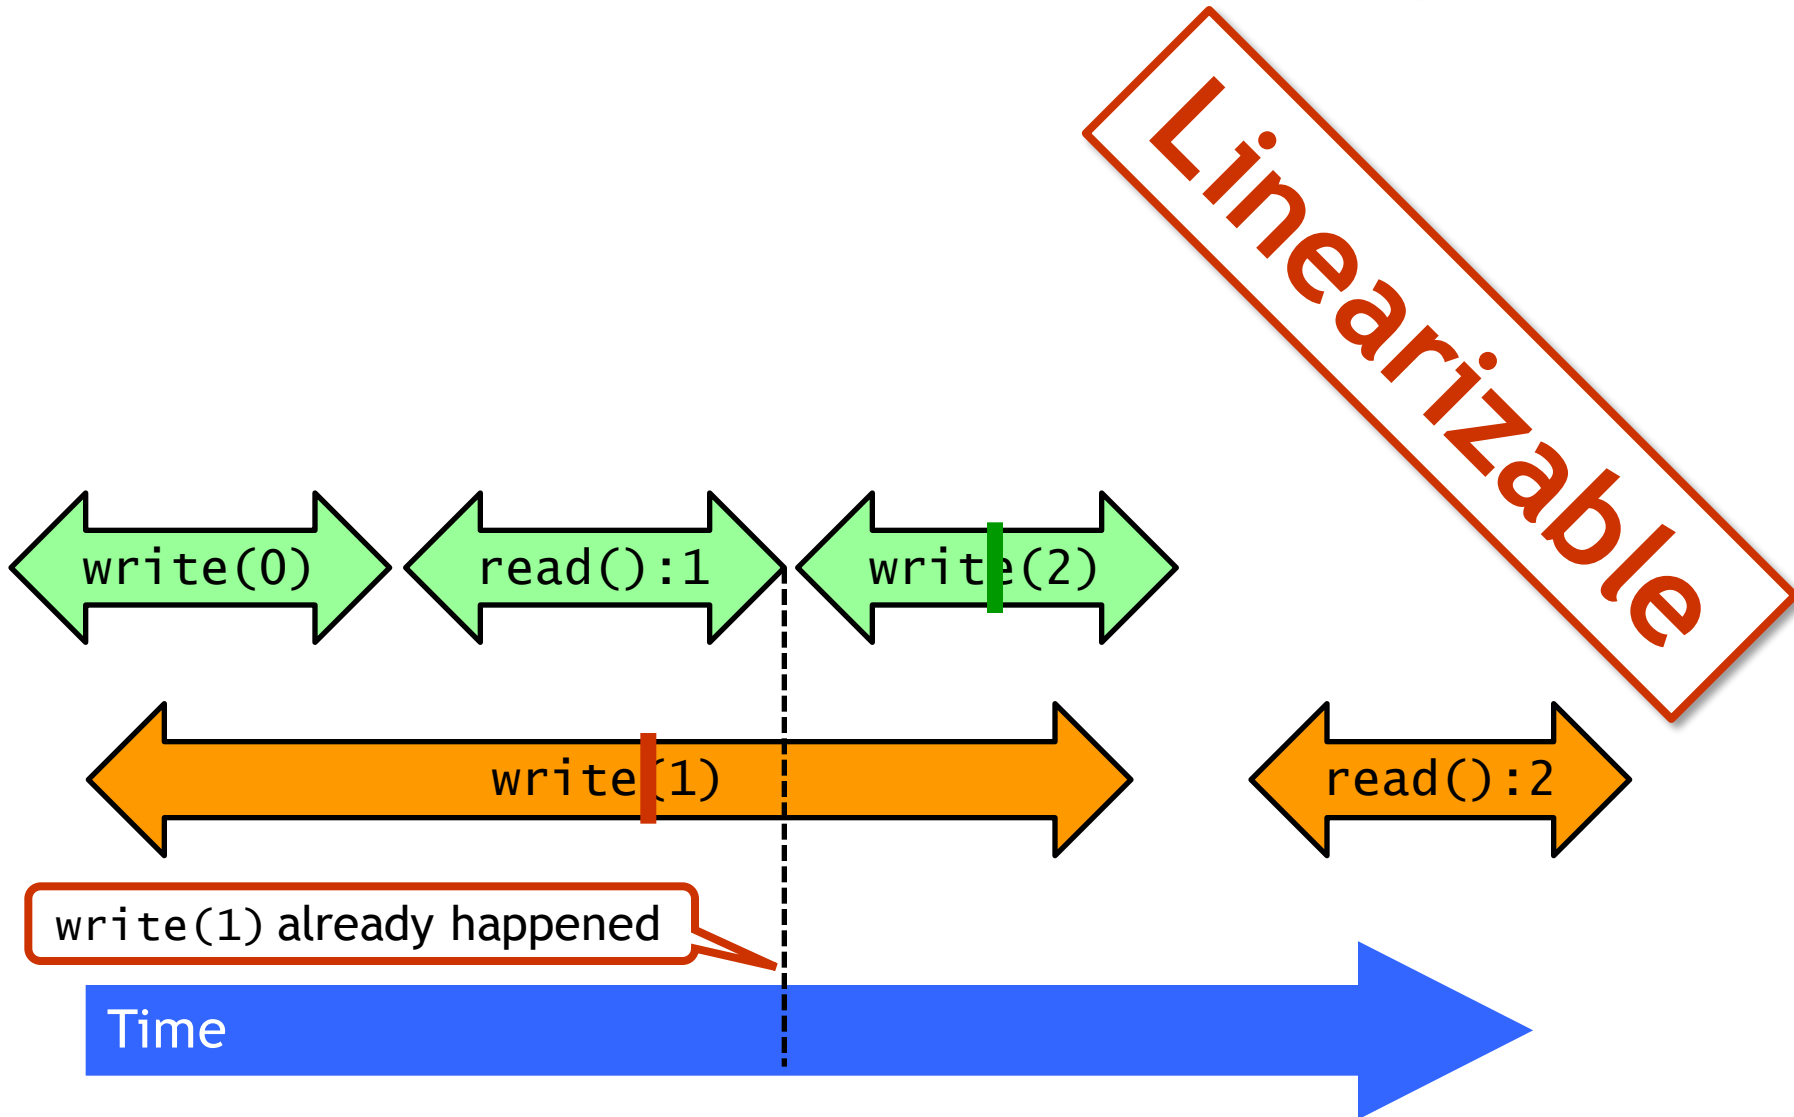
Read/Write Register Example



Read/Write Register Example



Read/Write Register Example



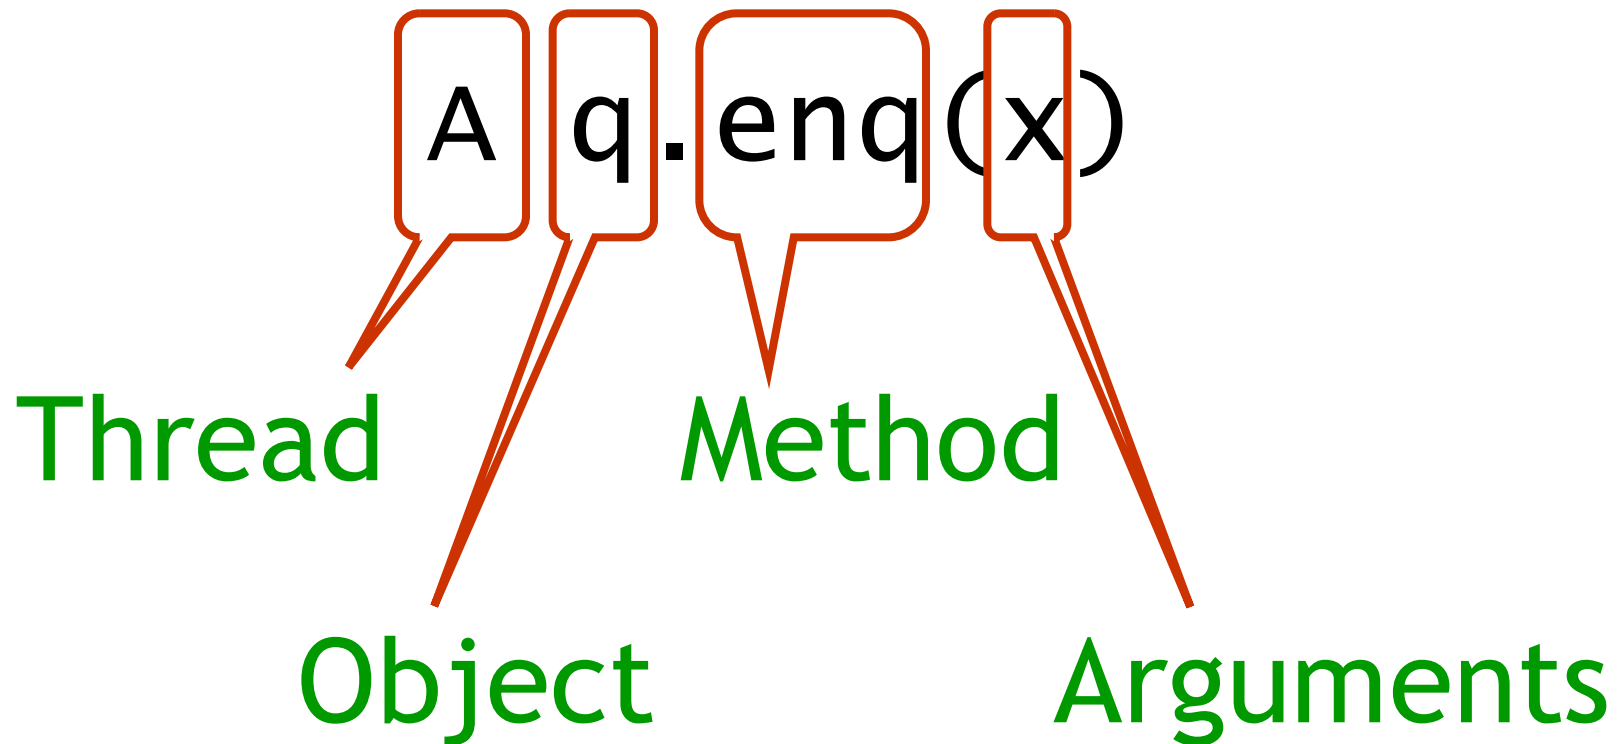
Talking About Executions

- Why do we need to consider executions?
 - Can't we specify the linearization point of each operation without describing an execution?
- Not Always
 - In some cases, linearization point depends on the execution
- Let's define a formal model of executions
 - Define precisely what we mean (ambiguity is bad)
 - Allow reasoning, formal or informal

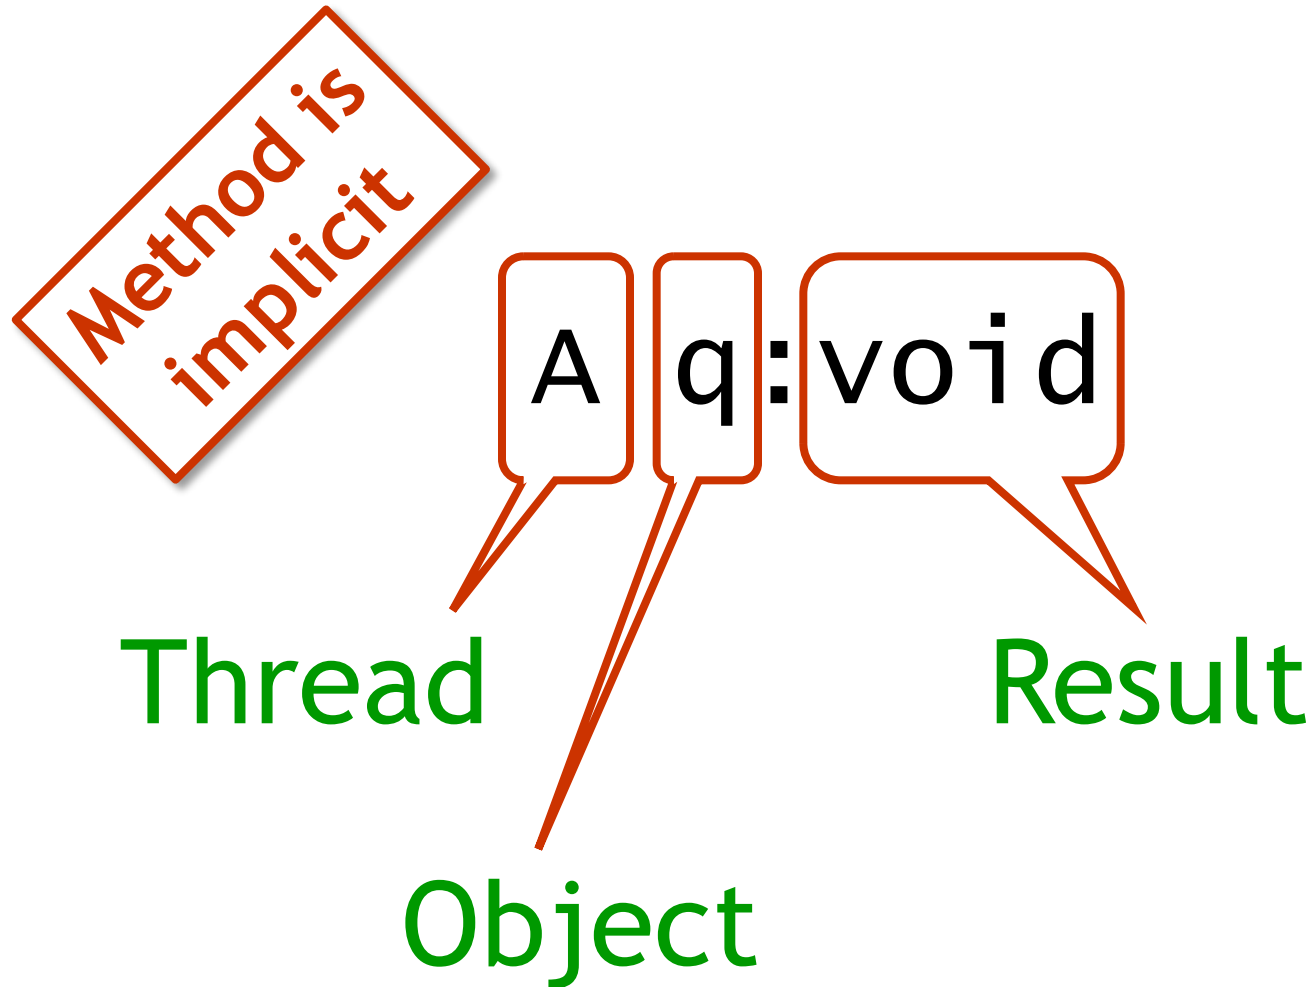
Split Method Calls into Two Events

- Invocation
 - Method name & arguments
`q.enq(x)`
- Response
 - Result or exception
`q.enq(x)` returns `void`
`q.deq()` returns `x`
`q.deq()` throws `Empty`

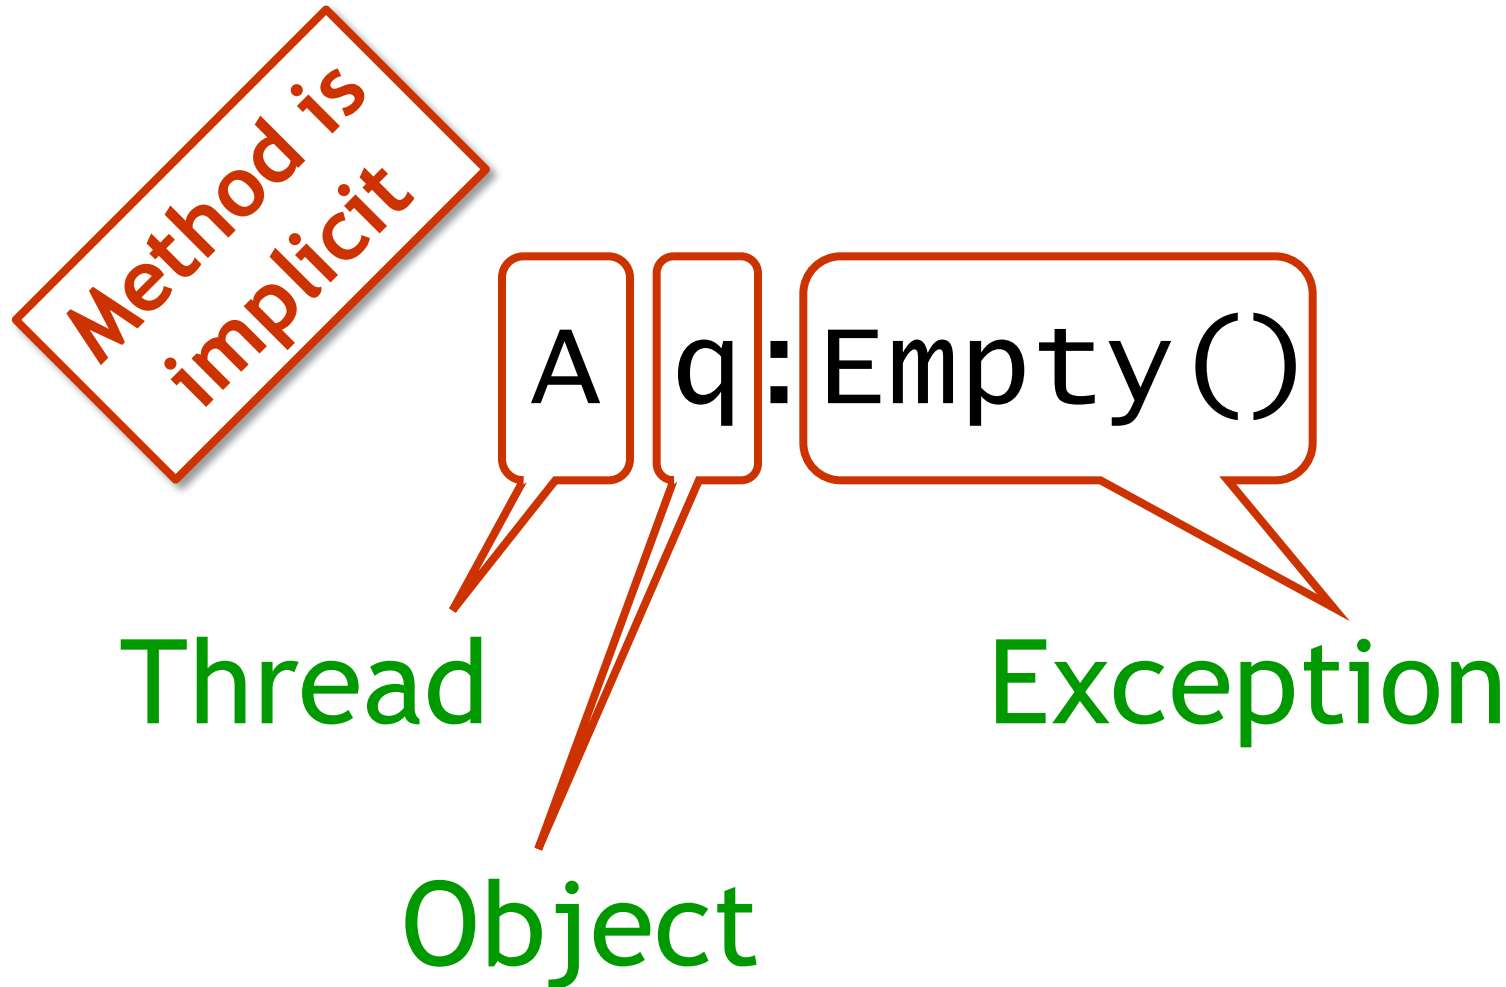
Invocation Notation



Response Notation



Response Notation



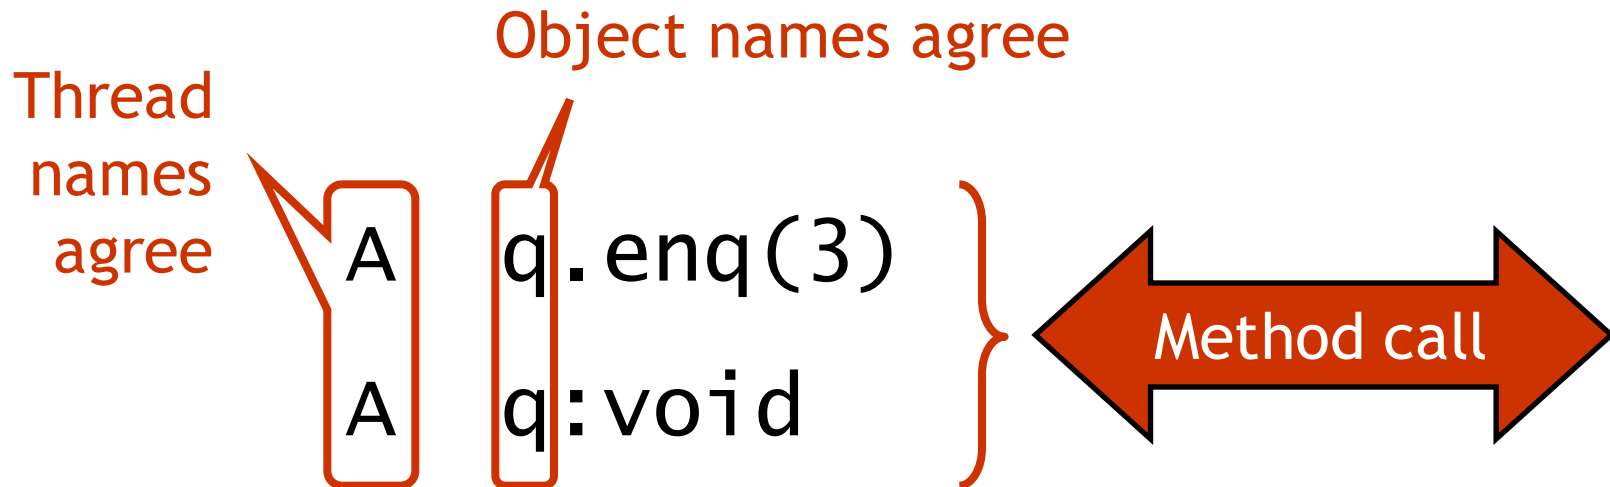
History: Describing an Execution

$$H = \left\{ \begin{array}{l} A \text{ } q.\text{enq}(3) \\ A \text{ } q:\text{void} \\ A \text{ } q.\text{enq}(5) \\ B \text{ } p.\text{enq}(4) \\ B \text{ } p:\text{void} \\ B \text{ } q.\text{deq}() \\ B \text{ } q:3 \end{array} \right.$$

Sequence of
invocations and
responses

Definition

- Invocation & response **match** if



Object Projections

$H =$	A q.enq(3) A q:void B p.enq(4) B p:void B q.deq() B q:3	$H q =$	A q.enq(3) A q:void B q.deq() B q:3
-------	--	---------	--

Thread Projections

$H =$
 A $q.enq(3)$
 A $q:void$
 B $p.enq(4)$
 B $p:void$
 B $q.deq()$
 B $q:3$

$H|B =$
 B $p.enq(4)$
 B $p:void$
 B $q.deq()$
 B $q:3$

Complete Subhistory

A q.enq(3)
 A q:void
 A q.enq(5)
 H = B p.enq(4)
 B p:void
 B q.deq()
 B q:3

An invocation is pending if it has no matching response

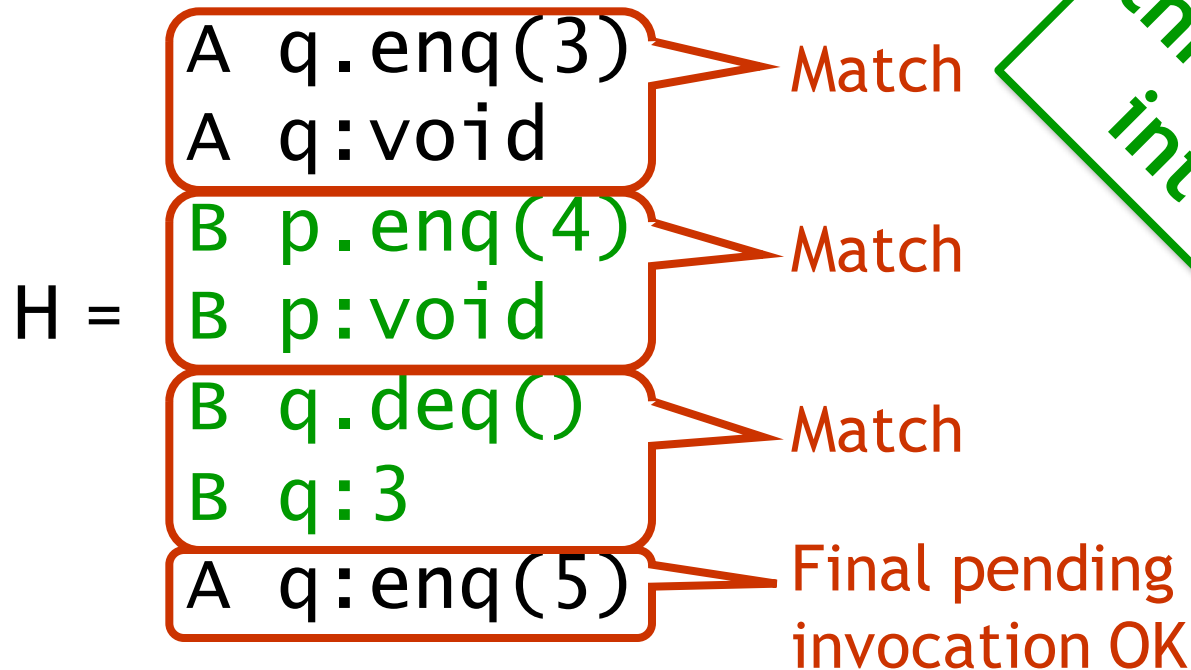
May or may not have taken effect
 ⇒ discard pending invocations

Complete Subhistory

A q.enq(3)
 A q:void

Complete(H) = B p.enq(4)
 B p:void
 B q.deq()
 B q:3

Sequential Histories



Method calls
 of different
 threads do not
 interleave

Well-Formed Histories

Per-thread projections
are sequential

$H =$
 A q.enq(3)
 B p.enq(4)
 B p:void
 B q.deq()
 A q:void
 B q:3

$H|B =$
 B p.enq(4)
 B p:void
 B q.deq()
 B q:3

$H|A =$
 A q.enq(3)
 A q:void

Equivalent Histories

Threads see the same
thing in both histories

$$\begin{cases} H|A = G|A \\ H|B = G|B \end{cases}$$

H =

```

A q.enq(3)
B p.enq(4)
B p:void
B q.deq()
A q:void
B q:3
  
```

G =

```

A q.enq(3)
A q:void
B p.enq(4)
B p:void
B q.deq()
B q:3
  
```

Sequential Specifications

- A sequential specification is some way of telling whether a
 - ...single-thread, single-object history
 - ...is legal
- Simple way is using
 - ...pre and post-conditions
 - But plenty of other techniques exist

Legal Histories

- A sequential (multi-object) history **H** is legal if
 - For every object **x**
 - **H | x** is in the sequential specification for **x**

Precedence

H =

 A q.enq(3)

 B p.enq(4)

 B p:void

 A q:void

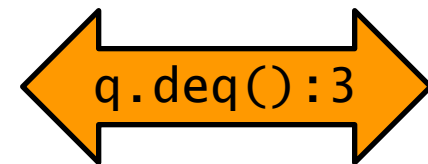
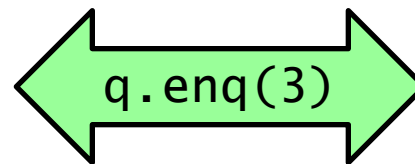
 B q.deq()

 B q:3

A method call precedes

 another if response event

 precedes invocation event



Non-Precedence

H =

 A q.enq(3)

 B p.enq(4)

 B p:void

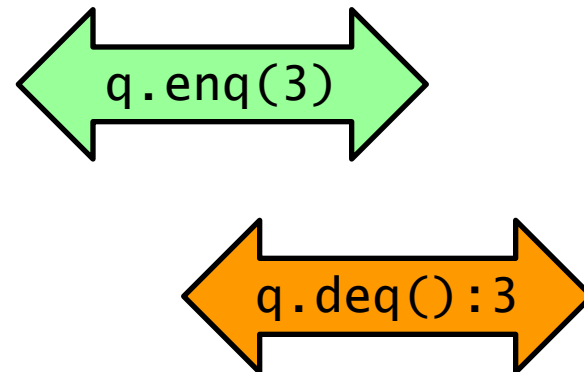
 B q.deq()

 A q:void

 B q:3

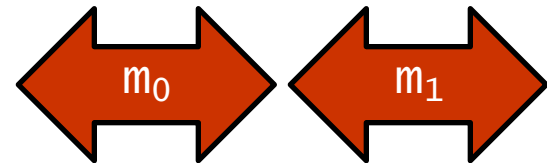
Some method calls

overlap one another



Notation

- Given
 - History H
 - Method executions m_0 and m_1 in H
- We say $m_0 \rightarrow_H m_1$, if
 - m_0 precedes m_1
- Relation $m_0 \rightarrow_H m_1$ is a
 - Partial order
 - Total order if H is sequential



Linearizability

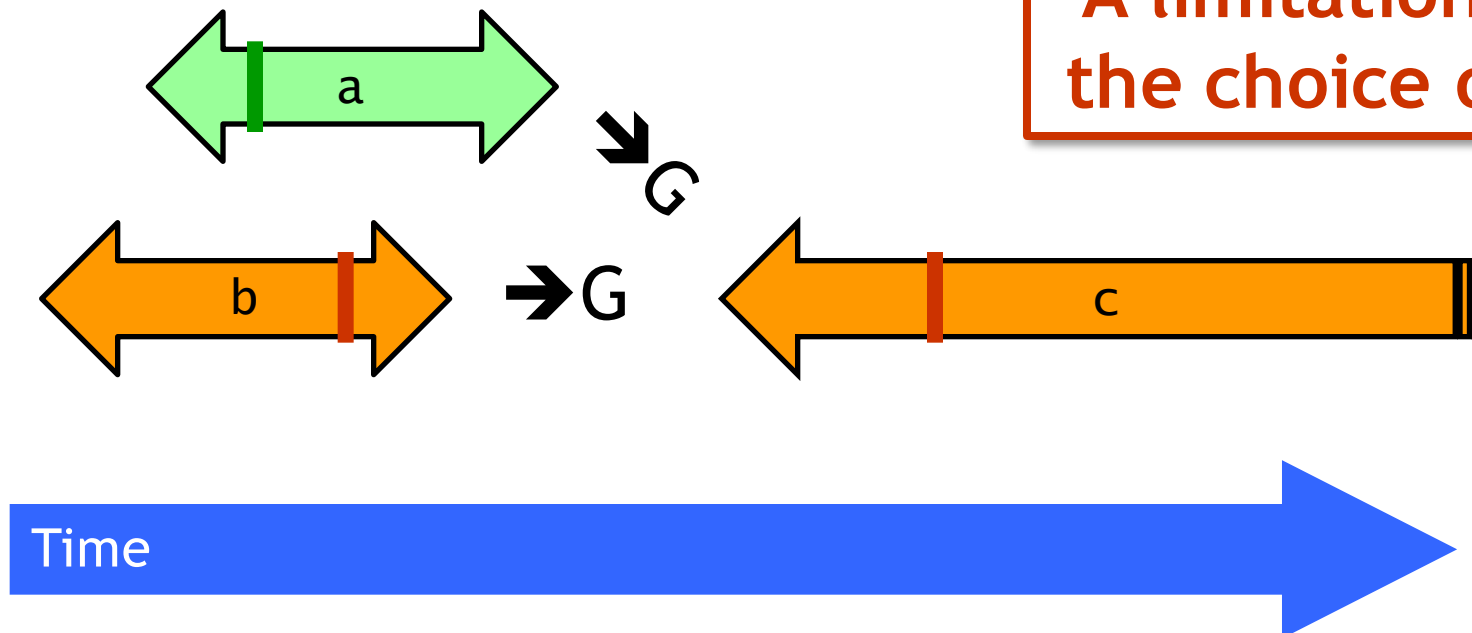
- History **H** is **linearizable** if it can be extended to **G** by
 - Appending zero or more responses to pending invocations
 - Discarding other pending invocations
- So that **G** is equivalent to
 - Legal sequential history **S**
 - Where $\rightarrow_G \subset \rightarrow_S$

What is $\rightarrow_G \subset \rightarrow_S$?

$$\rightarrow_G = \{a \rightarrow c, b \rightarrow c\}$$

$$\rightarrow_S = \{a \rightarrow b, a \rightarrow c, b \rightarrow c\}$$

A limitation on the choice of S!



Remarks

- Some pending invocations
 - Took effect, so keep them
 - Discard the rest
- Condition $\rightarrow_G \sqsubset \rightarrow_S$
 - Means that **S** respects “real-time order” of **G**

Example

Complete this pending invocation

A `q.enq(3)`

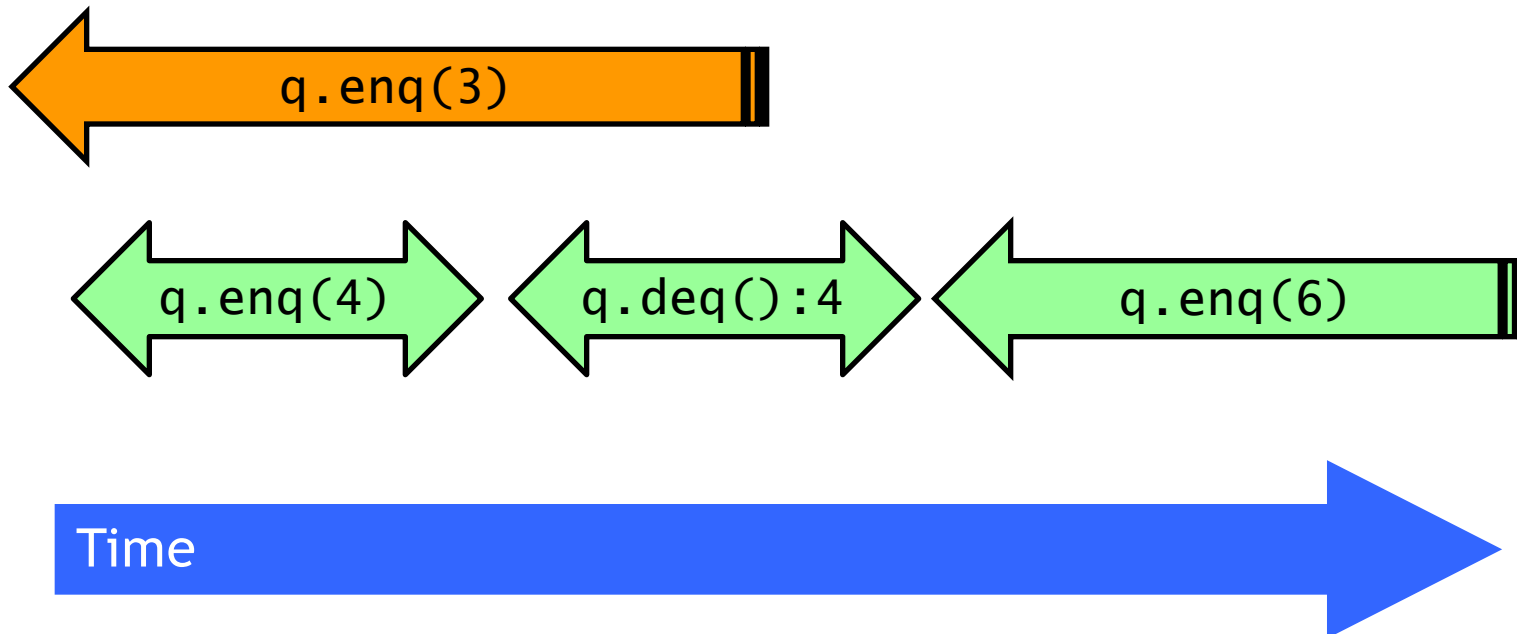
B `q.enq(4)`

B `q:void`

B `q.deq()`

B `q:4`

B `q:enq(6)`



Example

Complete this pending invocation

A `q.enq(3)`

B `q.enq(4)`

B `q:void`

B `q.deq()`

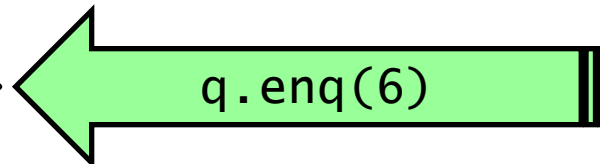
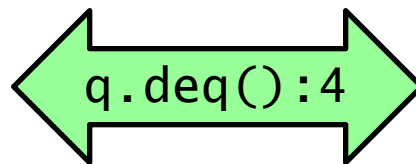
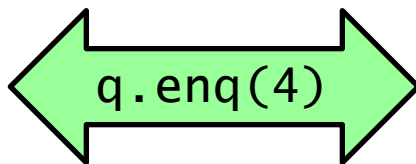
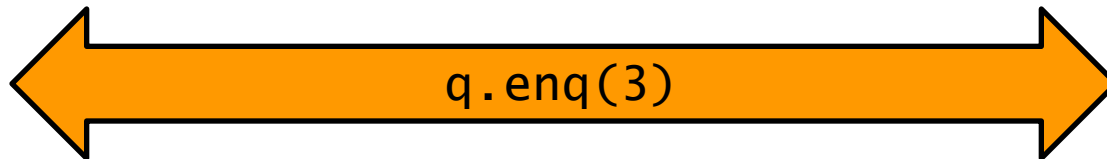
B `q:4`

Discard this one

Added response

B `q:enq(6)`

A `q:void`

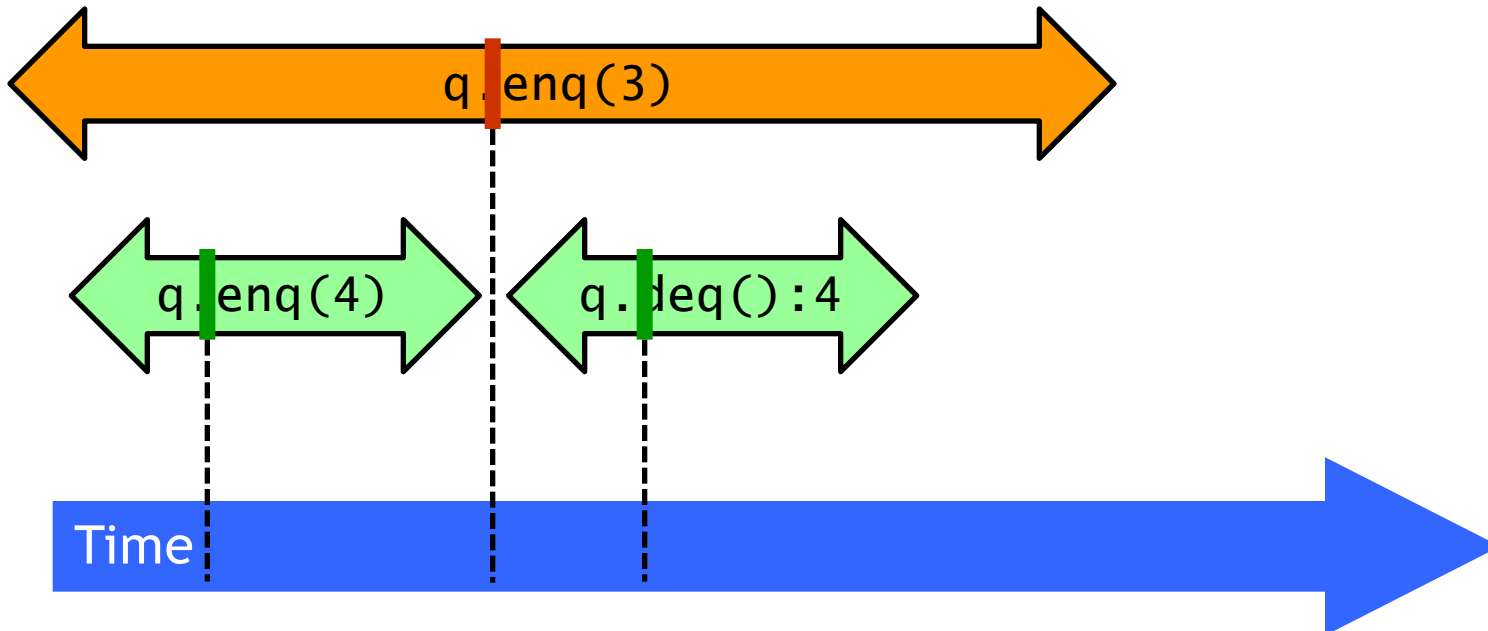


Example

A q.enq(3)
 B q.enq(4)
 B q:void
 B q.deq()
 B q:4
 A q:void

Equivalent
 Sequential
 history

B q.enq(4)
 B q:void
 A q.enq(3)
 A q:void
 B q.deq()
 B q:4



Locality Theorem

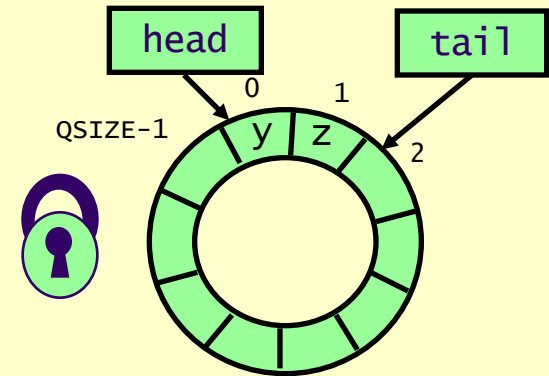
- History **H** is linearizable if and only if
 - For every object **x**
 - **H | x** is linearizable
- We care about objects only!
- Why Does Locality Matter?
 - Modularity
 - Can prove linearizability of objects in isolation
 - Can compose independently-implemented objects

Linearizability: Locking

```

public class Queue {
    int head = 0, tail = 0;
    Object[QSIZE] items;

    public synchronized
    void enq(Object x) {
        while (tail - head == QSIZE)
            this.wait();
        items[tail % QSIZE] = x;
        tail++;
        this.notifyAll();
    } ...
}
  
```



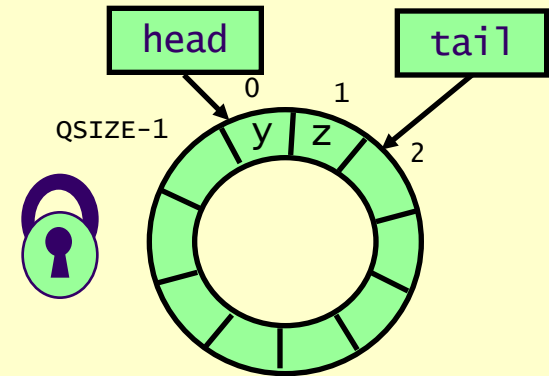
As we said, the linearization order is order lock acquired

Linearizability: Lock-free

```

public class LockFreeQueue {
    volatile int head = 0,
               tail = 0;
    Object[QSIZE] items;
    public void enq(Item x) {
        while (tail-head == QSIZE); // busy-wait
        items[tail % QSIZE] = x; tail++;
    }
    public Item deq() {
        while (tail == head); // busy-wait
        Item item = items[head % QSIZE];
        return item;
    }
}

```



Linearization order is order
head and tail fields modified

Strategy

- Identify one atomic step where method “happens”
 - Critical section
 - Machine instruction
- Does not always work
 - Might need to define several different steps for a given method

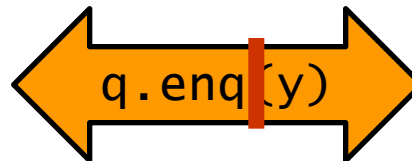
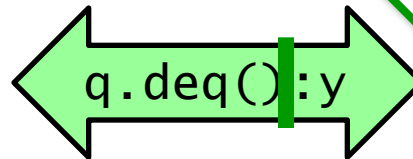
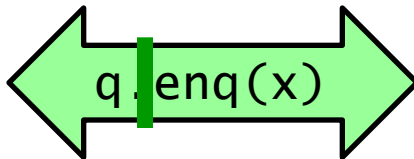
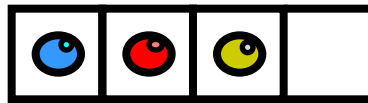
Alternative: Sequential Consistency

- History **H** is **sequentially consistent** if it can be extended to **G** by
 - Appending zero or more responses to pending invocations
 - Discarding other pending invocations
- So that **G** is equivalent to a
 - Legal sequential history **S**
 - ~~Where $\rightarrow G \subseteq \rightarrow S$~~ Differs from linearizability

Alternative: Sequential Consistency

- No need to preserve real-time order
 - Cannot re-order operations done by the same thread
 - Can re-order non-overlapping operations done by different threads
- Often used to describe multiprocessor memory architectures

Example



Not linearizable

Yet sequentially consistent!

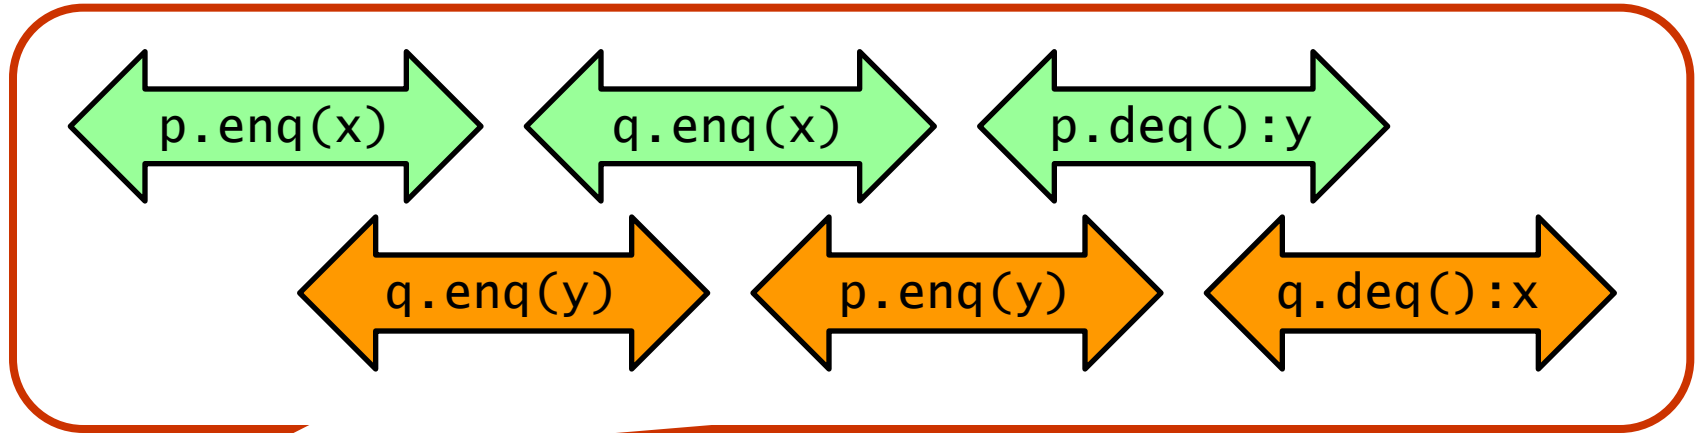
Time



Theorem

- Sequential consistency is not a local property
 (and thus we loose composability...)

FIFO Queue Example

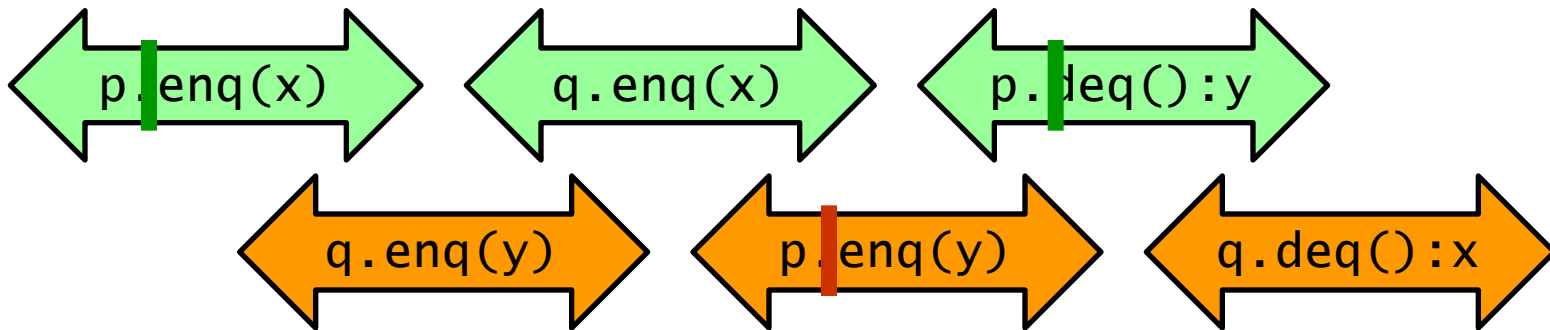


History H

Time



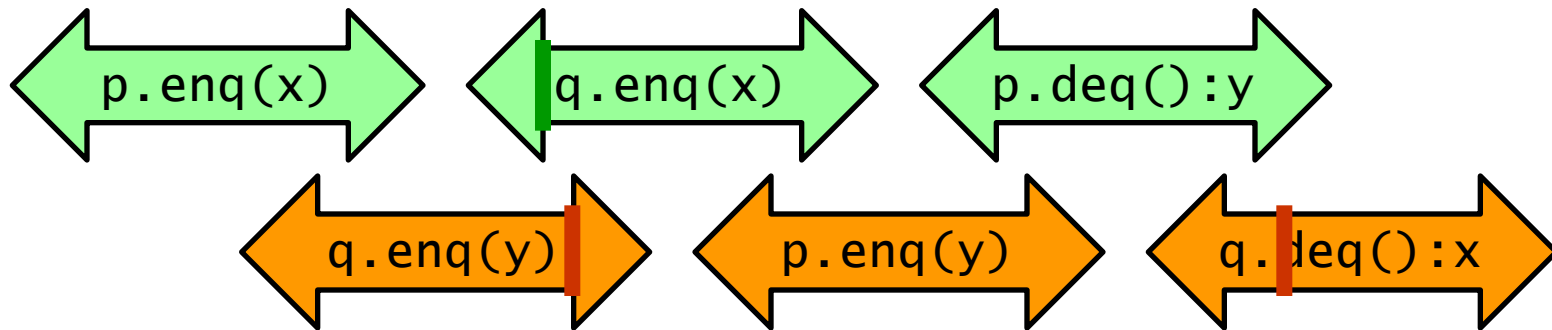
H | p Sequentially Consistent



Time



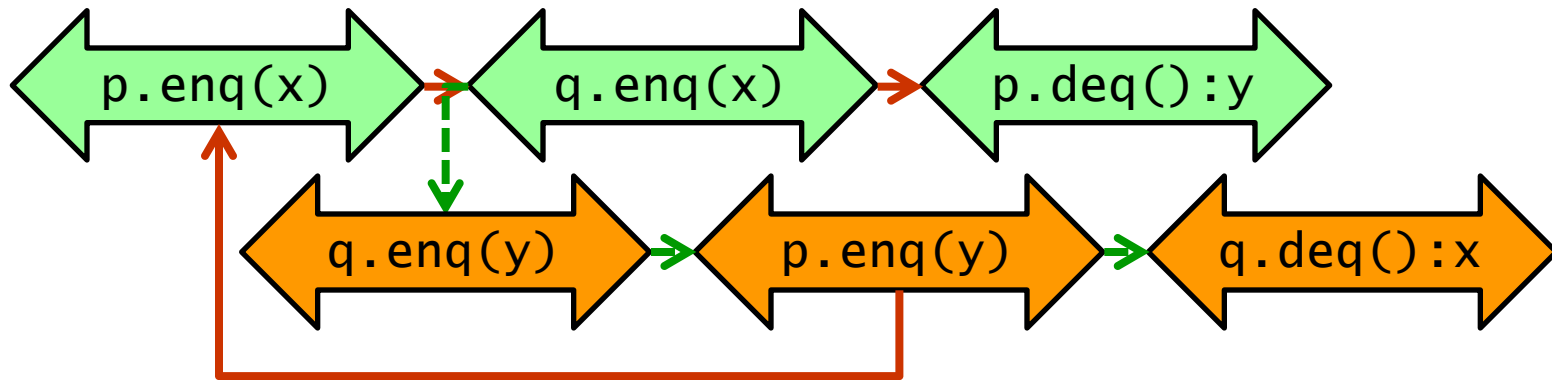
H | q Sequentially Consistent



Time



Ordering Imposed by **p** and **q**



Cannot satisfy both!

Time

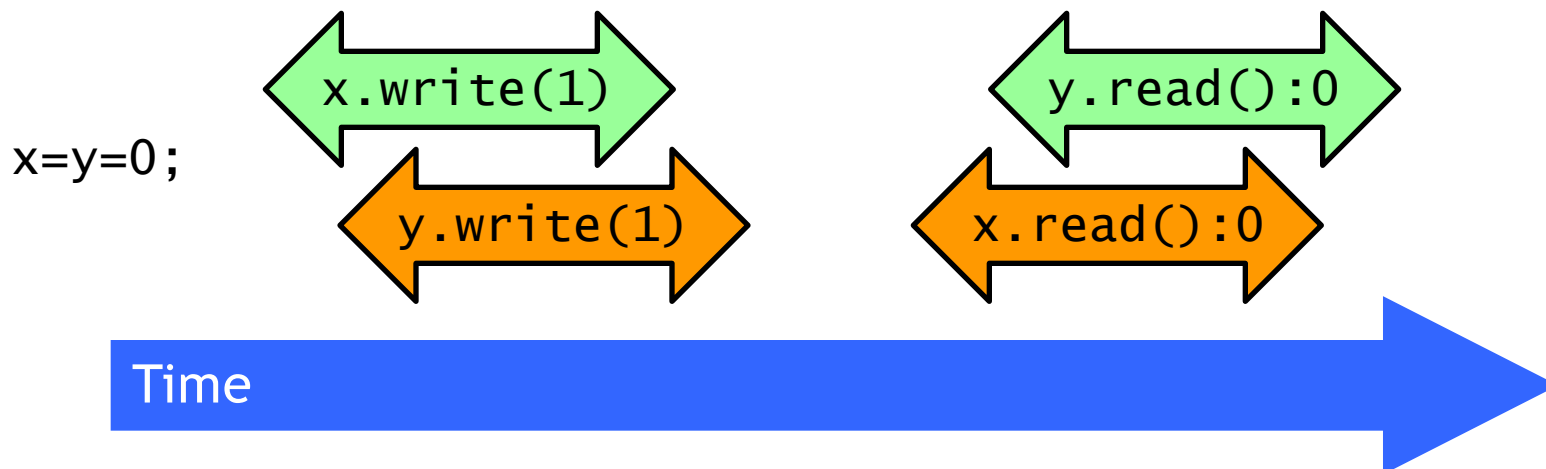


Fact

- Most hardware architectures do not support sequential consistency
- Because they think it is too strong
- Here is another story...

The Flag Example

- Each thread's view is sequentially consistent
 - It went first
- Entire history is not sequentially consistent
 - Cannot both go first
- Is this behavior really so wrong?



Opinion 1: It is Wrong!

- This pattern
 - Write mine, read yours
- Is exactly the flag principle
 - Beloved of Alice and Bob
 - Heart of mutual exclusion
 - Peterson
 - Bakery, etc.
- It is non-negotiable!

Opinion 2: But It Feels So Right...

- Many hardware architects think that sequential consistency is too strong
- Too expensive to implement in modern hardware
- OK if flag principle
 - Violated by default
 - Honored by explicit request

Memory Hierarchy

- On modern multiprocessors, processors do not read and write directly to memory
- Memory accesses are **very slow** compared to processor speeds
- Instead, each processor reads and writes directly to a **cache**
 - To read a memory location: load data into cache and read from cache
 - To write a memory location: update cached copy and **lazily** write cached data back to memory

While Writing to Memory

- A processor can execute **hundreds**, or even **thousands** of instructions
- Why delay on every memory write?
- Instead, write back in parallel with rest of the program

Revisionist History

- Flag violation history is actually OK
 - Processors delay writing to memory
 - Until after reads have been issued
- Otherwise unacceptable delay between read and write instructions
- Who knew you wanted to synchronize?

Synchronizing

- Writing to memory = mailing a letter
- Vast majority of reads & writes
 - Not for synchronization
 - No need to idle waiting for post office
- If you want to synchronize
 - Announce it explicitly
 - Pay for it only when you need it

Explicit Synchronization

- Memory barrier instruction
 - Flush unwritten caches
 - Bring caches up to date
- Compilers often do this for you
 - Entering and leaving critical sections
- Expensive

Volatile

- In Java, can ask compiler to keep a variable up-to-date with volatile keyword
- Also inhibits reordering, removing from loops, & other “optimizations”

Real-World Hardware Memory

- Weaker than sequential consistency
- But you can get sequential consistency at a price
- OK for expert, tricky stuff
 - Assembly language, device drivers, etc.
- Linearizability more appropriate for high-level software

Critical Sections

- Easy way to implement linearizability
 - Take sequential object
 - Make each method a critical section
- Like synchronized methods in Java
- Problems
 - Blocking
 - No concurrency

Summary

- Linearizability
 - Powerful specification tool for shared objects
 - Allows us to capture the notion of objects being “atomic”
 - Operation takes effect instantaneously between invocation and response
 - Uses sequential specification, locality implies composability
 - Good for high level objects

Summary

- Sequential Consistency
 - Not composable
 - Harder to work with
 - Good way to think about hardware models
- We will use **linearizability** in the remainder of this course unless stated otherwise