

# Concurrency: Multi-core Programming & Data Processing

## Linked Lists: Locking, Lock-Free, and Beyond

*Prof. P. Felber*

[Pascal.Felber@unine.ch](mailto:Pascal.Felber@unine.ch)


<http://iiun.unine.ch/>

*Based on slides by Maurice Herlihy and Nir Shavit*



Since you link the nodes in a linear fashion.  
Intuitively there is only limited thing you can  
do concurrently

# Last Lecture: Spin-Locks



# Today: Concurrent Objects

- Adding threads...
- Should not lower throughput
  - Contention effects
  - Mostly fixed by queue locks
- Should increase throughput Only possible if app is not inherently sequential
  - Not possible if inherently sequential
  - Surprising things are parallelizable

# Coarse-Grained Synchronization

- Each method locks the object
  - Avoid contention using queue locks
  - Easy to reason about
    - In simple cases
  - Standard Java model
    - Synchronized blocks and methods
- So, are we done?

# Coarse-Grained Synchronization

- Sequential bottleneck
  - All threads “stand in line”
- Adding more threads
  - Does not improve throughput
  - Struggle to keep it from getting worse
- So why even use a multiprocessor?
  - Well, some applications inherently parallel...

# This Lecture

- Introduce four “patterns”
  - Bag of tricks
  - Methods that work more than once
- For highly-concurrent objects
- Goal
  - Concurrent access
  - More threads, more throughput

e.g. central db is bottleneck. we want this to be highly parallel

# 1. Fine-Grained Synchronization

- Instead of using a single lock...
- Split object into **independently-synchronized components**
- Methods conflict when they access
  - The same component...
  - At the same time

## 2. Optimistic Synchronization

- Object = linked set of components
- Search without locking...
- If you find it, lock and check...
  - OK, we are done
  - Oops, try again
- Evaluation
  - Cheaper than locking
  - Mistakes are expensive

Try to optimistically go through structure. Only when you find what you're looking for, you lock. You need to do some checks on what you have locked. Only makes sense when that check is less expensive than locking.



# 3. Lazy Synchronization

Probably the most efficient

- Postpone hard work
- Removing components is tricky
  - Logical removal
    - Mark component to be deleted
  - Physical removal
    - Do what needs to be done

# 4. Lock-Free Synchronization

- Do not use locks at all
  - Use `compareAndSet()` and relatives...
- Advantages
  - Robust against asynchrony
- Disadvantages
  - Complex depends on the datastructure
  - Sometimes high overhead

The most parallel. Scales well. No locks. No mutex. Nothing to slow you down.

# Linked List

Insert/Search/Remove

Implement a set using a list. Set with no duplicated and sorted.

- Illustrate these patterns...
- Using a list-based **Set**
  - Common application
  - Building block for other apps

# Set Interface

- Collection of objects
- No duplicates
- Methods
  - `add()` a new object
  - `remove()` an object
  - Test if set `contains()` object readonly

# List-Based Sets

```


public interface Set {
  public boolean add(Object x);
  public boolean remove(Object x);
  public boolean contains(Object x);
}
  
```

Add object to set  
 Remove object from set  
 Is object in set?

# List Node

```
public class Node {  
    Object object; ➤ Object of interest  
    int key; ➤ Usually hash code  
    Node next; ➤ Reference to next node  
}
```

# The List-Based Set



Ordered + sentinel nodes  
 (min & max possible keys)

they have a special key

# Reasoning about Concurrent Sets

- Identify invariants
  - Properties that always holds
- True when object is created
- Truth preserved by each method
  - `add()`, `remove()`, `contains()`
  - Each step of each method
- Most steps are trivial
  - Usually one step tricky
  - Often linearization point



# Interference

- Proof that invariants are preserved works if methods considered are the only modifiers
- Language encapsulation helps
  - List nodes not visible outside class
- Freedom from interference needed even for removed nodes
  - Some algorithms traverse removed nodes
  - Careful with **malloc()** and **free()** !
- Garbage-collection helps here

# Blame Game


- Suppose
  - `add()` leaves behind 2 copies of `x`
  - `remove()` removes only 1
- Which one is incorrect?
  - If invariant says *no duplicates*
    - `add()` is incorrect
  - Otherwise
    - `remove()` is incorrect

# Set Invariant (Partly)


- Sentinel nodes
  - Tail reachable from head
- Sorted according to hash key
- No duplicates

# Sequential List Based Set

add(b)

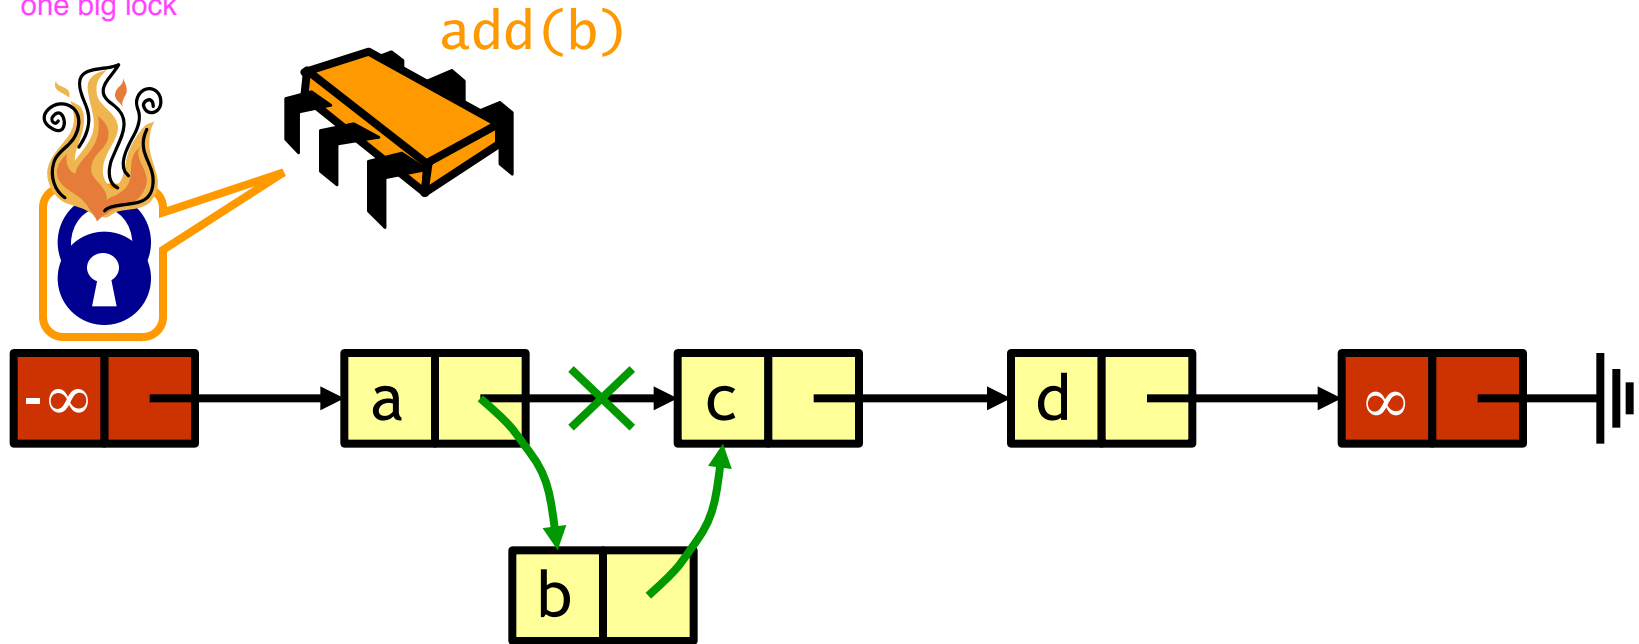


remove(b)



# Coarse-Grained Locking

We protect the whole list with one big lock



Simple but **hotspot + bottleneck**

No parallelism

# Coarse-Grained Locking

- Easy, same as synchronized methods
  - “One lock to rule them all...”
- Simple, clearly correct
  - Deserves respect!
- Works poorly with contention
  - Queue locks help
  - But bottleneck still an issue

# Fine-grained Locking


- Requires careful thought
  - “Do not meddle in the affairs of wizards, for they are subtle and quick to anger”
- Split object into pieces
  - Each piece has own lock
  - Methods that work on disjoint pieces need not exclude each other

For instance, lock individual nodes. It's not very easy to do. You need to hold multiple locks in the right order. Deadlock: Each thread has one lock and is waiting for other threads to release locks.

# Hand-over-Hand Locking


You need to make sure that you always keep at least one lock.

Felber Hand over movement



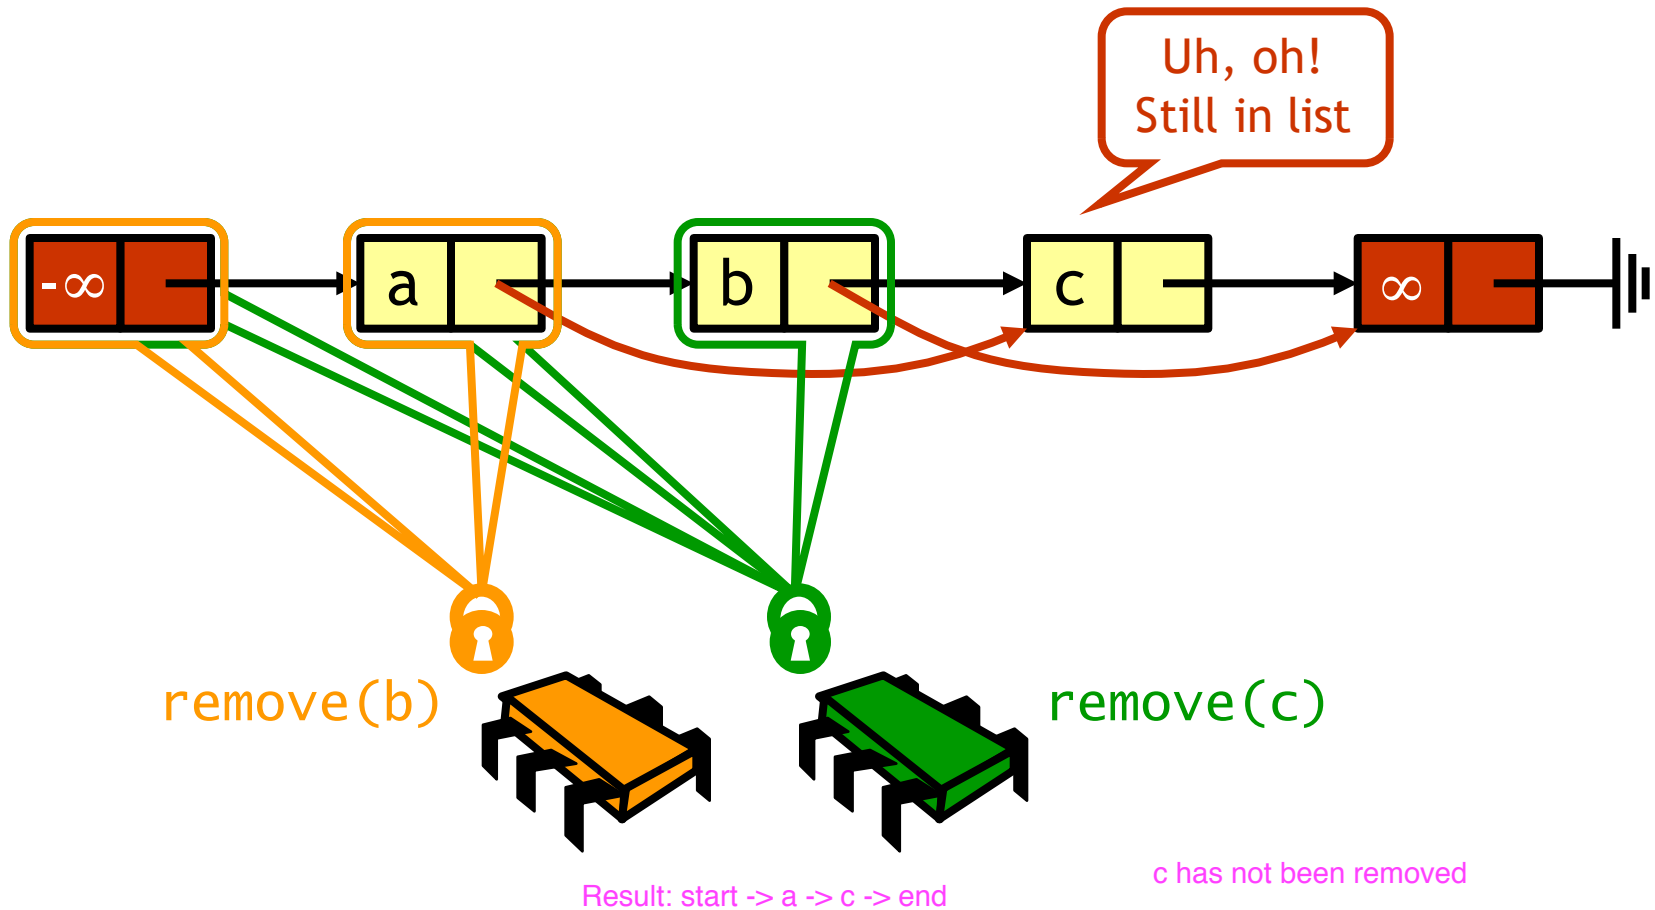


# Removing a Node




# Removing a Node

Two threads removing at the same time




# Problem

- To delete node **b**
  - Swing node **a**'s next field to **c**



- Problem is
  - Someone could delete **c** concurrently




so we need to hold more than a single lock to make this modification

# Insight

- If a node is locked
  - No one can delete node's *successor*
- If a thread locks
  - The node to be deleted
  - And its predecessor
  - Then it works


# Hand-over-Hand Again

so we lock a AND b



# Hand-over-Hand Again

see power point for animation



# Remove

```

public boolean remove(Object object) {
  int key = object.hashCode();
  Node pred, curr;
  try {
    ...
  } finally {
    curr.unlock();
    pred.unlock();
  }
}

```

Key used to order node  
 Predecessor and current nodes  
 Everything else  
 Make sure locks released

# Remove

```

try {
  pred = this.head;
  pred.lock();
  curr = pred.next;
  curr.lock();
  ...
} finally { ... }

```

Lock previous
   
 Lock current
   
 Traverse the list



# Remove: Searching

```

while (curr.key <= key) {
  if (object == curr.object) {
    pred.next = curr.next;
    return true;
  }
  pred.unlock();
  pred = curr;
  curr = curr.next;
  curr.lock();
}
return false;

```

Search key range (curr and pred locked)

If node found, remove it

Unlock predecessor and demote current (only one node locked!)

Find and lock new current

Otherwise not present

# Adding Nodes

- To add node **b**
  - Lock predecessor
  - Lock successor
- Neither can be deleted
  - (Is successor lock actually required?)

lock around insertion point

no, not necessarily. it's optional

# Drawbacks

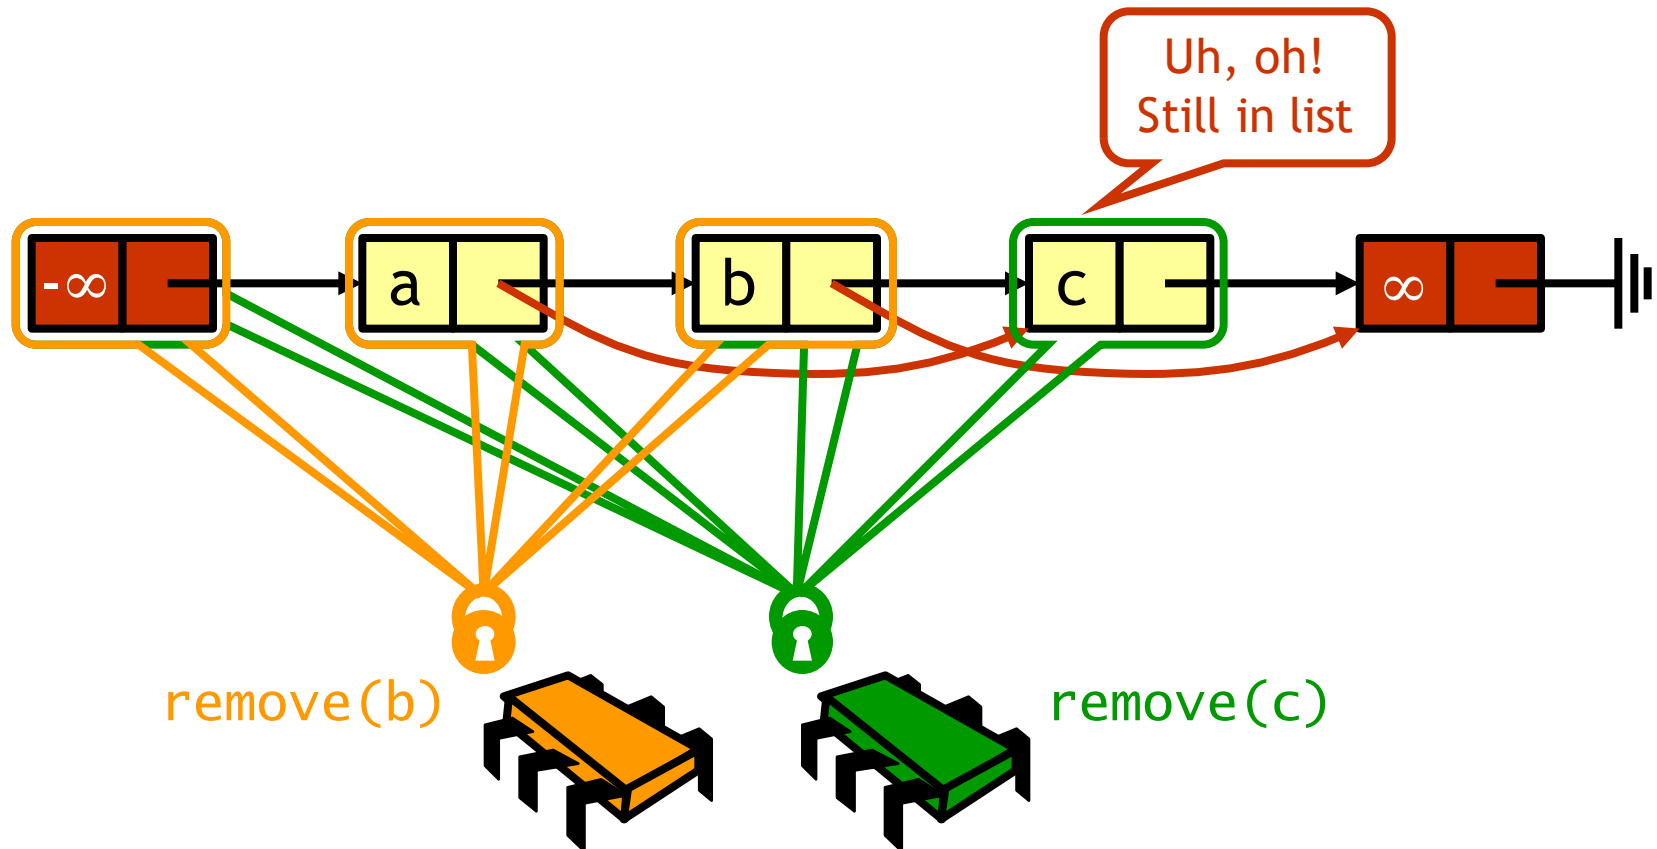
- Better than coarse-grained lock
  - Threads can traverse in parallel
- Still not ideal
  - Long chain of acquire/release
  - Threads cannot overtake one another
  - Inefficient

# Optimistic Synchronization

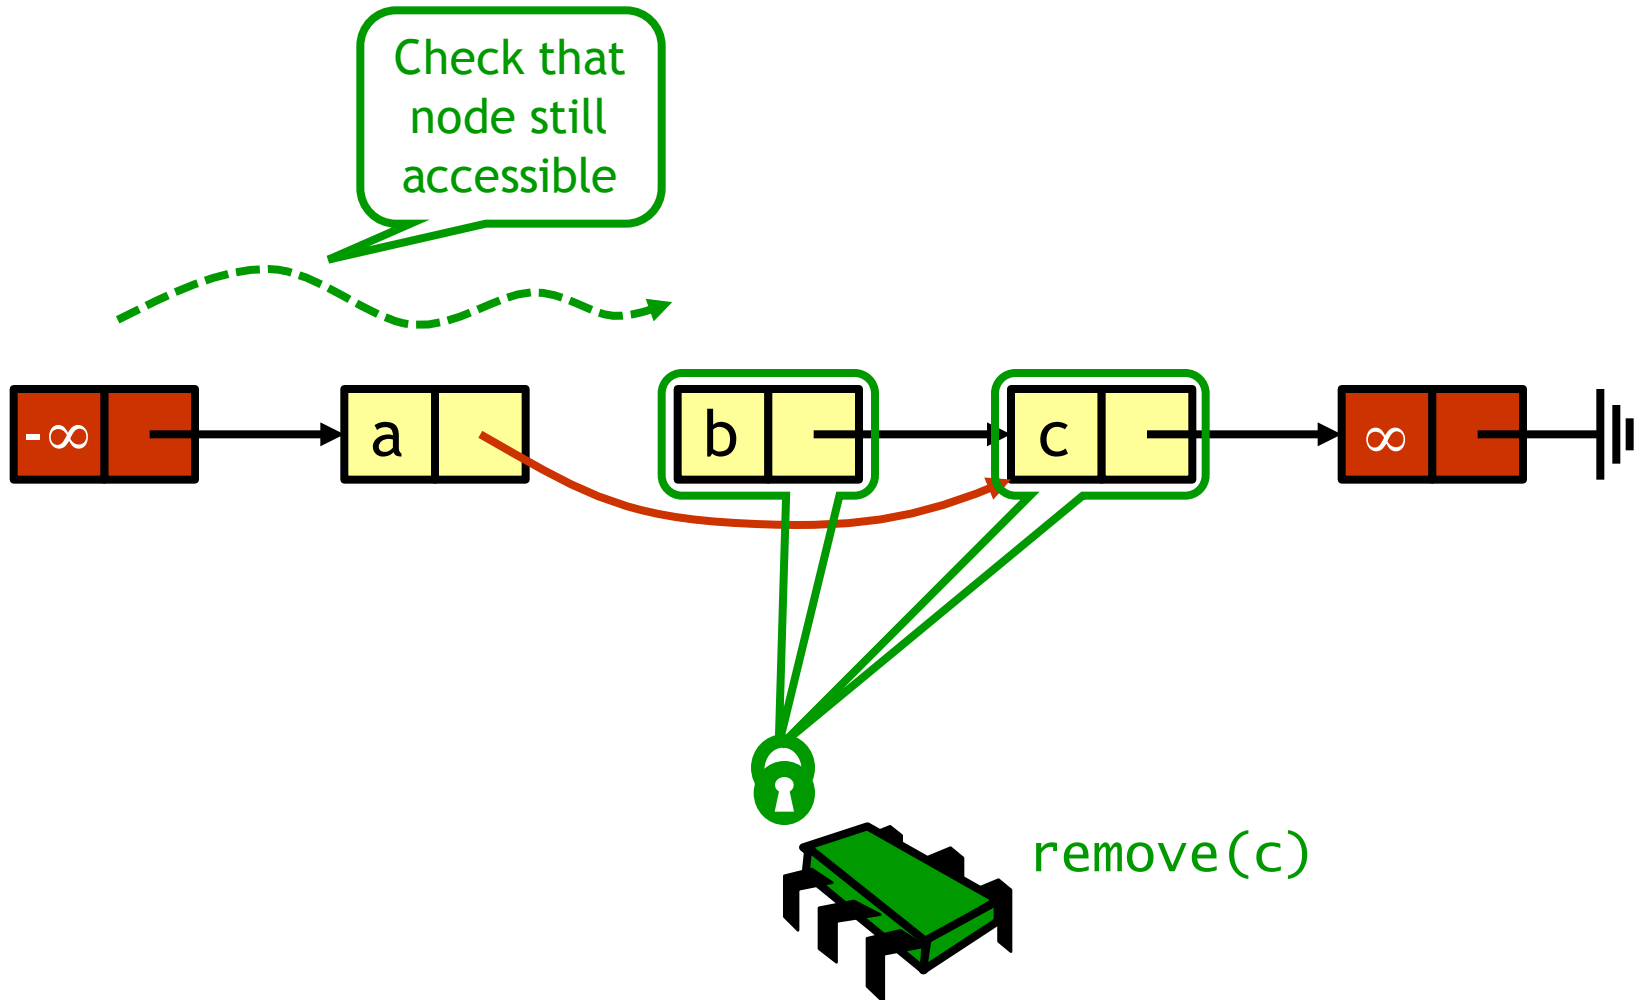
- Find nodes without locking
- Lock nodes
- Check that everything is OK

# What Could Go Wrong?

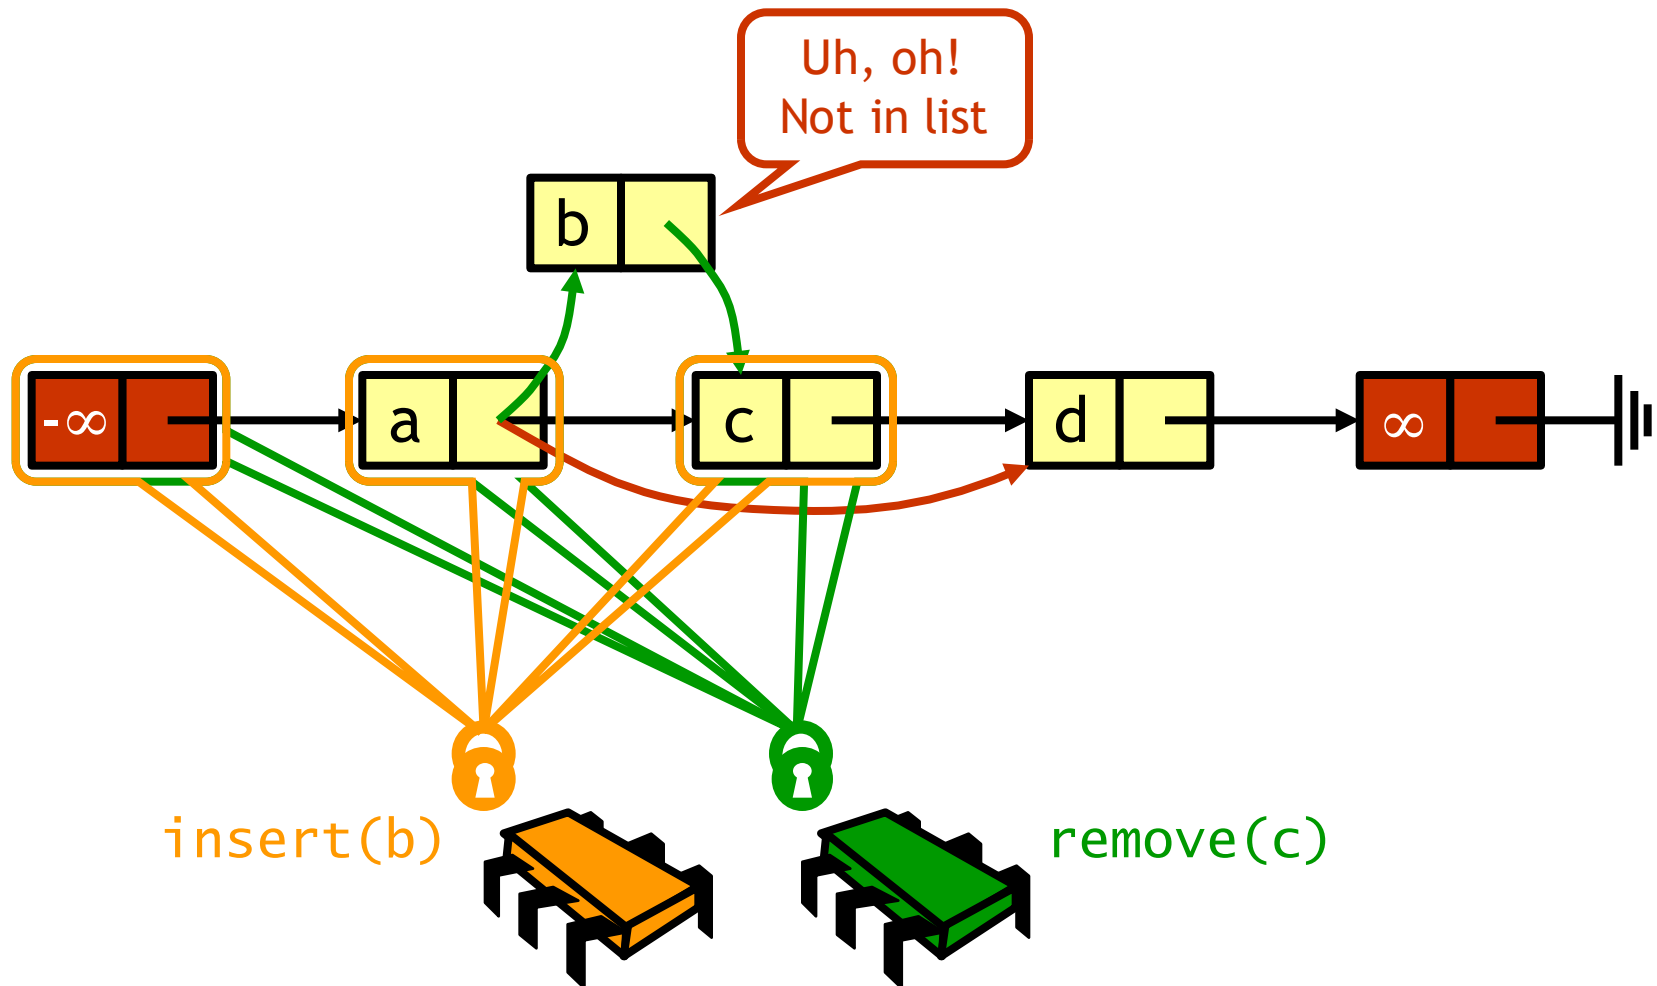
See powerpoint for animation



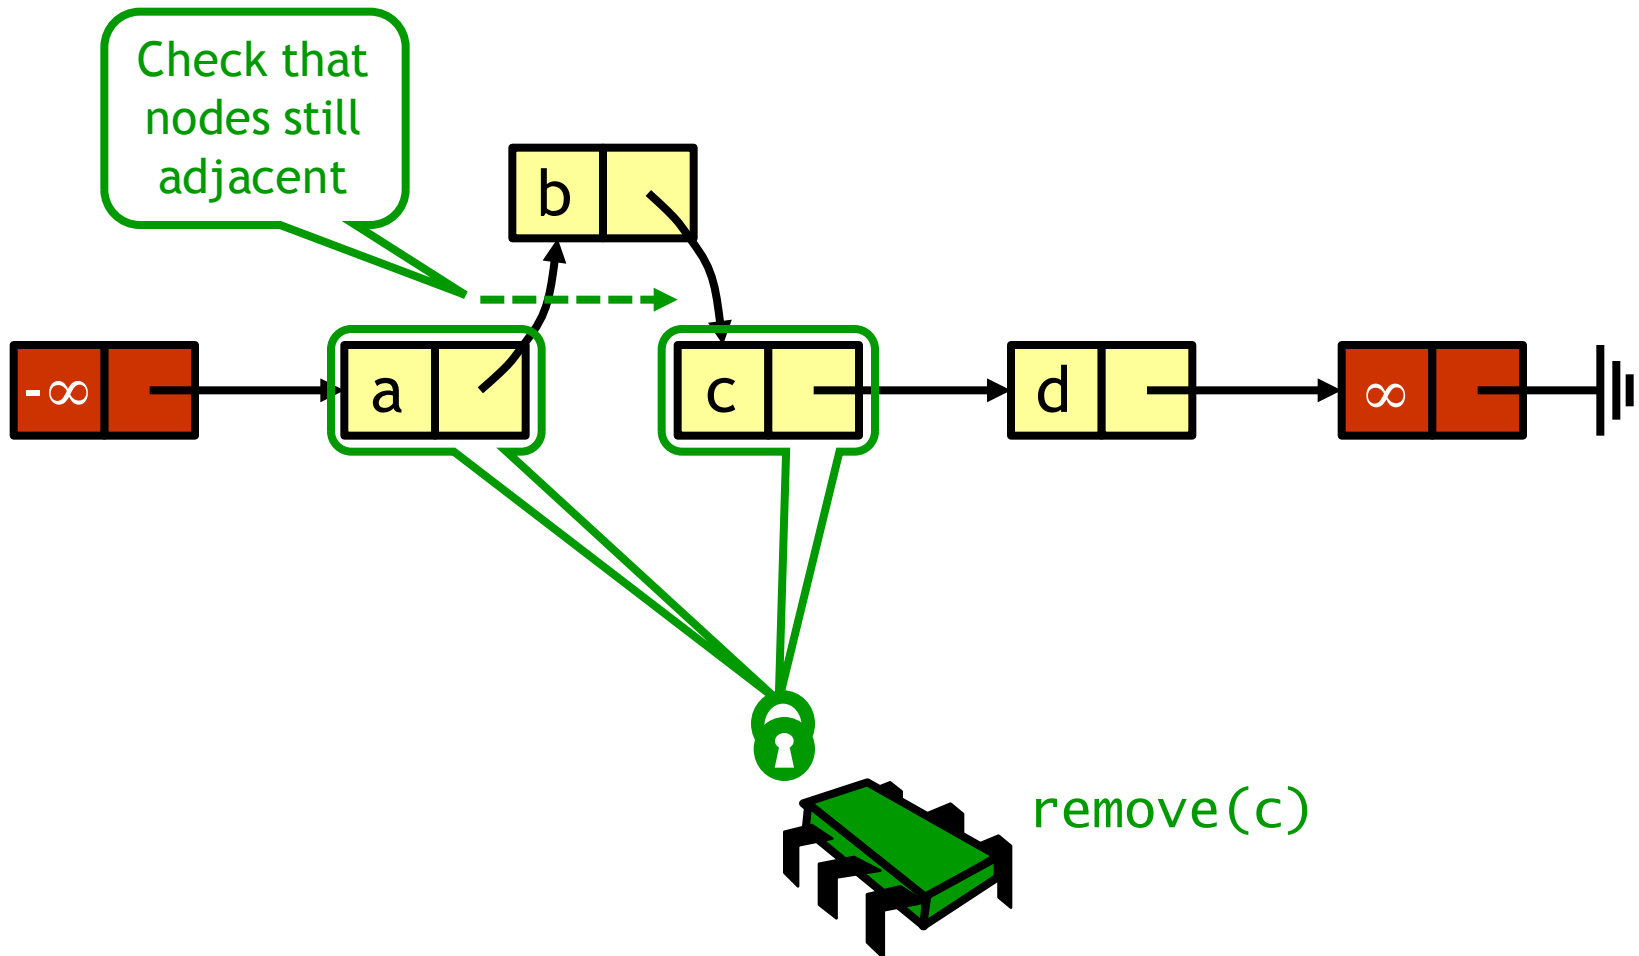
# What Could Go Wrong?



# What Could Go Wrong?




# What Could Go Wrong?






# Optimistic Fine Grained



- To insert **c**

1. Optimistically traverse list to find **b**
2. Lock **b.curr** then lock **b.succ**
3. **Re-Traverse** list to find **b** and verify **b.curr** precedes **b.succ**  
this means b hasn't been deleted, and successor is also there
4. Perform insertion and release locks

# Optimistic Fine Grained removal



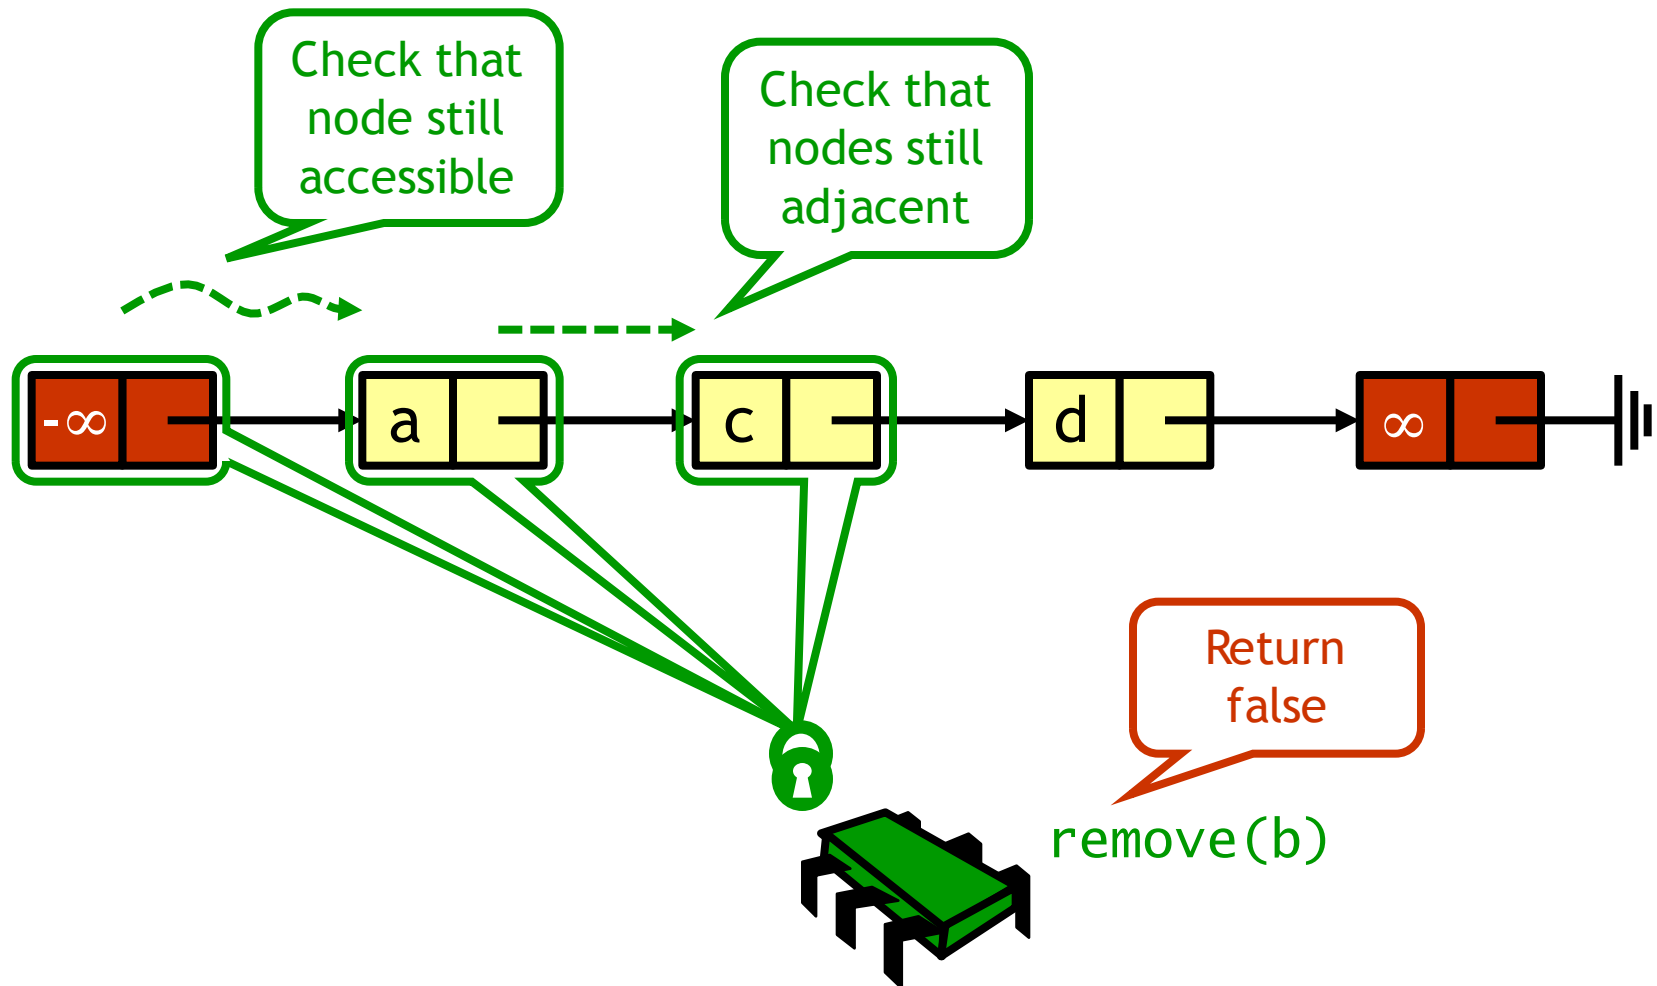
- To remove **c**

1. Optimistically traverse list to find **c**
2. Lock **c.pred** then lock **c.curr**
3. Re-Traverse list to find **c** and verify that **c.pred** precedes **c.curr**
4. Perform removal and release locks

# Correctness

- If
  - Nodes **b** and **c** both locked
  - Node **b** still accessible
  - Node **c** still successor to **b**
- Then
  - Neither will be deleted
  - OK to delete and return **true**

# Removing an Absent Node



# Correctness

- If
  - Nodes **a** and **c** both locked
  - Node **a** still accessible
  - Node **c** still successor to **a**
- Then
  - Neither will be deleted
  - No thread can add **b** after **a**
  - OK to return **false**

# Validation

```
private boolean
```

```
    validate(Node pred,
              Node curr) {
```

Predecessor & current nodes

```
        Node node = head;
```

Start at the beginning

```
        while (node.key <= pred.key) {
```

Search range of keys

```
            if (node == pred)
```

Predecessor reachable?

```
                return pred.next == curr;
```

Current node next?

```
                node = node.next;
```

Otherwise move on

```
        }
```

```
        return false;
```

Predecessor not reachable

```
    }
```

# Remove: Searching

```

public boolean remove(Object object) {
    int key = object.hashCode();
    while (true) {
        Node pred = this.head;
        Node curr = pred.next;
        while (curr.key <= key) {
            if (object == curr.object)
                break;
            pred = curr;
            curr = curr.next;
        }
    }
    ...
  
```

Search key

Retry on synchronization conflict

Examine predecessor and current nodes

Search by key

Stop if we find object

Move along

# On Exit from Loop

- If object is present
  - `curr` holds object
  - `pred` just before `curr`
- If object is absent
  - `curr` has first higher key
  - `pred` just before `curr`
- Assuming no synchronization problems



# Remove

```

try {
    pred.lock(); curr.lock();
    if (validate(pred, curr)) {
        if (curr.object == object) {
            pred.next = curr.next;
            return true;
        } else {
            return false;
        }
    }
} finally {
    pred.unlock(); curr.unlock();
} ...
  
```

Lock both nodes  
 Check for synchronization conflicts  
 noone has inserted/removed nodes  
 Object found, remove node  
 Object not found  
 Always unlock

# Summary: Optimistic List

- **Wait-free** traversal no thread blocks any other
  - May traverse removed nodes
  - **Must have non-interference** (natural in languages with GC like Java)
- Limited hotspots
  - Only at locked **add()**, **remove()**, **contains()** destination locations, not traversals
- But two traversals was good when traversing was cheaper than locks
  - Yet traversals are wait-free

# So Far, So Good

- Much less lock acquisition/release
  - Performance
  - Concurrency
- Problems
  - Need to traverse list twice
  - `contains()` acquires locks
    - Most common method call (90% in many applications)
- Optimistic works if
  - Cost of scanning twice without locks < cost of scanning once with locks


# Lazy List

- Like optimistic, except
  - Scan once
  - `contains()` never locks
- Key insight
  - Removing nodes causes trouble
  - Do it “lazily”

# Lazy List

- Remove Method
  - Scans list (as before)
  - Locks predecessor & current (as before)
- Logical delete
  - Marks current node as removed (new!)
  - Use additional mark bit in node
- Physical delete
  - Redirects predecessor's next (as before)

# Lazy Removal



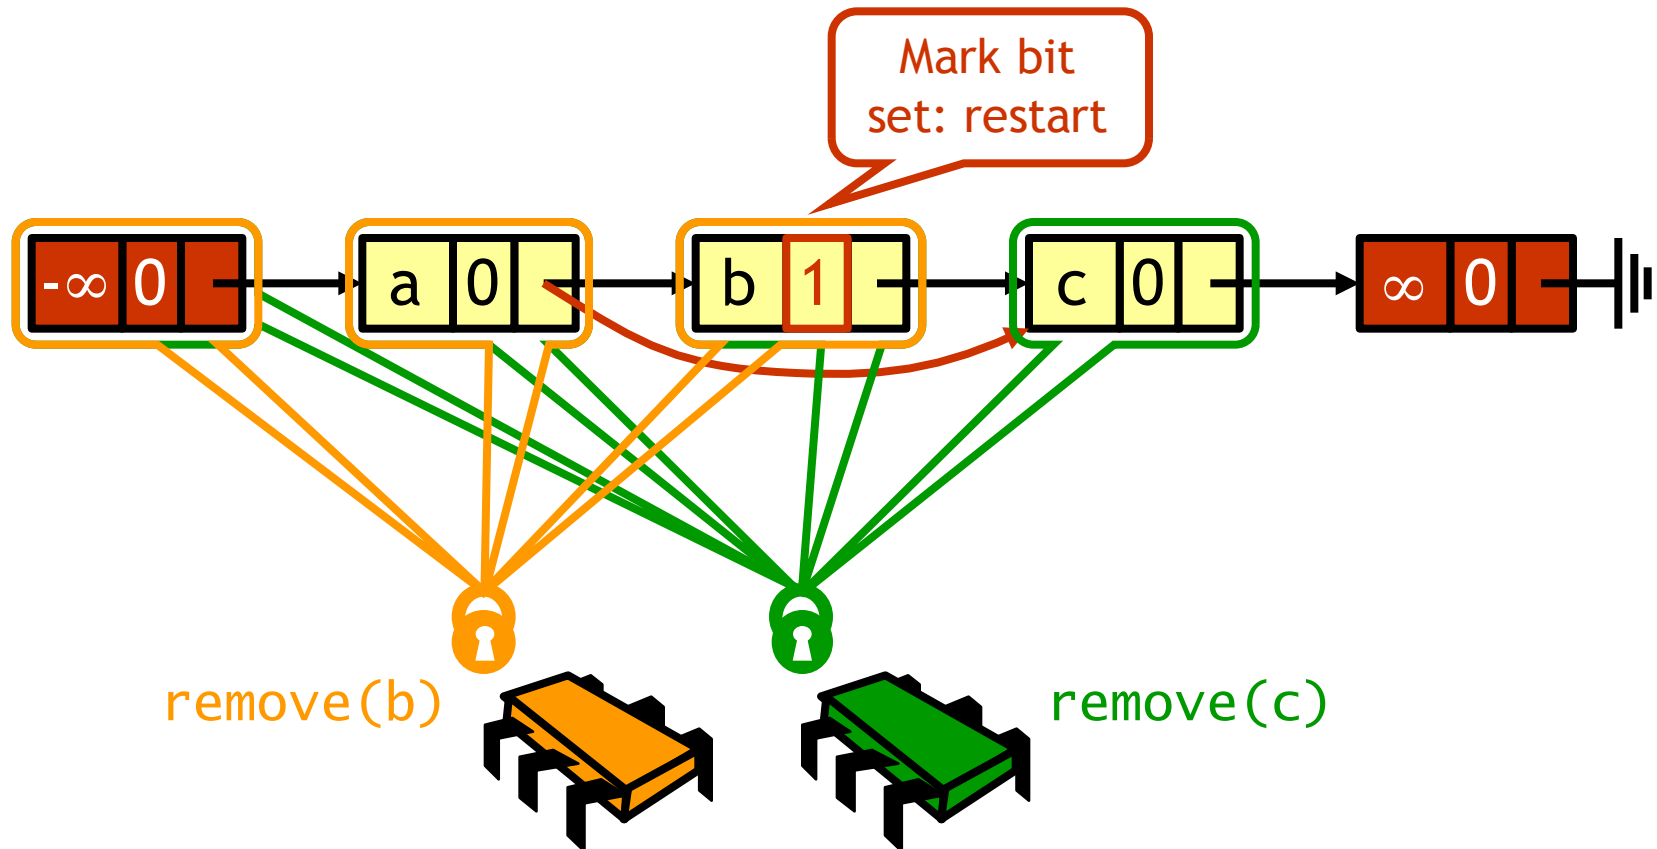
- To remove **c**

1. Optimistically traverse list to find **c**
2. Lock **c.pred** then lock **c.curr**
3. Verify marks and that **c.pred** precedes **c.curr**
4. Set mark bit (logical removal)
5. Perform physical removal and release locks

# Lazy List

- All Methods
  - Scan through locked and marked nodes
  - Removing a node does not slow down other method calls...
  - Must still lock **pred** and **curr** nodes for updates
- Validation
  - No need to rescan list!
  - Check that **pred** is not marked
  - Check that **curr** is not marked
  - Check that **pred** points to **curr**

# What Could Go Wrong?





# Validation

```
private boolean
  validate(Node pred,
           Node curr) {
  return
    !pred.marked &&
    !curr.marked &&
    pred.next == curr;
}
```

Predecessor not logically removed

Current not logically removed

Predecessor still points to current

# Remove

```

try {
    pred.lock(); curr.lock();
    if (validate(pred, curr) {
        if (curr.object == object) {
            curr.marked = true;
            pred.next = curr.next;
            return true;
        } else return false;
    }
} finally {
    pred.unlock(); curr.unlock();
} ...

```

Validate as before

Object found

Logical removal

Physical removal

# Contains

```

public boolean contains(Object object) {
    int key = object.hashCode();

```

```

    Node curr = this.head;

```

Start at the head

```

    while (curr.key <= key) {
        if (object == curr.object)
            break;
        curr = curr.next;
    }

```

Traverse without  
 locking  
 (nodes may have  
 been removed)

```

    return object == curr.object && !curr.marked;
}

```

Present and undeleted?

# Summary: Lazy List

- Wait-free traversal uses mark bit + fact that list is ordered
  - Not marked  $\Rightarrow$  in the set
  - Marked or missing  $\Rightarrow$  not in the set
- Lazy **add()**
- Lazy **remove()**
- Wait-free **contains()**

# Evaluation

- Good
  - `contains()` does not need to lock
    - In fact, it is wait-free!
    - Good because it is typically called often
  - Uncontended calls do not re-traverse
- Bad
  - Contended calls do re-traverse
  - Traffic jam if one thread delays

# Traffic Jam

- Any concurrent data structure based on mutual exclusion has a weakness
- If one thread
  - Enters critical section
  - And “eats the big muffin” (stops running)
    - Cache miss, page fault, put aside by scheduler...
  - Everyone else using that lock is stuck!

# Wait/Lock/Obstruction Freedom

Wait  
freedom

“All thread always  
makes progress”

Guarantees per-  
thread progress

VS.

Lock  
freedom

“Some thread always  
makes progress”

Guarantees system-  
wide progress

VS.

Obstruction  
freedom

“Any thread that runs by itself  
for long enough makes progress”

# Lock-Free Data Structures

practically interesting

- No matter what...
  - Some thread will complete method call
  - Even if others halt at malicious times
  - Weaker than wait-free, yet
- Implies that
  - You cannot use locks
  - Um, that is why they call it lock-free




# Lock-Free Lists


- Next logical step
- Eliminate locking entirely
- `contains()` wait-free and `add()` and `remove()` lock-free
- Use only `compareAndSet()` to atomically update links

# Adding a Node


am a still pointing to c -> set atomic





# Removing a Node



# What Could Go Wrong?



remove(b)




remove(c)

# What Could Go Wrong?

- Problem
  - Method updates node's next field after node has been removed
- Solution
  - Use **AtomicMarkableReference**
  - Atomically
    - Swing reference and update flag
  - Remove in two steps
    - Set mark bit in next field
    - Redirect predecessor's pointer

# Marking a Node

- **AtomicMarkableReference** class
  - In package `java.util.concurrent.atomic`
  - Holds a reference and a mark bit



# AtomicMarkableReference


Data type

```

public class AtomicMarkableReference <T> {
    public T get(boolean[] marked);
    public boolean compareAndSet(
        T expectedRef,
        T updateRef,
        boolean expectedMark,
        boolean updateMark);
    public boolean attemptMark(
        T expectedRef,
        boolean updateMark);
    ...
}
  
```

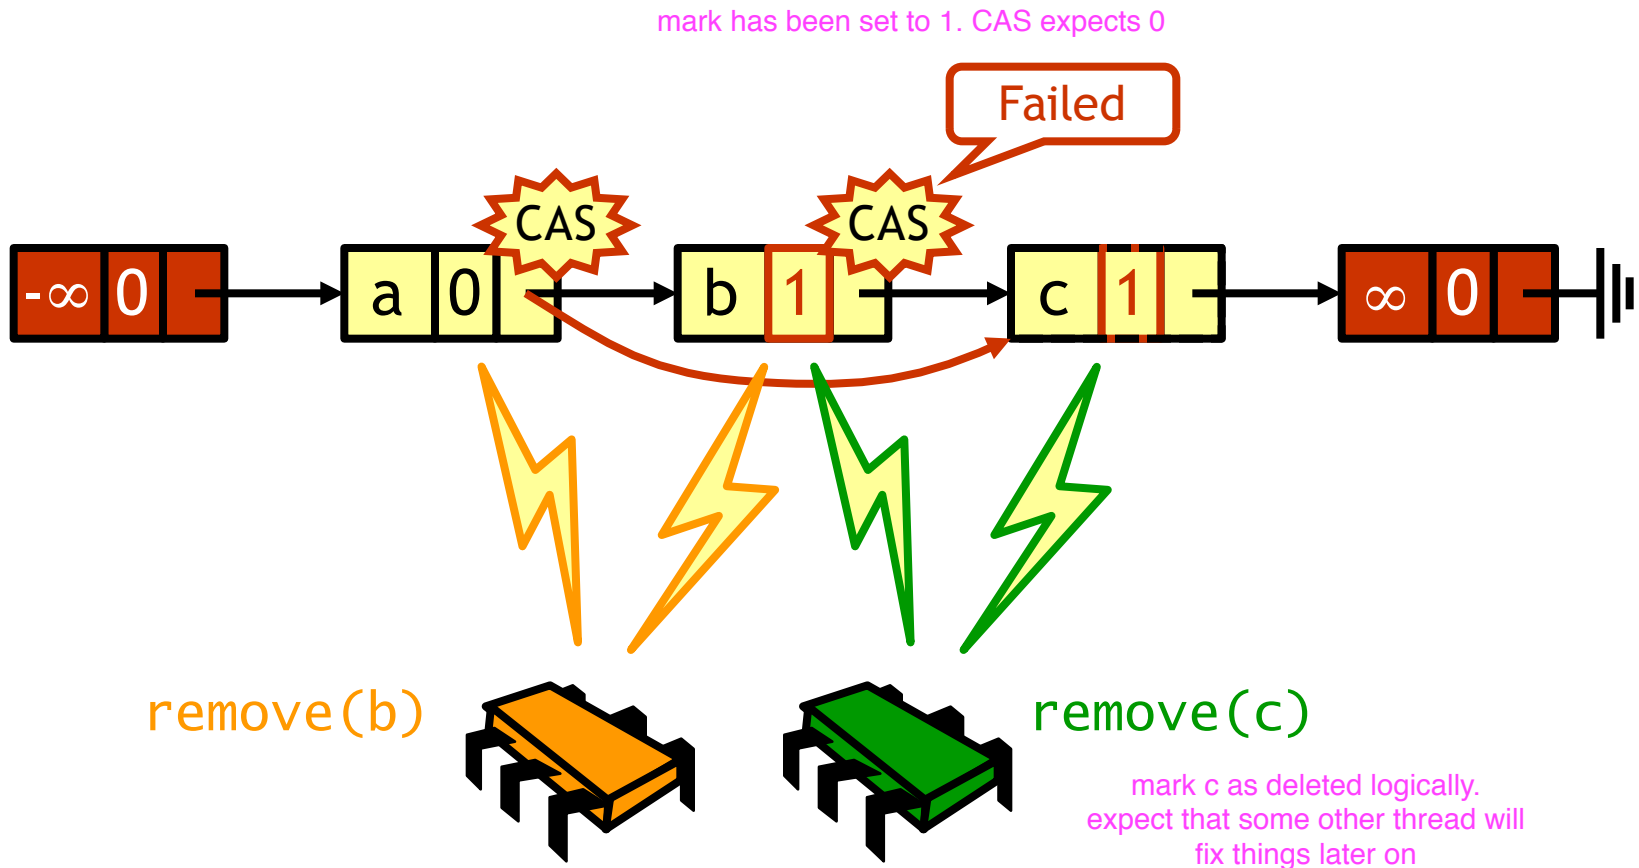
<T> {  
 Extract reference and mark (at index 0)  
 If this is the current reference...  
 ...then change to this new reference...  
 ...and this is the current mark...  
 ...and this new mark  
 If this is the current reference...  
 ...then change to this new mark

# Removing a Node






# What Could Go Wrong?



# Traversing the List

- What do you do when you find a “logically” deleted node in your path?
- **Finish the job**
  - CAS the predecessor's next field
  - Proceed (repeat as needed)

# Lock-Free Traversal



# The Window Class

```

class window {
  Node pred;
  Node curr;
  window(Node pred, Node curr) {
    this.pred = pred;
    this.curr = curr;
  }
}
  
```

A container for predecessor and current nodes

# Using the Find Method

...

```
Window window = find(head, object);
```

Find window

```
Node pred = window.pred;
```

```
Node curr = window.curr;
```

Extract pred  
and curr


...

# The Find Method


```

window window = find(head, b);
Node pred = window.pred;
Node curr = window.curr;
  
```

Object in list:



Object not in list:



# Remove

```

public boolean remove(T object) {
    boolean b;
    while (true) {
        Window window = find(head, object);
        Node pred = window.pred, curr = window.curr;
        if (curr.object != object)
            return false;
        Node succ = curr.next.getReference();
        b = curr.next.attemptMark(succ, true);
        if (!b) continue;
        pred.next.compareAndSet(curr, succ, false, false);
        return true;
    }
}
  
```

**Keep trying** (points to the `while (true)` loop)

**Find neighbors** (points to the `Window window = find(head, object);` line)

**Not there** (points to the `if (curr.object != object)` condition)

**Try to mark node as deleted** (points to the `b = curr.next.attemptMark(succ, true);` line)

**If it fails, retry, otherwise job done** (points to the `if (!b) continue;` line)

**Try to advance reference (if it fails, someone else did or will advance it)** (points to the `pred.next.compareAndSet(curr, succ, false, false);` line)

# Add

```

public boolean add(T object) {
    while (true) {
        Window window = find(head, object);
        Node pred = window.pred, curr = window.curr;
        if (curr.object == object)
            return false;
        Node n = new Node(object);
        n.next = new AtomicMarkableReference(curr, false);
        if (pred.next.compareAndSet(curr, n, false, false))
            return true;
    }
}

```

Already there

Create new node

Install new node, else retry loop



# Wait-Free Contains

```

public boolean contains(T object) {
    boolean marked[] = new boolean[1];
    int key = object.hashCode();
    Node curr = this.head;
    while (curr.key <= key) {
        if (object == curr.object)
            break;
        curr = curr.next;
    }
    curr.next.get(marked);
    return (object == curr.object && !marked[0]);
}

```

Only difference from lazy list  
is that we get and check mark

# Lock-Free Find

```

public window find(Node head, T object) {
    Node pred, curr, succ; int key = object.hashCode();
    boolean[] marked = { false }; boolean b;
    retry: while (true) {
        pred = head; curr = pred.next.getReference();
        while (true) {
            succ = curr.next.get(marked);
            while (marked[0]) { ... }
            if ((curr.key == key && curr.object == object)
                || curr.key > key)
                return new window(pred, curr);
            pred = curr; curr = succ;
        }
    }
}
  
```

**Restart if list changes while traversed**

**Start from head**

**Move down the list**

**Get successor and mark**

**Try to remove deleted nodes**

**Otherwise advance window and loop again**

**If found object or greater key, return pred and curr**

# Lock-Free Find

```

...
while (marked[0]) {
    b = pred.next.compareAndSet(curr, succ, false, false);
    if (!b) continue retry;
    curr = succ;
    succ = curr.next.get(marked);
}
...
  
```

Try to snip out node


If predecessor's next field changed must retry whole traversal

Otherwise move on to check if next node deleted

# Summary: Lock-Free List

- **AtomicMarkableReference** atomically updates mark and reference
  - Prevents manipulation of logically-removed next pointer
- Lock-free **add()** and **remove()**
  - Remove performs logical removal, may leave node
- Wait-free **contains()** traverses both marked and removed nodes
- Physically remove marked nodes in **find()**

# Performance



# Summary

- Four “generic” approaches to concurrent data structure design
  - Fine-grained locking
  - Optimistic synchronization
  - Lazy synchronization
  - Lock-free synchronization

# “To Lock or Not to Lock”

- Locking vs. non-blocking
  - Extremist views on both sides
- Nobler to compromise, combine locking and non-blocking
  - Example: Lazy list combines blocking `add()` and `remove()` and a wait-free `contains()`
  - Blocking/non-blocking is a property of a method