

Concurrency: Multi-core Programming & Data Processing

Concurrent Queues and Stacks

Prof. P. Felber

Pascal.Felber@unine.ch

<http://iiun.unine.ch/>

Based on slides by Maurice Herlihy and Nir Shavit



Last Lecture

- Five concurrent data structure designs
 - Coarse-grained locking
 - Simple but hotspot + bottleneck
 - Fine-grained locking
 - All delayed by front thread: hotspots + bottleneck
 - Optimistic synchronization
 - Limited hotspots but two traversals
 - Lazy synchronization
 - Lazy add/remove + wait-free contains
 - Lock-free synchronization
 - Lock-free add/remove + wait-free contains

Another Fundamental Problem

- We told you about
 - Sets implemented using linked lists
- Next: **queues** and **stacks**
 - Ubiquitous data structure
 - Often used to buffer requests...
- Queue/stacks belongs to broader pool class
- Pool: similar to **set** but
 - Allows duplicates (**multiset**)
 - No membership test (no **contains()**)

Pool Flavors

- Bounded
 - Fixed capacity, good when resources an issue
- Unbounded
 - Holds any number of objects
- Blocking (remove from empty pool, add to full pool)
 - Caller waits until state changes
- Non-blocking
 - Method throws exception

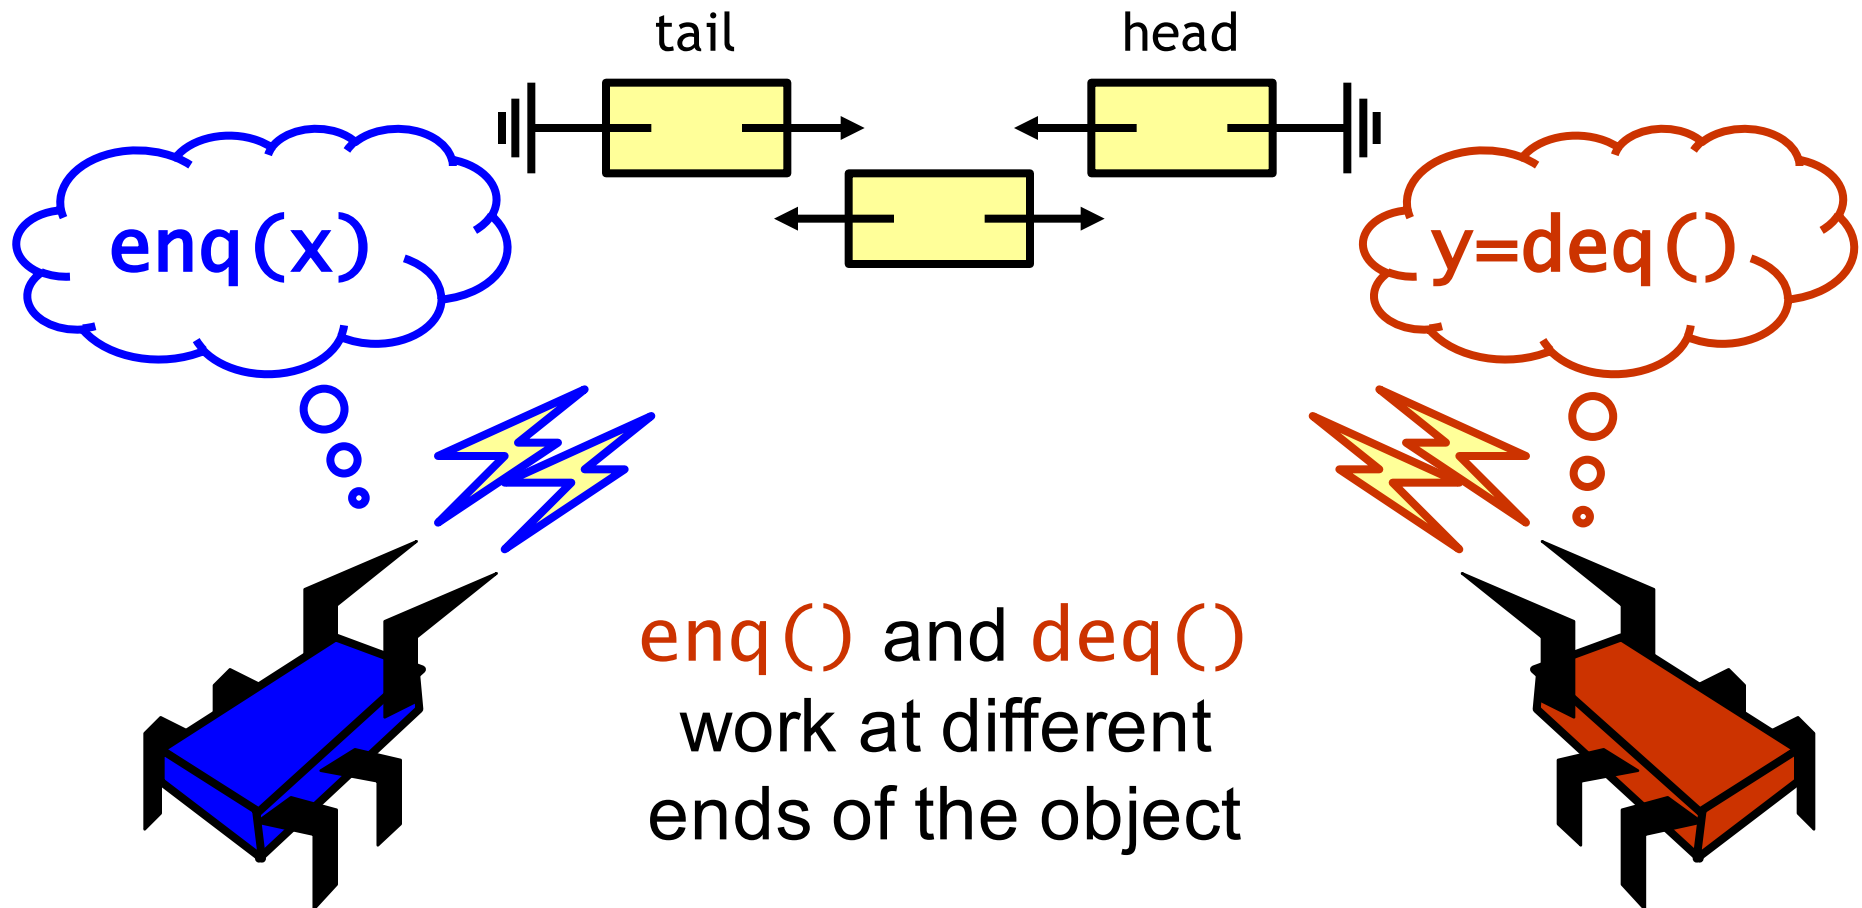
Queues & Stacks

- `add()` and `remove()`
 - Queue: `enq()` and `deq()`
 - Stack: `push()` and `pop()`
- A queue is a pool with FIFO order on enqueues and dequeues
- A stack is a pool with LIFO order on pushes and pops


This Lecture

- Bounded, blocking, lock-based queue
- Unbounded, non-blocking, lock-free stack
- Elimination-backoff stack

Queue: Concurrency



Queue: Concurrency



Java Monitor Locks

- The Java **ReentrantLock** is a monitor
 - Allows blocking on a condition rather than spinning
- Threads
 - Acquire lock
 - Release lock
 - Wait on a condition
 - Wake up threads waiting on condition

Java Monitor Locks

```

public interface Lock {
    void lock(); ➔ Acquire lock
    void lockInterruptibly()
        throw InterruptedException;
    boolean tryLock();
    boolean tryLock(long time, TimeUnit unit);
    Condition newCondition(); ➔ Conditions to wait on
    void unlock(); ➔ Release lock
}
  
```

Java Lock Conditions

```
public interface Condition {
```

```
    void await()  
        throws InterruptedException;
```

Release lock and
wait on condition

```
    boolean await(long time, TimeUnit unit)  
        throws InterruptedException;
```

```
    ...
```

```
    void signal();
```

Signal release of next thread in line

```
    void signalAll();
```

Signal release of all awaiting threads

```
}
```

The `await()` Method

`c.await();`

- Releases lock on `c` and sleeps (gives up processor)
 - Move to “waiting room” and wait to be awoken
- Upon being awoken
 - Reacquires lock and continue execution
- The awaiting thread must hold the lock prior to the call

The `signal()` Method

`c.signal();`

- Awakens one waiting thread
 - Which will reacquire lock
 - Possibly competing with other threads
- Then returns
- The signaling thread must hold the lock prior to the call

The `signalAll()` Method

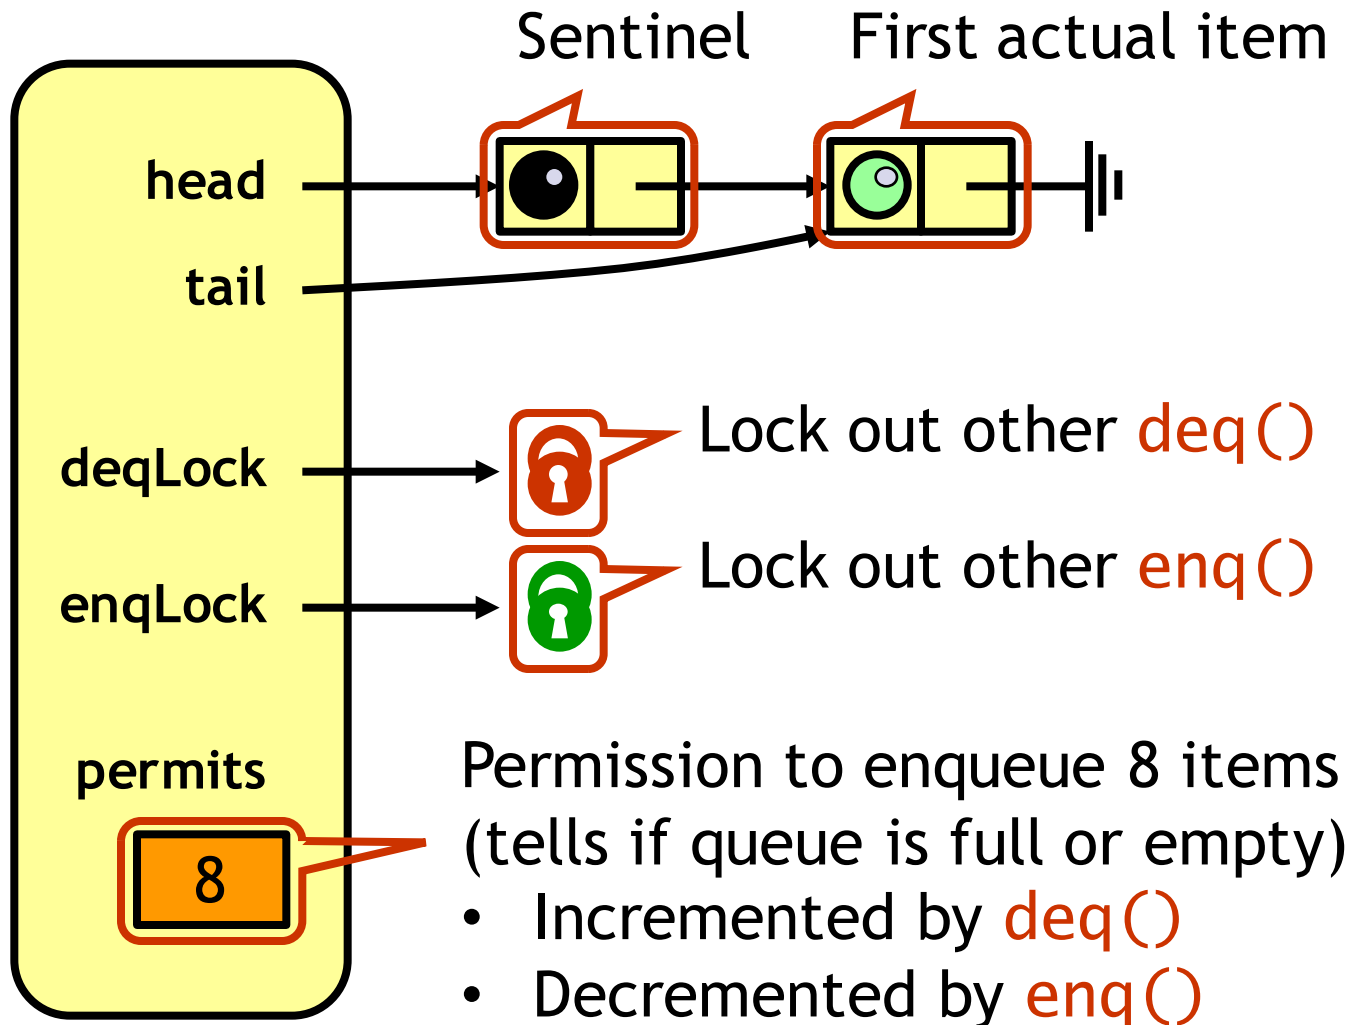
`c.signalAll();`

- Awakens all waiting threads
 - Which will reacquire lock
 - Possibly competing with other threads
- Then returns
- The signaling thread must hold the lock prior to the call


Java Synchronized Monitor

- Methods defined on class **Object**
 - **await()** → **wait()**
 - **signal()** → **notify()**
 - **signalAll()** → **notifyAll()**
- Can be called only from synchronized blocks or methods (lock must be held)

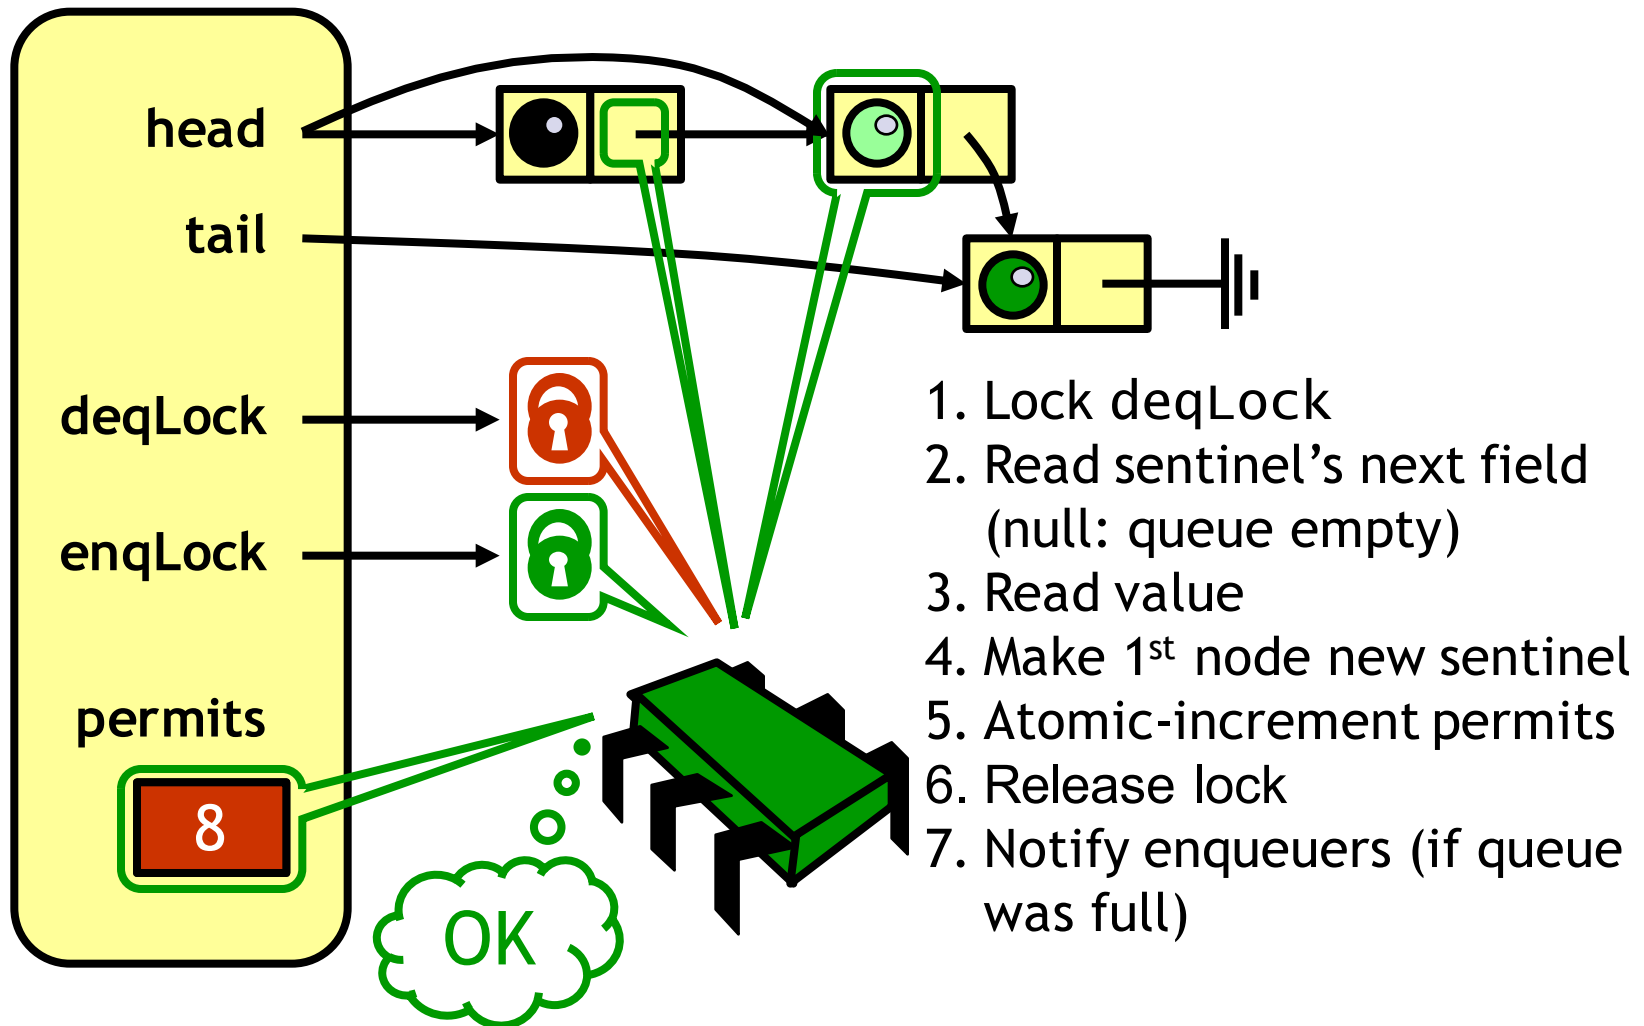
Bounded Queue



Enqueuer



Dequeuer



Bounded Queue

```

public class BoundedQueue <T> {
    ReentrantLock enqLock = new ReentrantLock();
    ReentrantLock deqLock = new ReentrantLock();
    Condition notFullCondition =
        enqLock.newCondition();
    Condition notEmptyCondition =
        deqLock.newCondition();
    AtomicInteger permits;
    Node head, tail;
    int capacity;
}

```

Enqueue & dequeue locks

Condition for threads to wait while enqueueing

Condition for threads to wait while dequeueing

Number of permits from 0 to capacity

Head and tail

Capacity of the queue

Enqueue (part I)

```

public void enq(T x) {
    boolean mustWakeDequeuers = false;
    enqLock.lock();
    try {
        while (permits.get() == 0)
            notFullCondition.await();
        Node n = new Node(x);
        tail.next = n; tail = n;
        if (permits.getAndDecrement() == capacity)
            mustWakeDequeuers = true;
    } finally {
        enqLock.unlock();
    } ...

```

Lock enqueue lock

If no permit, wait until notFullCondition becomes true then check permits again

Add a new node

If I was the enqueuer that changed queue state from empty to non-empty, need to wake dequeuers

Release enqueue lock

Enqueue (part II)

```
public void enq(T x) {
```

```
    ...
```

```
    if (mustWakeDequeuers) {  
        deqLock.lock();
```

To let the dequeuers know
that the queue is non-empty,
acquire dequeue lock

```
    try {
```

```
        notEmptyCondition.signalAll();
```

Signal all dequeuers waiting
that they can attempt
to re-acquire
dequeue lock

```
    } finally {
```

```
        deqLock.unlock();
```

Release dequeue lock




```
    }
```

```
}
```

```
}
```

Deque (part I)

```

public T deq() {
    boolean mustWakeEnqueuers = false;
    T v;
    deqLock.lock();
    try {
        while (head.next == null)  Is queue empty?
            notEmptyCondition.await();
        v = head.next.value;  Read value of first node
        head = head.next;  Make first node new sentinel
        if (permits.getAndIncrement() == 0)
            mustWakeEnqueuers = true;
    } finally { deqLock.unlock(); } ...

```

Deque (part II)

```

public T deq() {
    ...
    if (mustWakeEnqueuers) {
        enqLock.lock();
        try {
            notFullCondition.signalAll();
        } finally {
            enqLock.unlock();
        }
    }
    return v;
}

```

Return value after waking up enqueueers

The Shared Counter

- The `enq()` and `deq()` methods
 - Do not access the same lock concurrently...
 - But they still share a counter...
 - Which they both increment or decrement on every method call
 - Can we get rid of this bottleneck?

Split the Counter

- The **enq()** method
 - Decrements only
 - Cares only if value is zero
- The **deq()** method
 - Increments only
 - Cares only if value is capacity


Split Counter

- Enqueueer decrements **enqSidePermits**
 - Initially set to capacity
- Dequeueer increments **deqSidePermits**
 - Initially 0
- When enqueueer runs out
 - Locks **deqLock** (holds both locks!)
 - Transfers permits from dequeue to enqueue side
- Intermittent synchronization
 - Not with each method call


A Concurrent Stack

- `add()` and `remove()` of stack are called `push()` and `pop()`
- A stack is a pool with LIFO order on pushes and pops


Unbounded Lock-Free Stack



Unbounded Lock-Free Stack



Unbounded Lock-Free Stack



Push

```

public class LockFreeStack {
    AtomicReference top = new AtomicReference();
    public boolean tryPush(Node node) {
        Node n = top.get();
        node.next = n;
        return top.compareAndSet(n, node);
    }
    public void push(T value) {
        Node node = new Node(value);
        while (!tryPush(node))
            backoff();
    }
}
  
```

Try to push node

Read top value...

...to be new node's successor

Try to swing top to point at new node

Create new node...

...then try to push (upon failure, back off before retrying)

Pop

```
public class LockFreeStack {
```

```
...
```

```
public T pop() {
```

```
    Node n;
```

```
    while (true) {
```

If stack is empty, throw exception

```
        if ((n = top.get()) == null)
            throw new EmptyStackException();
```

```
        if (top.compareAndSet(n, n.next))
```

Try to pop
top node

```
            return n.value;
```

Success: return popped value

```
        backoff();
```

Failure: back off

```
    }
```

```
}
```

```
}
```


Lock-Free Stack




- Good
 - No locking
- Bad
 - Contention on top (add backoff)
 - No parallelism
 - ABA problem (more on that later)
- Is a stack inherently sequential?

Elimination-Backoff Stack

- How to “turn contention into parallelism”
 - Replace regular exponential-backoff...
 - ...with an alternative elimination-backoff mechanism


Observation

- After any equal number of pushes and pops, the stack stays the same




Idea: Elimination Array

- Pick a location at random in elimination array and try to match a push with a pop




Push Collides With Pop

- If push collides with pop, no need to access the stack



Push Collides With Pop

- If push does not collide with pop, access the stack



Summary

- We saw both
 - Lock-based
 - Lock-free
- Implementations of
 - Queues
 - Stacks
- Do not be quick to declare a data structure inherently sequential
 - Elimination-backoff stack can exploit parallelism