

Concurrent Systems — Exam

June 2012

Name: _____

Duration: 120 minutes — No document authorized

1.

a) What is false sharing?

b) We say that a concurrent system is “asynchronous”. What does that mean?

c) What is the difference between “safety” and “liveness” properties? Illustrate these concepts with an example.

d) What does Amdahl's law say? (You can give an informal explanation.)

e) Describe a typical cache hierarchy as found in a multi-core processor?

f) What are the semantics of read/write locks? What fairness properties are desirable for such locks?

2.

What are the “compareAndSet” and “getAndIncrement” operations? Describe their semantics informally or using pseudo-code.

Can we implement “getAndIncrement” using “compareAndSet”?

If the answer is “no”, explain why.

If the answer is “yes”, write the corresponding code and indicate if the implementation is lock-free, and/or wait-free.

3.

Consider the following 2-thread mutual exclusion algorithm seen in the course:

```
class MyLock implements Lock {
    private volatile boolean[] flag = new boolean[2];

    public void lock() {
        // Convention: i is local thread, j is other thread
        flag[i] = true;
        while (flag[j]) {}
    }

    public void unlock() {
        flag[i] = false;
    }
}
```

Does this algorithm ensure mutual exclusion?

What problem does it suffer from?

How can this problem be avoided?

4.

Consider a shared stack s and two threads τ_1 and τ_2 . Are the following histories linearizable and/or sequentially consistent? If so, write the equivalent sequential history.

Note: It could be helpful to draw a graphical representation of the histories.

a)

τ_1 $S.push(a)$

τ_1 $S:void$

τ_2 $S.push(b)$

τ_1 $S.pop()$

τ_2 $S:void$

τ_1 $S:b$

τ_2 $S.pop()$

τ_2 $S:a$

b)

τ_1 $S.push(a)$

τ_1 $S:void$

τ_2 $S.push(b)$

τ_1 $S.pop()$

τ_2 $S:void$

τ_1 $S:a$

τ_2 $S.pop()$

τ_2 $S:b$

c)

τ_1 $S.push(a)$

τ_1 $S:void$

τ_2 $S.pop()$

τ_1 $S.push(b)$

τ_2 $S:b$

τ_2 $S.pop()$

τ_2 $S:a$

d)

τ_1 $S.push(a)$

τ_1 $S:void$

τ_2 $S.pop()$

τ_2 $S:b$

τ_1 $S.push(b)$

τ_2 $S.pop()$

τ_2 $S:b$

Consider the following code that implements the well-known “double-checked locking” software design pattern.¹ Try to understand and explain the rationale of this approach. In particular, why is it better than a simpler implementation in which the `getHelper()` method would be synchronized?

This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

¹ http://en.wikipedia.org/wiki/Double-checked_locking

6.

Write a simple *lock-based* queue that supports concurrent enqueues and dequeues (i.e., it must be possible for 2 threads to operate simultaneously on both ends of the queue if it is not empty). You are free to choose the design of your implementation (blocking or non-blocking, bounded or unbounded).

[illegible]