

Concurrent Systems — Exam

June 2013

Name: _____

Duration: 120 minutes — No document authorized

1.

a) Explain informally the notion of linearizability. How does it differ from sequential consistency?

b) Explain the principle of the MESI protocol. As a reminder, MESI stands for Modified, Exclusive, Shared, Invalid.

c) Why is TTAS (test-and-test-and-set) more efficient on many types of architectures than TAS (test-and-set) for implementing a spin lock?

d) What does the `volatile` keyword guarantee in Java?

e) Rank the three following progress conditions from the strongest to the weakest: *lock-freedom*, *obstruction-freedom*, *wait-freedom*.

f) What is the main disadvantage of work shedding as compared to work stealing?

2.

What are the “getAndSet” and “getAndDecrement” operations? Describe their semantics informally or using pseudo-code.

Can we implement “getAndDecrement” using “getAndSet”?

If the answer is “no”, explain why.

If the answer is “yes”, write the corresponding code and indicate if the implementation is lock-free, and/or wait-free.

3.

Consider the following 2-thread mutual exclusion algorithm seen in the course:

```
public class MyLock implements Lock {
    private volatile int victim;

    public void lock() {
        // Convention: i is local thread
        victim = i;
        while (victim == i) {}
    }

    public void unlock() {}
}
```

Does this algorithm ensure mutual exclusion?

What problem does it suffer from?

How can this problem be avoided?

4.

Considering a shared queue Q and two threads τ_1 and τ_2 , write four sample histories that respectively have the following properties:

- a) Linearizable and sequentially consistent
- b) Not linearizable but sequentially consistent
- c) Linearizable but not sequentially consistent
- d) Neither linearizable nor sequentially consistent

If there are no valid histories with some of the requested properties, explain why.

a)

b)

c)

d)

Consider the bounded lock-free queue below. We assume that integer overflows are not an issue.

How can we fix it while preserving the lock freedom property of the implementation?

[illegible]

6.

A barrier is a synchronization aid that allows a set of threads to all wait for each other to reach a common point. A barrier is initialized with a thread count (**n**). All threads that call the **cross()** method block until **n** threads have called the method. At that point, all threads are allowed to proceed.

A barrier can be easily implemented in Java using a lock and a condition variable, whose interfaces are summarized below as a reminder:

```
interface Lock {
    void lock();
    void unlock();
    Condition newCondition();
}

interface Condition {
    void await();
    void signal();
    void signalAll();
}
```

Complete the code below to implement a barrier in Java.

```
class Barrier {
```

```

    _____
    _____
    _____

```

```
    public Barrier(int n) {
```

```

    _____
    _____
    _____

```

```
    }
    public void cross() {
```

```

    _____
    _____
    _____
    _____
    _____
    _____
    _____
    _____
    _____

```

```
    }
}
```